

NCA Assistant

User Manual

Version 1.9

Non-Compartmental Pharmacokinetic Analysis

Bioequivalence • Drug Interaction Studies • Power & Sample Size

Radboud Applied Pharmacometrics

Radboudumc, Nijmegen, The Netherlands

*Designed and developed by Rob ter Heine (hospital pharmacist and clinical pharmacologist,
Radboudumc)*

Manual version 1.9, for NCA Assistant v1.8.0

Not validated for your use: qualifying your installation is your responsibility (see Intended Use and Your Data).

Table of Contents

Table of Contents.....	2
How to Use This Manual.....	5
Intended Use and Your Data	6
What the App Does Not Do	6
Getting Started.....	8
Opening the Public App.....	8
Installing the App on Your Own Computer	8
The Hub Screen	8
Getting Help Inside the App	9
Saving Your Work	9
Tutorials	11
Tutorial 1: A First Analysis Without a File.....	11
Tutorial 2: Uploading Data with BLQ Results	11
Tutorial 3: All Subjects at Once (Theophylline)	12
Tutorial 4: Bioequivalence	12
Tutorial 4b: Covariates in a Parallel-Group Study	13
Tutorial 4c: Reference-Scaled Bioequivalence (Highly Variable Drug)	15
Tutorial 5: Planning a Study	17
Tutorial 6: Partial AUCs.....	17
Tutorial 7: Figures	17
Answer Key	18
Preparing Your Data and Study Types	20
General Format Rules	20
Scenario 1: Single Subject	20
Scenario 2: Multiple Subjects, Same Dose	21
Scenario 3: Multiple Subjects, Different Doses	21
Scenario 4: Standard 2-Period Crossover (BE)	21
Scenario 5: Replicate Designs and Paired Comparison	22
Scenario 6: Parallel-Group BE	22
Scenario 7: Steady-State PK.....	22
Scenario 8: Drug–Drug Interaction Study.....	23
Scenario 9: CDISC ADNCA Dataset.....	23
Scenario 10: IV Bolus and IV Infusion.....	24
Scenario 11: Single Ascending Dose and Dose Proportionality.....	24
Scenario 12: Multiple Dosing and Accumulation	24
Scenario 13: Food Effect.....	25
Scenario 14: Metabolites and Several Analytes.....	25
Scenario 15: Absolute and Relative Bioavailability	25

Studies with Three or More Treatments	25
BLQ Handling	25
Common Mistakes.....	26
Using the App	28
Path 1: Plan a Study.....	28
Path 2: Upload & Check Data	29
Path 3: Visualize Data	30
Path 4: One Subject at a Time	32
Path 5: All Subjects (Batch).....	32
Path 6: Bioequivalence.....	34
Analysing Drug–Drug Interaction Studies	38
Complete Analysis Record	39
Reading Your Results	41
The Data Quality Report.....	41
The Parameter Table	41
Summary Statistics.....	41
The Bioequivalence Table and Forest Plot	42
Judging the Terminal Phase.....	43
Checking Plausibility	44
Working in a Regulated Environment	45
Before You Rely on the App.....	45
How a Controlled Installation Works	46
Signing In and Staying Signed In	48
A Study, Step by Step	49
Reviewing the Audit Trail and Inspections	53
Good Practice: Data You Can Trust.....	55
What the App Does and What Your Organisation Does	55
Upgrades, Data Protection and Problems.....	56
Tutorial 8: From Analysis to Approved Record.....	57
Key Points	57
Installing NCA Assistant for Regulated Use.....	59
What You Need Before You Start	59
Step 1: Plan and Approve	60
Step 2: Prepare the Server.....	61
Step 3: Install the Release and the Validated Packages	62
Step 4: Configure Shiny Server and nginx	62
Step 5: Create the Controlled Directory and the Settings	63
Step 6: Protect the Audit Trail and the Records	63
Step 7: Issue the Accounts.....	64

Step 8: Qualify the Installation.....	64
Step 9: Go Live.....	65
Running the Installation.....	65
Changes and Upgrades	66
When Something Goes Wrong.....	66
Retiring an Installation.....	66
Checklist.....	67
Introduction to Non-Compartmental Analysis	68
What Is Non-Compartmental Analysis?	68
The Concentration–Time Curve	68
Key Pharmacokinetic Parameters	68
Why NCA Is Used So Widely	69
When to Use NCA vs. Compartmental Modelling	69
Theoretical Background	70
Trapezoidal Methods for AUC Calculation	70
Terminal Phase Characterisation and Half-Life Estimation.....	70
BLQ Handling: Pharmacological Rationale	71
Steady-State Considerations.....	72
Partial AUCs: Interpolation and the End of the Interval.....	72
Statistical Methods.....	73
The Two One-Sided Tests (TOST) Procedure.....	73
The ANOVA Model for Crossover Designs	73
Fixed-Effects vs. Mixed-Effects Models	73
Confidence Interval Construction	73
Design-Specific Models.....	74
Reference-Scaled Approaches (EMA ABEL and FDA RSABE).....	74
Drug–Drug Interaction Analysis.....	75
Power and Sample Size for Bioequivalence	76
Partial AUCs in Bioequivalence.....	76
Regulatory Guidelines	77
Appendices	78
Appendix A: Output Parameters.....	78
Appendix B: Glossary.....	82
Appendix C: Recognised Column Names.....	87
Appendix D: Troubleshooting	87
Appendix E: Setting Up and Administering a Controlled Installation.....	90
Appendix F: References.....	93
Appendix G: Document Control.....	95

How to Use This Manual

Few readers need all of this manual. Find your row below and start there; the rest is reference.

- **New to pharmacokinetics or NCA**, for example as a PhD student: read Introduction to Non-Compartmental Analysis first, then do Tutorials 1 to 3. Look up any unfamiliar term in Appendix B, the glossary.
- **Analysing a study**: Preparing Your Data and Study Types, the path you use in Using the App, and Reading Your Results. Tutorials 2, 3 and 6 show the steps.
- **Bioequivalence or drug interaction statistics**: Path 6 in Using the App, Statistical Methods, and Tutorials 4 and 6.
- **Reviewer or inspector on a controlled installation**: Working in a Regulated Environment, and Tutorial 8.
- **System owner or IT**: Installing NCA Assistant for Regulated Use, with Appendix E as its reference.
- **Quality assurance**: Intended Use and Your Data, Working in a Regulated Environment, and Steps 1 and 8 of Installing NCA Assistant for Regulated Use.

Labels of buttons and fields are written as the app shows them, so you can match the text to the screen.

Intended Use and Your Data

NCA Assistant is for pharmacokineticists doing non-compartmental analysis, bioequivalence testing and study planning. Average bioequivalence for every design, plus EMA ABEL and FDA RSABE for replicate designs. The FDA scaled test for narrow therapeutic index drugs is for planning only. Studies run in several groups (ICH M13A, 2.2.3.5) are not modelled; a group column in the file is ignored. The public instance and a standard installation have no audit trail, electronic signature or access control; a controlled installation on your organisation's server adds them (Working in a Regulated Environment). The public instance on shinyapps.io is for evaluation and training, with synthetic or pseudonymised data. For regulated work, install a tagged release on your own system and qualify it there with the validation package. Responsibility for the analysis and its conclusions stays with the user. Not for dosing decisions for individual patients.

Your data. On the public instance, uploads are processed on shinyapps.io servers run by Posit PBC (USA). Upload only synthetic, example or anonymised data there. Pseudonymised trial data are still personal data under the GDPR. Sending them to a third-party host needs agreements your organisation must have in place, and may breach sponsor confidentiality. For real study data, run the app on your own computer. A controlled installation shows its own notice instead: there, uploads stay on your organisation's server.

Status. NCA Assistant is not validated for your use. The validation package in the repository (a URS, an IQ/OQ/PQ protocol and a test script with 627 automated and 84 manual tests) is a starting point: the developer's test run shows that the code does what the tests check, on the developer's computer. Qualifying the installation you use is your responsibility; the chapters Working in a Regulated Environment and Installing NCA Assistant for Regulated Use explain how.

Supplier, support and warranty. The app is free, open-source software under the GNU General Public License, version 3 or later, written by Rob ter Heine (Radboud Applied Pharmacometrics, Radboudumc, Nijmegen). It comes without any warranty and without a support agreement. Problems and questions go to the issue tracker at github.com/Robterheine/NCAassistant/issues.

What the App Does Not Do

Use other software for the following. The app refuses some of these, but not all, so check before you start.

Not supported	What to use or do instead
The FDA scaled test for narrow therapeutic index drugs, and the FDA model for replicate designs	Software that implements them. Plan a Study plans the NTID test. For a replicate study the app offers EMA ABEL and FDA RSABE.
Studies with three or more treatments in one analysis (for example a three-way food-effect or DDI crossover, a Williams design)	Analyse two treatments at a time (see Studies with Three or More Treatments), knowing the residual then comes only from those periods, or use software that fits the full design.
Dose-proportionality statistics (the power model)	Use the dose-normalised parameters and plots here, and fit the power model elsewhere.
Absolute or relative bioavailability between routes or strengths in one comparison	Run the NCA per route with dose normalisation and compute F outside the app.

Urine data (amount excreted, renal clearance, excretion rate)	Other NCA software.
Sparse or destructive sampling (one sample per subject, Bailer's method)	Other NCA software, or a population approach.
Therapeutic drug monitoring data with irregular dosing, superposition, several infusion durations in one run, baseline-corrected endogenous compounds	A population PK model, or other NCA software.
Dosing decisions for individual patients	Not an intended use.

Getting Started

Opening the Public App

Go to robterheine.shinyapps.io/NCAassistant in Chrome, Firefox, Edge or Safari. There is nothing to install. The public instance is for evaluation and training: read Intended Use and Your Data before you upload anything. It disconnects after a period without activity, and it keeps nothing between sessions.

Installing the App on Your Own Computer

For confidential data and for regulated work, run the app on your own computer. It needs R 4.1 or newer; RStudio is optional but makes the steps easier.

1. Install R from cran.r-project.org, and optionally RStudio Desktop from posit.co.
2. Download a release from github.com/Robterheine/NCAassistant/releases: under Assets, choose "Source code (zip)", and unzip it. Use a tagged release rather than the latest code, so you know which version you run.
3. In RStudio, open `install_and_run.R` from the unzipped folder and click Source. Without RStudio, open a terminal in that folder and run `Rscript install_and_run.R`.
4. The first time, the script installs about twenty R packages from CRAN, which can take several minutes. Later starts take a few seconds.
5. The app opens in your web browser. To stop it, press Esc in the RStudio console, or Ctrl+C in the terminal.

For a qualified installation, add `--validated` (`Rscript install_and_run.R --validated`), or run `NCA_VALIDATED <- TRUE` before sourcing the script. It then installs the package versions the release was validated with, listed in `validation/renv.lock`, instead of the latest ones from CRAN. Run it from the folder that contains `app.R`.

If the installation fails. Read the last lines the script printed: they name the package that failed. On Windows, a package that needs compiling may need Rtools (cran.r-project.org/bin/windows/Rtools); on a Mac, the Xcode command-line tools (`xcode-select --install` in a terminal). Then run the script again; packages that installed are not installed twice.

The Hub Screen

The hub has six cards, numbered in the order of a typical study. Plan a Study needs no data. Upload & Check Data is where every analysis starts; the other four paths use the data processed there.

Card	Needs uploaded data	What you get
1. Plan a Study	No	Sample size or power, with plots
2. Upload & Check Data	(starts here)	A quality report and the processed data for paths 3 to 6
3. Visualize Data	Yes	Concentration-time figures and a Figure Record

4. One Subject at a Time	No (manual entry) or yes	NCA of one profile, with the half-life inspector
5. All Subjects (Batch)	Yes	NCA of every profile, summary statistics, a Complete Analysis Record
6. Bioequivalence	Yes	NCA, ANOVA, confidence intervals, a verdict and a Complete Analysis Record

Once data are processed, the navigation bar shows a badge such as "12 subjects, 132 obs". Home, in the bar under the navigation, brings you back to the hub. Processing a new file clears every result of the previous one, and a result is also cleared when you change a setting it depends on: run the analysis again after a change. Settings stay when you leave a page and come back.

NCA Assistant Statistical Methods Data Guide About User Manual v1.8.0

Non-Compartmental Analysis Assistant

Pharmacokinetic analysis made accessible.
Rob ter Heine, Radboud Applied Pharmacometrics

1. Plan a Study

No data needed

Calculate how many subjects you need for a bioequivalence or pharmacokinetic study, and estimate the probability of success (power).

- Standard bioequivalence and the EMA and FDA scaled methods, for crossover, parallel and replicate designs.

2. Upload & Check Data

Start here if you have data

Upload your concentration-time data (CSV or Excel), map the columns, handle values below the limit of quantification, and run an automatic quality check.

- Auto-detects column names, study design, and data problems.

3. Visualize Data

Explore your data

Create publication-ready concentration-time plots: individual spaghetti plots and geometric mean summary curves with error bars.

- Export to PNG, PDF, or SVG at journal-submission resolution.

4. One Subject at a Time

Inspect and analyse PK profiles subject by subject. Ideal for dose-escalation studies and quality review.

Interactive half-life adjustment for each profile.

5. All Subjects (Batch)

Run PK analysis on all subjects at once. Get summary statistics, parameter tables, and export results.

Includes population summary and individual review.

6. Bioequivalence

Compare two formulations (Test vs. Reference). Runs NCA, computes 90% confidence intervals, and concludes equivalence.

Complete BE workflow from data to confidence intervals and forest plot.

New to PK analysis? Check the [Data Preparation Guide](#) first, then start with Upload & Check Data.

*Look for the **i** buttons throughout the app — they explain every setting in plain language.*

Figure 1. The hub, with the six paths numbered in the order of a typical study.

Getting Help Inside the App

Every setting that needs explaining has an **i** button next to it; click it for a short explanation in plain language. The top navigation bar has three more sources. The Data Guide shows how to lay out the data for each study type, with example files to download. Statistical Methods describes every calculation, with references, and ends with Example Methods Paragraphs to adapt for a report or thesis. About lists the packages, their versions and licences, and the version history.

Saving Your Work

The app keeps nothing between sessions. Before you close the browser tab, or leave the public app idle, download what you need: the results as Excel or CSV, the figure, or the Complete Analysis Record, which holds the data, the settings and a script that reproduces the analysis. After a disconnect, upload the file again and click Process Data. On a controlled installation, your records and the audit trail are kept on the server; the uploaded file and the results are not.

Tutorials

Ten short tutorials, from a first analysis without a file to a bioequivalence study with partial AUCs, with baseline covariates (Tutorial 4b) and with a reference-scaled approach (Tutorial 4c), and one on reviewing a record on a controlled installation (Tutorial 8, in the chapter Working in a Regulated Environment). Each gives the steps with the labels the app shows, then questions. The answers are in the Answer Key at the end of this chapter, rounded as the app displays them. The example files are in the Data Guide (top navigation bar), under Example Datasets. The first time, do the tutorials in order: later ones assume you know the upload steps. Do the tutorials on the public app, on your own computer or on a training installation, never on a controlled production server, where every practice run would be recorded.

New to NCA? Read Introduction to Non-Compartmental Analysis first: the tutorials assume you know what C_{max}, AUC and a half-life are.

Tutorial 1: A First Analysis Without a File

1. On the hub, click 4. One Subject at a Time.
2. Under **Where is your data?**, choose **Enter data manually (type values directly)**. The Time points and Concentrations boxes already hold an example profile: eight samples from 0 to 24 h, one value per line.
3. Click **Use This Data**.
4. Under **Analysis Settings**, keep Dose 100, Dose unit mg, Time unit h, Conc unit ng/mL and **Route of administration Oral / IM / SC (extravascular)**. Click **Run PK Analysis**.
5. Read the **PK Parameters** card, then the **Half-Life Calculation** card, which shows the points used for the terminal phase.

Questions: 1) What are C_{max} and T_{max}? 2) What is the AUC to last point? 3) What is the half-life, and on how many points does it rest? 4) What is the apparent clearance CL/F, and in which unit?

Tutorial 2: Uploading Data with BLQ Results

This tutorial uses example_blq.csv: eight subjects after a 100 mg oral dose, with results below the limit of quantification written as "<0.5".

1. On the hub, click 2. Upload & Check Data. Keep **What kind of file?** at **Simple table (one row per sample)**. Under **Or use an example**, choose example_blq.csv and click **Load example**. To see how the file is laid out, click **Download** next to it.
2. Check the column mapping: Subject ID is Subject, Time is Time, Concentration is Concentration.
3. Leave **LLOQ value** at 0 and click **Process Data**. The app stops: it found BLQ text but no LLOQ. It shows "LLOQ auto-detected: 0.5" with a button **Set LLOQ to 0.5**.
4. Click **Set LLOQ to 0.5**. The button only fills in the field; check that it reads 0.5, keep **BLQ handling rule** at the default (Rule 1), and click **Process Data** again. The quality report now lists the 22 BLQ entries as a warning.

5. Go to 5. All Subjects (Batch). Keep **Dose information** at **Same dose for all subjects**, enter 100 as **Dose given to each subject**, and check the units (mg, h, ng/mL).

6. Under **Partial AUCs (optional)**, set **Number of intervals** to 1, with Start 0 and End 1. Click **Run Analysis on All Subjects**.

Questions: 1) Which subjects have a lag time, and how long is it (Lag Time, in **All Parameters (per profile)**)? 2) Subject 5 has a BLQ result at 8 h between two measurable ones. What is its time of last measurable concentration? 3) What is subject 2's Partial AUC 0–1? Now go back to Upload & Check Data, choose Rule 6 (Before first quantifiable → LLOQ/2, rest → zero), click **Process Data**, and run the batch analysis again. What is subject 2's Partial AUC 0–1 now, and why did it change while the lag time did not?

Tutorial 3: All Subjects at Once (Theophylline)

1. Under **Or use an example**, choose example_theoph.csv and click **Load example**. Map Subject ID to Subject, Time to Time and Concentration to conc. The app also maps the Dose column, which holds mg/kg. Leave **LLOQ value** at 0 and click **Process Data**; the badge shows 12 subjects, 132 obs.

2. Go to 5. All Subjects (Batch). Because a Dose column is mapped, **Dose information** is set to **Doses differ by subject or period (from Dose column in data)**. Switch it to **Same dose for all subjects** and enter 320 as **Dose given to each subject**, with Dose unit mg, Conc unit mg/L and Time unit h. The Dose panel warns that these doses look like doses per kg. They are: the file gives mg/kg. Questions 1 to 7 use the nominal 320 mg that every subject was meant to receive, and question 8 shows what the individual doses change.

3. Keep **Route of administration** at **Oral / IM / SC (extravascular)** and **Trapezoidal method** at **Linear-up / Log-down**. Click **Run Analysis on All Subjects**.

Questions, with where to look: 1) Mean Cmax (**Summary Statistics**, Arithmetic Mean)? 2) Geometric mean AUC to last point (Summary Statistics, Geometric Mean)? 3) Median Tmax? 4) Which subject has the longest half-life, and how long is it (**All Parameters (per profile)**, sort by Half-Life)? 5) Geometric CV of AUC to infinity (observed) (Summary Statistics, Geometric CV (%))? 6) Is that longest half-life reliable? Tick **Show all parameters** and look at Points Used for Half-Life and AUC % Extrapolated (observed). 7) The note above the results names a subject with a pre-dose concentration above 5% of Cmax. Which subjects have a pre-dose concentration above zero, and what does that mean? 8) Optional: on the Upload page, choose Wt under Dose per kg: weight column and click Process Data. In All Subjects (Batch), Dose information now reads Doses differ by subject or period; set Conc unit to mg/L and run the analysis again. What is subject 9's CL/F now?

Tutorial 4: Bioequivalence

1. Load example_be_crossover.csv the same way (**Or use an example**, then **Load example**). Map Subject ID, Time and Concentration, and also Treatment / Formulation to Treatment, Period to Period and Sequence to Sequence. Click **Process Data**.

2. Go to 6. Bioequivalence. In **Step 1: PK Analysis Settings**, enter Dose 100 with units mg, h and ng/mL, and keep the route at **Oral / IM / SC (extravascular)**.

3. In **Step 2: Bioequivalence Settings**, check that **Study design** is **2x2 standard crossover (TR | RT)**, that **Reference treatment** is Reference, and that **Statistical model** is

Fixed effects (all terms fixed). Keep **Log-transform (recommended)**, the **Parameters to compare** (Cmax and AUC to last point), **Confidence interval** 90 and the limits 80 and 125.

4. Click **Run Complete BE Analysis**.

Questions: 1) What are the ratio and 90% CI for Cmax (tab **Test vs. Reference Comparison**, columns Estimate and 90% CI Lower and Upper)? 2) And for AUC to last point? 3) Is the product bioequivalent for both, and what does the note above the table say? 4) What is the within-subject CV for Cmax? It is not in the table on screen: click **Download CI Table (CSV)** and take the Residual Variance (MSE), then $CV = 100 \cdot \sqrt{(\exp(MSE) - 1)}$; or open Plan a Study, where a button offers the CV of this analysis.

Tutorial 4b: Covariates in a Parallel-Group Study

This tutorial uses example_be_parallel_covariates.csv: 40 subjects in two parallel groups of 20, one oral dose of 100 mg, samples up to 48 h, and the baseline columns Age, Weight and Sex. Suppose the protocol named body weight as a covariate before the study began. A covariate is a characteristic of each subject that may move the PK parameter. Use baseline values only, known before the first dose: a value measured after dosing can be changed by the treatment itself. Choose the covariates before you look at any result.

1. On the hub, click **2. Upload & Check Data**. Under **Or use an example**, choose example_be_parallel_covariates.csv and click **Load example**. Check that **Subject ID** is Subject, **Time** is Time, **Concentration** is Conc, **Treatment / Formulation** is Treatment and **Dose** is Dose. Leave **Period**, **Sequence** and **Dose per kg: weight column** empty: a column mapped to one of these roles cannot be a covariate. Click **Process Data**.

2. Go to **6. Bioequivalence**. In **Step 1: PK Analysis Settings** keep **Dose information** at **Different doses (from Dose column)**, with units mg, h and ng/mL and the route **Oral / IM / SC (extravascular)**.

3. In **Step 2: Bioequivalence Settings**, check that **Study design** is **Parallel groups** and that **Reference treatment** is Reference. Open **Covariates (optional)**. Under **Baseline characteristics to adjust for**, choose Weight and nothing else. The line below the box reads "Weight : numeric", and the section title shows the badge "1 selected".

4. Click **Run Complete BE Analysis**.

5. Read the line above the table (Adjusted for: Weight), the table with its two **Unadjusted 90% CI** columns, and the card **Group balance on the covariates**.

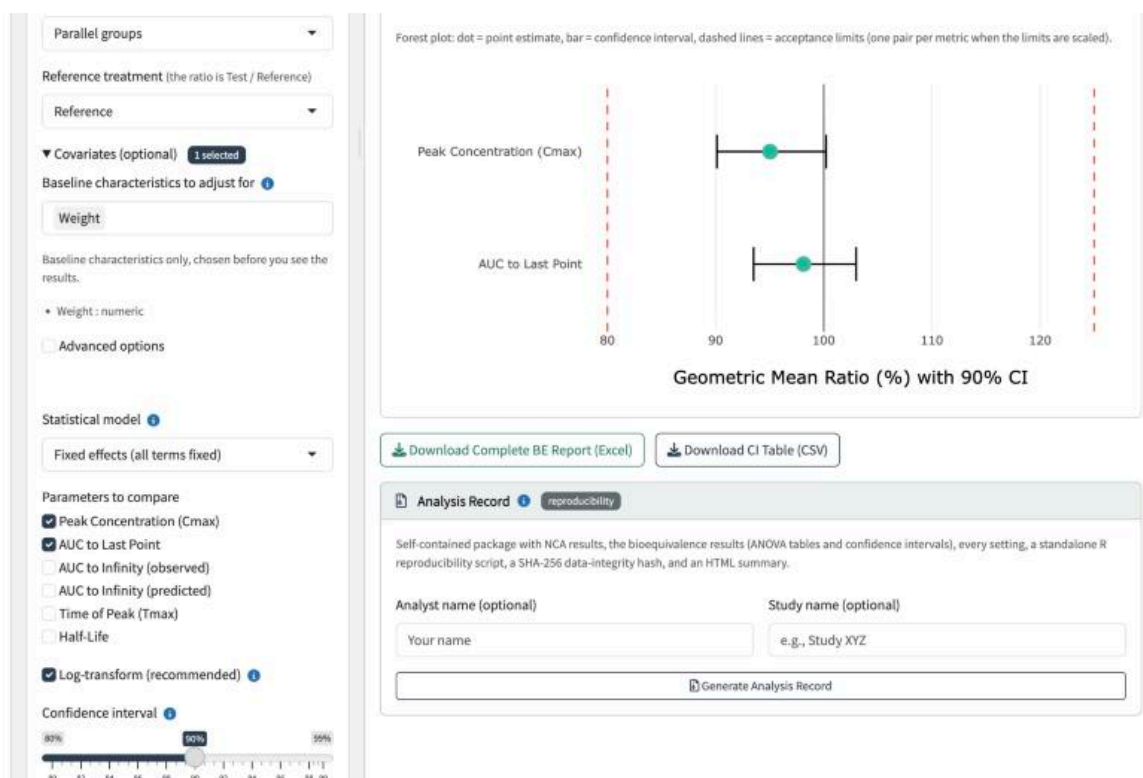


Figure 4b.1. The Covariates (optional) section with Weight chosen, beside the forest plot of the adjusted analysis (Tutorial 4b).

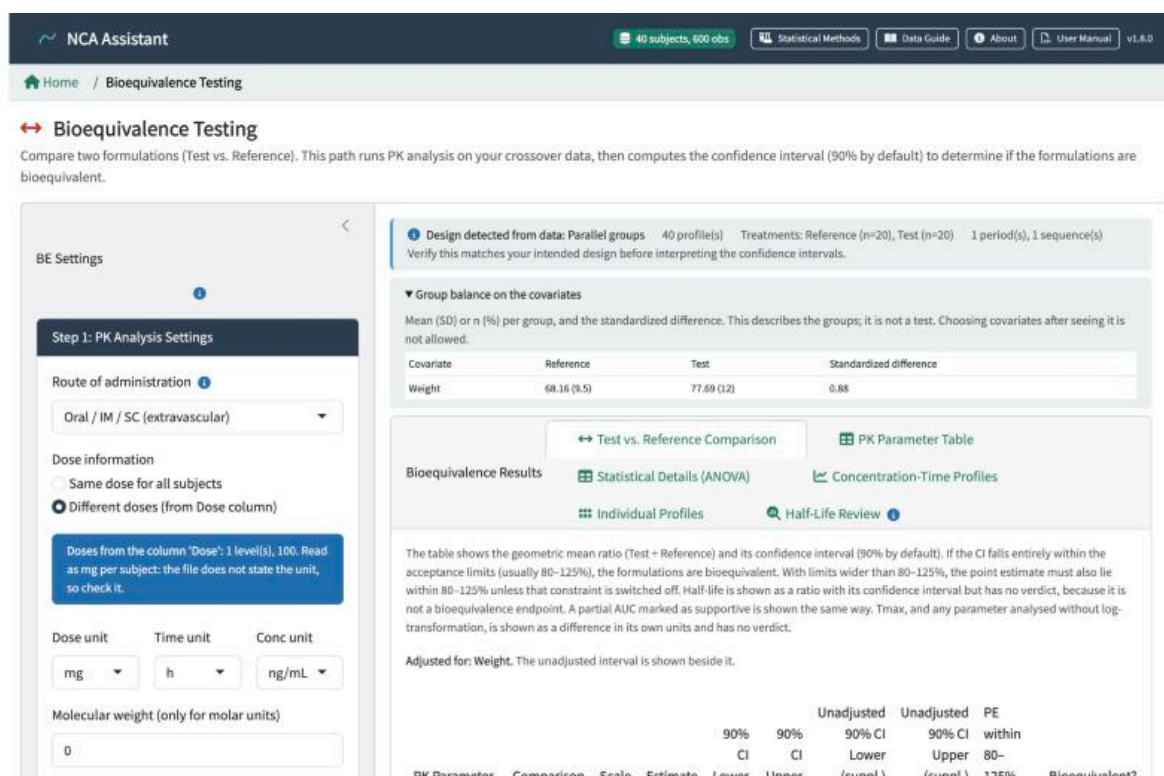


Figure 4b.2. The group balance card, and the adjusted table with the unadjusted interval beside it (Tutorial 4b).

Questions: 1) What are the ratio and the 90% CI for Cmax adjusted for Weight, and what do the unadjusted columns show? 2) Is the product bioequivalent for Cmax with the adjustment, and without it? And for AUC to last point? 3) What is the mean weight in each group, and how large is the standardized difference? 4) By how much does Cmax change per kilogram

of body weight? Click **Download Complete BE Report (Excel)** and open the sheet BE_Covariates. 5) How many residual degrees of freedom does the model have, and how does the CV fall with the adjustment? Use the columns Degrees of Freedom, Residual Variance and Unadjusted residual variance in **Download CI Table (CSV)**, and $CV = 100 \cdot \sqrt{(\exp(MSE) - 1)}$. 6) Why is it wrong to pick Weight after you have seen this table?

Tutorial 4c: Reference-Scaled Bioequivalence (Highly Variable Drug)

This tutorial uses example_be_replicate_hvd.csv: 32 subjects in a 2×2×4 full replicate (TRTR | RTRT), one oral dose of 100 mg, samples up to 48 h. The Reference is highly variable for C_{max}. A product is called highly variable when the within-subject variability of the Reference exceeds about 30%. With fixed limits of 80–125% such a product needs many subjects, even when Test and Reference are the same product. EMA ABEL and FDA RSABE scale the limits, or the criterion, to the variability of the Reference. The approach is a declaration made in the protocol before the study: the app never switches by itself. Below you run all three approaches one after another, to compare them. In a real study you run and report only the one the protocol names.

1. Load example_be_replicate_hvd.csv as in Tutorial 4b. Check that **Subject ID**, **Time**, **Concentration** (Conc), **Treatment / Formulation**, **Period**, **Sequence** and **Dose** are mapped, and click **Process Data**.
2. Go to **6. Bioequivalence** and use the dose settings of Tutorial 4b. Check that **Study design** is **2×2×4 full replicate (TRTR | RTRT)**, **Reference treatment** is Reference, **Statistical model** is **Fixed effects (all terms fixed)**, the parameters are C_{max} and AUC to last point, the confidence interval is 90 and the limits are 80 and 125. Keep **Acceptance approach** at **Standard (average bioequivalence)** and click **Run Complete BE Analysis**. Note the ratio, the 90% CI and the verdict of both parameters.
3. Set **Acceptance approach** to **FDA RSABE (reference-scaled)**. The results are cleared. Click **Run Complete BE Analysis** again. Read the table, including the columns Route, s_WR (FDA contrasts), Limits and Criterion bound, and the lines under it that say how each metric was judged.

2x2x4 full replicate (TRTR | RTRT)

Reference treatment (the ratio is Test / Reference)

Reference

Acceptance approach ⓘ

FDA RSABE (reference-scaled)

Note: By default this app performs average bioequivalence (ABE) with the acceptance limits you enter. For a replicate design you can instead choose EMA ABEL or FDA RSABE under Acceptance approach. That choice is your declaration before the run, and the app never switches by itself. The results also show CV_{WR} and the limits it would imply, for information. The FDA scaled test for narrow therapeutic index drugs is in Plan a Study only.

Statistical model ⓘ

Fixed effects (all terms fixed)

Parameters to compare

☒ Peak Concentration (Cmax)

☒ AUC to Last Point

☐ AUC to Infinity (observed)

☐ AUC to Infinity (predicted)

☐ Time of Peak (Tmax)

☐ Half-Life

☒ Log-transform (recommended) ⓘ

Confidence interval ⓘ

90%

Download Complete BE Report (Excel)

Download CI Table (CSV)

Figure 4c.1. The Acceptance approach selector on a replicate design, set to FDA RSABE, with the note that the choice is a declaration before the run (Tutorial 4c).

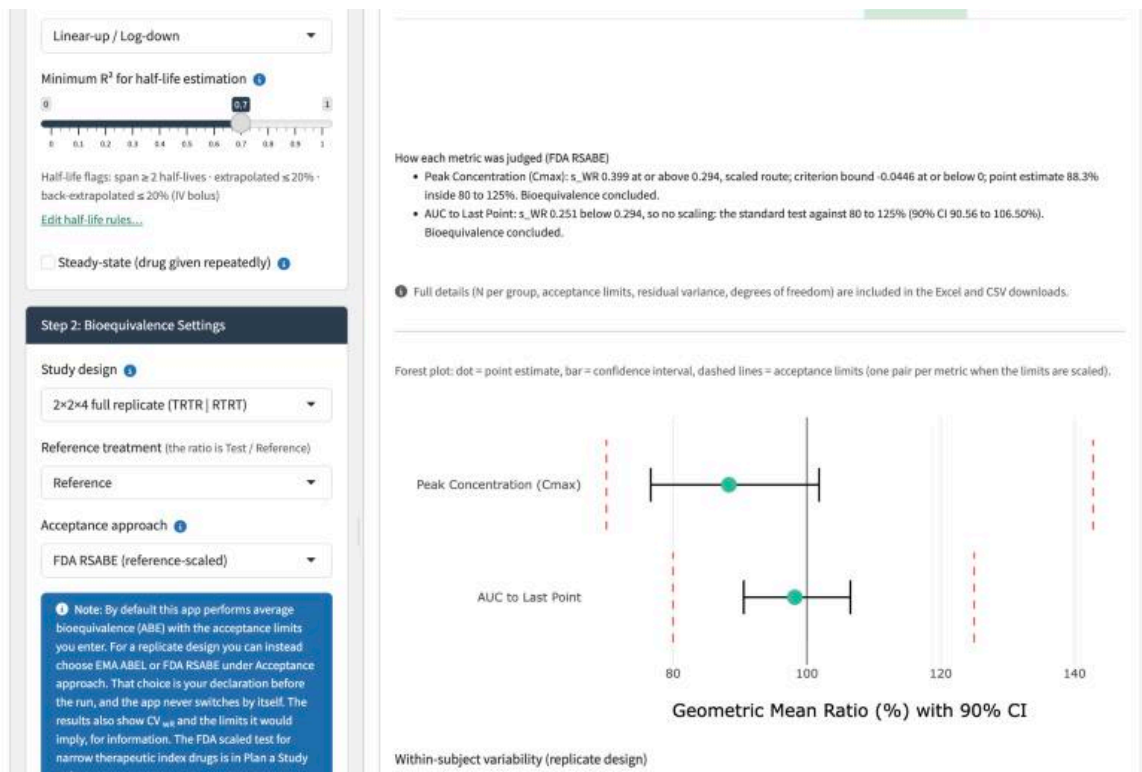


Figure 4c.2. The lines that say how each metric was judged under RSABE, and the forest plot with one pair of limits per metric (Tutorial 4c).

4. Set **Acceptance approach** to **EMA ABEL (Cmax, widened limits)** and run once more. Read the same table and lines.

Questions: 1) What are the ratio, the 90% CI and the verdict for C_{max} and for AUC to last point with the standard approach? 2) With RSABE: which route did C_{max} take, what are s_{WR}, the limits and the criterion bound, and what is the verdict? 3) Which route did AUC to last point take with RSABE, and why? 4) With ABEL: what are CV_{WR} and the limits for C_{max}, and which limits apply to AUC? 5) Which two conditions must hold for RSABE to conclude bioequivalence? 6) If the protocol named RSABE, which result do you report, and why may you not report the most favourable one?

Tutorial 5: Planning a Study

1. On the hub, click 1. Plan a Study.

2. Step 1: What do you want to calculate?: How many subjects do I need? (find sample size). Step 2: What kind of study?: Standard bioequivalence (80–125% limits, any regimen). Step 3: Study design: 2×2 standard crossover (TR | RT).

3. Under Parameters, enter **Within-subject variability (CV %)** 25, **Expected Test/Reference ratio (%)** 95 and **Target power** 80. Click **Calculate**.

Questions: 1) How many subjects are needed? 2) Switch Step 1 to **What is the chance of success for my planned study? (calculate power)**, enter 30 as **Total number of subjects in your study**, and click **Calculate**. What is the power? 3) Switch back to the sample size, set the CV to 35 and click Calculate. How many subjects are needed now? A result is cleared when you change an input: click Calculate again after every change.

Tutorial 6: Partial AUCs

Use the crossover data of Tutorial 4 (sampled at 0, 0.25, 0.5, 1, 1.5, 2, 3, 4, 6, 8, 12 and 24 h).

1. In 5. All Subjects (Batch), enter Dose 100 (mg, h, ng/mL). Under **Partial AUCs (optional)**, set **Number of intervals** to 2. Interval 1: Start 0, End 1, and tick **Also report C_{max} and T_{max} in this interval**. Interval 2: Start 4, and tick **End at the last measurable concentration (t)**. Click **Run Analysis on All Subjects**.

2. Go to 6. Bioequivalence, with the settings of Tutorial 4. Each path has its own intervals, so enter them again in its Partial AUCs panel, with the C_{max} box ticked for interval 1. Set **Role in bioequivalence** to **Pivotal (verdict)** for interval 1 and **Supportive (no verdict)** for interval 2. The partial metrics appear under **Parameters to compare**, already ticked. Click **Run Complete BE Analysis**.

Questions: 1) Work out Partial AUC 0–1 of subject 1 under Test by hand, from the first four samples (0, 12.3, 45.6 and 89.2 ng/mL), and compare it with the batch table. 2) What are the ratio and 90% CI for Partial AUC 0–1? 3) What does the table show for Partial AUC 4–t, and why? 4) What is the ratio for C_{max} 0–1?

Tutorial 7: Figures

These three examples use Visualize Data (card 3). No analysis settings are needed.

A quick look after upload. Load example_theoph.csv (Tutorial 3) and open 3. Visualize Data. The individual profiles plot shows 12 lines, one per profile, on a linear axis. Switch **Y-axis scale** to Log₁₀: a straight tail points to mono-exponential elimination.

A summary plot for a crossover. Load `example_be_crossover.csv` with the Treatment column mapped (Tutorial 4), open Visualize Data and choose **Summary Plot**. It overlays one geometric mean curve per treatment. This is a visual check, not a test: only the confidence interval in Bioequivalence says whether the products are bioequivalent. For a report, go to **Export Figure** in the sidebar, choose PDF, 7 × 5 in and 300 DPI, and click **Download Figure**.

Dose-normalised profiles. With `example_theoph.csv`, whose Dose column differs per subject, tick **Normalise to dose (show C/Dose)** and colour the lines by Subject. Change the Y-axis label to "C/Dose (mg·L⁻¹ per mg/kg)". With dose-proportional pharmacokinetics the normalised profiles overlap.

Answer Key

Tutorial 1. 1) C_{max} 28.3 ng/mL at 1 h. 2) 156.8 ng/mL·h. 3) 4.607 h, from the last 5 points (adjusted R² 0.9724). 4) 616.7 L/h: 100 mg divided by an AUC to infinity of 162.2 ng/mL·h, converted to litres.

Tutorial 2. 1) Subjects 2 and 5 (0.5 h), and 4 and 7 (0.25 h). 2) 24 h: under Rule 1 the BLQ result at 8 h becomes 0, but the measurable values after it still count, and the fall to zero adds its area to AUC. 3) 2.238 ng/mL·h under Rule 1 and 2.394 under Rule 6. Rule 6 sets the BLQ samples during the lag to LLOQ/2, which adds area before the first measured concentration. The lag time stays 0.5 h, because a value set by a BLQ rule is not a measured concentration.

Tutorial 3. 1) 8.759 mg/L. 2) 98.65 mg/L·h. 3) 1.135 h. 4) Subject 1, 14.3 h. 5) 28.43%. 6) Not very: it rests on 3 points, and 31.5% of AUC to infinity is extrapolated (flagged in amber). Theophylline's half-life is usually about 8 h; report this value with a flag, or leave it out of the descriptive statistics as your analysis plan says. 7) Subjects 1 (0.74 mg/L, 7% of its C_{max}), 7 and 10. Drug before the dose means carry-over from earlier dosing or an endogenous level; in a bioequivalence study ICH M13A excludes a period with a pre-dose value above 5% of C_{max} from the primary analysis. The single dose of 320 mg is also not right for every subject: subject 9 received 3.10 mg/kg × 86.4 kg = 268 mg, so its CL/F is about 19% too high. 8) 2.747 L/h, against 3.281 L/h with 320 mg. Subject 9 received 268 mg, so the nominal dose overstated its CL/F by 19%.

Tutorial 4. 1) 108.62% (106.27–111.02%). 2) 111.25% (108.95–113.60%). 3) Yes, both intervals lie within 80.00–125.00%. The note says that the study has 6 evaluable subjects, fewer than the 12 ICH M13A asks for, so the verdict is a statistical result only: this study would not be accepted as a pivotal study. 4) The residual variance is 0.000315, so the CV is about 1.8%. The example data are simulated with very little noise.

Tutorial 4b. 1) 95.04% (90.12–100.22%). The unadjusted columns show 77.76–92.12%; only the limits are shown, not the unadjusted ratio. 2) With the adjustment: yes. Without it: no, the lower limit of 77.76% is below 80%. AUC to last point passes either way (adjusted 98.13%, 93.51–102.99%; unadjusted interval 85.20–96.80%). 3) 68.2 kg on Reference and 77.7 kg on Test (68.16 and 77.69), a standardized difference of 0.88. That is a large imbalance. In these simulated data weight drives the volume of distribution, so the heavier Test group had lower concentrations before the adjustment. 4) The Weight coefficient is -0.0122 on the log scale, which is about 1.2% less C_{max} for each kilogram. 5) 37 residual degrees of freedom (40 - 2 - 1). The CV falls from 16.0% (unadjusted residual variance

0.0252) to 9.1% (residual variance 0.00823). 6) Choosing a covariate after you have seen the results lets you search for the interval that passes, and the type I error is no longer controlled. The declaration has to come first. The Analysis Record lists the covariates, but the app cannot know when you chose them.

Tutorial 4c. 1) Cmax 88.35% (77.52–100.69%): not bioequivalent, because the lower limit is below 80%. AUC to last point 98.20% (90.56–106.50%): bioequivalent. 2) The scaled route. s_WR is 0.399, at or above 0.294; the limits are 70.04–142.78%; the criterion bound is –0.0446, at or below 0; the point estimate of 88.35% lies within 80–125%. RSABE is met. The interval in the table, 76.67–101.81%, comes from the FDA contrasts and differs a little from the standard 77.52–100.69%. 3) The standard route: s_WR is 0.251, below 0.294, so nothing is scaled and the ordinary test against 80–125% applies. The interval of 90.56–106.50% lies within those limits. 4) CVwR is 41.5%, above 30%, and the limits are 73.84–135.43%. The 90% CI of 77.52–100.69% lies within them and the point estimate lies within 80–125%, so Cmax is bioequivalent. AUC keeps 80–125%, because the EMA widens Cmax only. 5) The criterion bound must be at or below 0 (it is the 95% upper confidence bound of the scaled criterion), and the point estimate must lie within 80–125%. 6) The one the protocol declared before the study. Running several approaches and keeping the one that passes defeats the purpose of a pre-specified analysis; you ran all three here only to see how they differ.

Tutorial 5. 1) 28 subjects. 2) 83.4%. With two subjects more than the 28 needed for 80%, the power is a little higher. 3) 52 subjects. A higher CV widens the confidence interval, so more subjects are needed to keep it within the limits.

Tutorial 6. 1) 42.48 ng/mL·h (1.5375 + 7.2375 + 33.7); the concentrations rise throughout, so both trapezoidal methods give the same value. 2) 113.55% (110.90–116.25%), bioequivalent. 3) 112.54% (110.10–115.05%) without a verdict, because the interval is supportive. 4) 110.69% (108.31–113.11%).

Tutorial 8. 1) Two runs on the same data, the second with "r2adj_threshold: 0.7 → 0.9"; the label says there were 2 runs before the record. 2) The times and concentrations were typed in by hand, so the reviewer must check every value against the source document. 3) Three: all runs on the same data before the record are listed, including those that led to the rejected record. 4) The analyst's name and user ID with the UTC time of creation, and the meaning "I have checked the settings of this analysis against the protocol, its results and its audit trail, and I approve this record." 5) The Approve and Reject buttons do not appear for one's own record, and the server refuses such a signature.

Preparing Your Data and Study Types

General Format Rules

The app expects data in **long format**: one row per sample, with a subject identifier, the sampling time and the measured concentration. Accepted files are CSV (.csv), text (.txt, .tsv) and Excel (.xlsx, .xls). The app suggests a column mapping from the column names (Appendix C); check it before you process the data. When another column could also be the time or the concentration, a note names it. CDISC ADNCA datasets are uploaded differently, see Scenario 9.

Time is the time since the dose of that profile, with the pre-dose sample at 0. One pre-dose sample per profile drawn just before the dose is analysed at time 0, and its uploaded time is kept: a negative time no earlier than 10% of the profile's duration, such as -0.25 h in a 24 h profile. Any other negative time is refused, because it would make the lag time negative and add the time before the dose to AUC.

Actual or nominal times. Use the actual sampling times for the NCA, as ICH M13A (2.2.2.2) asks: a sample taken at 1.1 h enters the calculation at 1.1 h. Nominal times (1 h) belong in the mean plot of Visualize Data, which averages per time point. Keep a second copy of the file with nominal times for that figure.

Decimal Separators and Regional Formats

Choose the delimiter (comma, semicolon or tab) and the decimal point (period or comma) in the upload settings. With a decimal comma, values such as 12,5 are read correctly, also in a concentration column that contains BLQ text such as "<0,5". Excel files need no settings, because Excel stores numbers as numbers. In a decimal-comma file a point can only separate thousands, so 12.500 is read as 12500. Any other point in a number, including 0.250, is refused, and the message names the column.

Units and Doses

Use one concentration unit, one time unit and one dose unit for the whole file, and select them in the app. They set the conversion factor for clearance and volume, so a wrong unit gives wrong CL/F and Vz/F while Cmax and AUC look normal. A unit column with more than one unit (even ng/mL and µg/L) is refused; convert the values first. Molar concentrations (for example nmol/L) with a mass dose (mg) need the molecular weight, which you enter in the analysis settings. When the file states its units in a unit column (such as ConcUnit or Time_Unit, or the CDISC variables AVALU and RRLTU), the app pre-selects them and refuses a selection that contradicts them.

A Dose column holds the amount given, in the dose unit you select, or a dose per kg. For a dose per kg, choose the body-weight column (kg) under **Dose per kg: weight column** on the Upload page, and the app uses dose × weight as the amount. Read as mg, a dose in mg/kg gives clearance and volume per kilogram labelled L/h and L; the Dose panel in the analysis says which unit it assumes and warns when the doses look per kg.

Scenario 1: Single Subject

To upload a file, include a Subject column; it may hold the same value on every row. Without it the upload cannot tell which rows belong together, and the app picks another column as

Subject. Manual entry in One Subject at a Time needs only Time and Concentration (Tutorial 1).

Subject	Time	Concentration
P1	0	0
P1	0.5	4.2
P1	1	18.5
P1	2	12.1
P1	4	6.3
P1	8	2.8
P1	24	0.4

Scenario 2: Multiple Subjects, Same Dose

Subject	Time	Concentration
S001	0	0
S001	1	22.5
S001	24	0.4
S002	0	0
S002	1	28.1
S002	24	0.6

Scenario 3: Multiple Subjects, Different Doses

Add a Dose column if subjects received different doses (dose escalation, weight-based dosing), or if a subject received different doses in different periods. The app takes the dose of each profile (subject, treatment and period) from this column, so every period is analysed with its own dose. For dose proportionality, see Scenario 11. With weight-based dosing the column may hold the dose per kg if you also map the body-weight column (see Units and Doses).

Scenario 4: Standard 2-Period Crossover (BE)

Six columns: Subject, Treatment, Period, Sequence, Time, Concentration. Time restarts at 0 in each period: it is the time since the dose of that period, not since the first dose of the study (the app refuses profiles that do not start near time zero). The ratio is Test ÷ Reference. In Bioequivalence you choose which treatment is the Reference; the app pre-selects it only when the name is unambiguous (R, Ref, Reference, Comparator, Innovator, Originator or RLD). Map the Period column. Without it a crossover is still analysed but gets no verdict, because the period effect cannot be separated from the treatment effect. Columns named APERIOD, Occasion or OCC are recognised as Period. A Visit column is not, so map it yourself if it holds the period. Subject IDs must be unique in the whole study: if each sequence numbers its subjects from 1, different people would be merged into one subject, so the app refuses the file.

If the protocol asks for a partial AUC, take a sample at every cutoff time. A cutoff between two samples is interpolated, so its value then depends on the trapezoidal method. The interval itself comes from the protocol or the product-specific guidance, never from the data.

Subject	Treatment	Period	Seq	Time	Conc
S01	Test	1	TR	0	0
S01	Test	1	TR	1	24.5
S01	Reference	2	TR	0	0
S01	Reference	2	TR	1	22.8

Scenario 5: Replicate Designs and Paired Comparison

Same six-column format. Replicate designs: 2×2×3 full replicate (sequences TRT and RTR), 2×3×3 partial replicate (TRR, RTR, RRT) and 2×2×4 full replicate (TRTR, RTTR). Period goes to 3 or 4 and a treatment appears more than once per subject; each administration is analysed as its own profile. Paired comparison: all subjects received the treatments in the same order, so Sequence has one value (or can be left out). For a highly variable drug, the Acceptance approach selector in Path 6 offers EMA ABEL (all three replicate designs) and FDA RSABE (2×3×3 and 2×2×4); example_be_replicate_hvd.csv (Tutorial 4c) shows both.

Scenario 6: Parallel-Group BE

Four columns: Subject, Treatment, Time, Concentration. No Period or Sequence needed. Baseline characteristics can be added as extra columns (for example Age, Weight, Sex), with the same value in every row of a subject, and used as covariates in Path 6. Keep numbers as plain numbers: text such as “72 kg” or a decimal comma is refused. The example is example_be_parallel_covariates.csv (Tutorial 4b).

Scenario 7: Steady-State PK

Same layout as Scenario 2, with one dosing interval per subject. Time 0 is just before the dose of that interval, and the pre-dose concentration is above zero. Include the pre-dose sample and sample up to the end of the dosing interval (τ). Tick “Steady-state (drug given repeatedly)” and enter τ in the app. If a profile has no pre-dose sample at time 0, the app assumes 0 there and says so, because AUC τ is then too low.

Two concentrations describe the low point of the interval. C_{min} is the lowest observed concentration from 0 to τ , the pre-dose sample included; with an absorption lag it can fall after the dose. C τ is the observed sample nearest to τ within a trough window (by default $\pm 10\%$ of τ , at most 1 time unit; you can change it under the dosing interval), and its time is in the column Time of C τ . It is never interpolated, and it is left empty, with a note, when no sample falls within the window. ICH M13A lists both for steady-state studies, and the EMA modified-release guideline asks for C τ . Both can be compared in Bioequivalence. AUC to infinity, its % extrapolated and the volume at steady state are single-dose quantities and are left empty.

A worked example. A twice-daily drug, 100 mg every 12 hours ($\tau = 12$ h), sampled at steady state. The trough was drawn 15 minutes late:

Subject	Time	Concentration (ng/mL)
S1	0	20
S1	0.5	48
S1	1	81
S1	2	92
S1	4	71
S1	6	52
S1	8	37
S1	12.25	21

With Steady-state ticked and $\tau = 12$, the trough at 12.25 h falls within the trough window (± 1 h), so C_{τ} is 21 ng/mL, with 12.25 h in the column Time of C_{τ} and a note that it is the sample nearest to τ . $C_{\min}$ is 20 ng/mL, the pre-dose sample. AUC_{τ} is 622.7 ng/mL·h, interpolated to 12 h between the samples at 8 and 12.25 h, so nothing is extrapolated (AUC Within Dosing Interval % Extrapolated is 0). C_{avg} is $622.7 / 12 = 51.9$ ng/mL, the fluctuation is $(92 - 20) / 51.9 \times 100 = 138.7\%$, and CL/F is $100 \text{ mg} / AUC_{\tau} = 160.6 \text{ L/h}$. Had sampling stopped at 8 h, C_{τ} would stay empty and part of AUC_{τ} would be extrapolated with λ_z . The app would say so in a note, and flag the profile if more than 20% were extrapolated.

A partial AUC at steady state has to lie within 0 and τ . AUC_{τ} is completed to τ with λ_z when the last sample comes earlier; a partial AUC you enter yourself is not, so it stays empty for such a profile.

Scenario 8: Drug–Drug Interaction Study

A drug interaction (DDI) study measures the effect of a co-administered drug (the perpetrator) on the pharmacokinetics of the victim drug. Most are crossovers in which each subject receives the victim drug alone and with the perpetrator, or fixed-order studies in which the victim is given alone first and with the perpetrator after a washout.

Use the same six columns as a crossover BE study (Scenario 4). Label the victim drug alone as “Reference” and the victim drug with the perpetrator as “Test”, or choose the Reference treatment in the app. The ratio is then (with perpetrator) $\div$ (alone), the fold-change in exposure. For a single-dose victim, add AUC to infinity to the parameters to compare, as ICH M12 (5.1.1) asks; at steady state compare AUC_{τ} , $C_{\max}$ and $C_{\min}$ (ICH M12, 5.1.1). Analysing Drug–Drug Interaction Studies, in the chapter on using the app, explains how to read the result.

Scenario 9: CDISC ADNCA Dataset

If your data come as a CDISC ADNCA dataset (columns such as USUBJID, AVAL, PARAMCD, NRRLT, ARRLT and ANL01FL), do not convert them by hand. The steps:

1. On the Upload page, set What kind of file? to CDISC ADNCA dataset and upload the .csv or .xlsx file.
2. Read the summary, and choose the time to use: ARRLT or MRRLT (actual times) for the NCA, NRRLT (nominal times) for a mean figure; AFRLT, time since the first dose, is refused. With several analytes or matrices, choose the one to analyse; Set pre-dose times (negative ARRLT) to 0 does what it says.

3. Click Convert dataset. The app applies ANL01FL, drops records with PCSTAT "NOT DONE", copies PCLLOQ into the LLOQ field, and refuses datasets it cannot convert safely, such as derived records (DTYPE), several units, or AVAL missing without a BLQ result.
4. Check the column mapping and the LLOQ, then click Process Data.

The choices are stored in the Analysis Record. SAS transport (.xpt) files are not read: converters/ADNCA_TO_FLAT.md in the repository gives a two-line R recipe to convert them to CSV. The same conversion is available outside the app as converters/adnca_to_flat.R.

Scenario 10: IV Bolus and IV Infusion

For an IV bolus, choose IV Bolus (injected into vein at once) as the route. The concentration at time 0 is C₀, the value just after the injection, not a pre-dose value: the app sets aside samples at or before time 0 and back-extrapolates C₀ from the first two concentrations when they decline. Clearance and volume are reported without /F. AUC % Back-Extrapolated shows how much of AUC lies before the first sample; above about 20% the AUC is unreliable.

For an IV infusion, time 0 is the start of the infusion, and you enter the Infusion duration in the time unit of the data. One duration applies to every profile in a run, so split a file with infusions of different lengths and analyse each part separately. MRT is corrected for the infusion time (minus half the duration). Bioequivalence does not take an infusion.

Scenario 11: Single Ascending Dose and Dose Proportionality

Give every subject a Dose column and put the dose level in a Treatment column (for example "10 mg", "30 mg", "100 mg"). All Subjects (Batch) then gives summary statistics per dose group. Tick Calculate dose-normalised parameters to get C_{max}, AUC and the partial AUCs divided by each profile's dose, and use Visualize Data with Normalise to dose (show C/Dose) to compare the profiles. Overlapping dose-normalised profiles suggest dose proportionality but do not test it. The power model ($\ln AUC = a + \beta \cdot \ln \text{Dose}$, with a confidence interval for β ; Smith et al., 2000) is not in the app; fit it with other software on the exported table.

Scenario 12: Multiple Dosing and Accumulation

A multiple-dose file usually holds a first-dose profile (Day 1) and a steady-state profile (Day N) for every subject. Code the day as Period, with time reset to 0 at each of the two doses, and the day also in a Treatment column (for example "Day 1" and "Day 7"). The steady-state setting applies to every profile in a run, so:

1. For the Day N parameters, analyse the Day N profiles as steady state, with τ .
2. For the accumulation ratio, use Bioequivalence with Paired comparison, Day N as Test and Day 1 as Reference, Steady-state ticked with the same τ , and compare AUC Within Dosing Interval. Day 1 then gives AUC from 0 to τ after the first dose, and the ratio with its confidence interval is the accumulation ratio. The note about a missing pre-dose sample does not apply to Day 1. Sample Day 1 up to τ as well. Its AUC from 0 to τ is then a single-dose partial AUC; if the last Day 1 sample comes before τ , the app extrapolates it with λz and reports the share in AUC Within Dosing Interval % Extrapolated.

3. Document that steady state was reached with consecutive trough samples, as ICH M13A (2.2.2.2) asks.

Scenario 13: Food Effect

A two-way food-effect study is a 2×2 crossover: the fed administration as Test and the fasted one as Reference (Scenario 4). Compare C_{max}, AUC to last point and AUC to infinity, and report the median difference in T_{max}. A 90% confidence interval within 80.00–125.00% supports the absence of a food effect. A three-way study (for example Test fasted, Test fed and Reference fasted) has three treatments, see Studies with Three or More Treatments.

Scenario 14: Metabolites and Several Analytes

Analyse one analyte per run. Filter a flat file to one analyte first, or choose it during the ADNCA upload. A flat file with a column such as Analyte, PARAMCD or Matrix that changes within a subject is refused. A metabolite-to-parent ratio needs both in molar units: enter the molecular weight of each in its own run, and compute the ratio from the exported tables.

Scenario 15: Absolute and Relative Bioavailability

Bioequivalence compares two treatments with one route and without dose normalisation, so it cannot compare IV with oral dosing, or strengths with different doses. Run All Subjects (Batch) per route with dose normalisation, and compute $F = (AUC_{oral} / dose_{oral}) \div (AUC_{IV} / dose_{IV})$ from the exported table. For a confidence interval of F, use other software.

Studies with Three or More Treatments

Bioequivalence takes exactly two treatments. For a three-way crossover (food effect, a three-arm DDI study, a Williams design), filter the file to two treatments per comparison and analyse each pair on its own. The residual variance then comes only from the periods of those two treatments, not from the ANOVA of the whole design that a statistical analysis plan usually specifies, so the intervals differ from a full-design analysis. When the plan asks for the full model, use other software.

BLQ Handling

Rule	In the app	What it does	When to use it
Rule 1 (default)	Standard: pre-first = 0, post-last = missing	BLQ before the first and between measurable values = 0; after the last = missing	The usual convention for NCA and bioequivalence (see below)
Rule 2	All BLQ → zero	Every BLQ = 0	When the protocol sets all BLQ values to zero
Rule 3	All BLQ → excluded	Every BLQ = missing; the curve then starts from (0, 0) and runs across the removed samples	Sensitivity analyses

Rule 4	Before and between measurable → LLOQ/2, after the last → missing, pre-dose → 0	BLQ before and between measurable values = LLOQ/2; after the last = missing; pre-dose = 0; a profile without any measurable value = missing, apart from its pre-dose sample	Sensitivity analyses; needs a justification in the protocol
Rule 5	Before peak → zero, after peak → excluded	BLQ before C _{max} = 0; after C _{max} = missing	When the protocol specifies it
Rule 6	Before first quantifiable → LLOQ/2, all other BLQ → zero (also in a profile with no quantifiable value)	BLQ after dosing and before the first measurable value = LLOQ/2; pre-dose, later BLQ and a profile without any measurable value = 0	When the protocol imputes during absorption; it adds area during a lag

Every assay has a lower limit of quantification (LLOQ): the lowest concentration it can measure reliably. Samples below it are reported as BLQ (below the limit of quantification). The app recognises “<0.5” (also “<0,5”), BLQ, BQL, BLOQ, ND (not detected) and NQ (not quantifiable), in any capitals. Set the LLOQ above 0 so the chosen rule is applied to these entries; each rule works per profile, in time order. NS (no sample), N/A and “missing” are treated as missing values, not as BLQ. If you click Process Data with the LLOQ left at 0, the app stops. When the entries look like “<0.5”, it shows the value it found and a button Set LLOQ to 0.5, which fills in the field (Tutorial 2). When they give different limits, say “<0.5” and “<0.1”, it lists both and offers no button, since it cannot tell which is right: take the LLOQ from the assay report. Different limits usually mean a data-entry error or a mixed export, and the quality report warns about them.

ICH M13A (2.2.2.2) treats BLQ values as zero in PK parameter calculations and leaves them out of the half-life. The app never uses a value set by a rule for the half-life or the lag time. Rule 2 sets every BLQ value to zero. Rule 1 does so before and between measurable values but leaves values after the last one missing, which gives the same AUC to last point. With Rules 3 to 6 the Bioequivalence page reminds you that the protocol has to justify the rule.

In the processed data the app adds a BLQ_flag column marking every value the rule set. It appears in the data preview, and the partial AUC notes and the bioequivalence table use it to report how much of a result rests on such values. A column of that name in your file is replaced.

The rule matters most for partial AUCs. An early window often holds BLQ samples, which Rule 1 sets to 0 and Rule 6 to LLOQ/2, and a late window over a plateau near the LLOQ can rest mainly on values the rule produced. The app says so when more than half of the samples in a window come from the BLQ rule, or when fewer than three measurable concentrations are left in it.

Common Mistakes

Wide format: subjects as columns instead of rows. **Time not reset:** in a crossover, time restarts at 0 in every period. **Negative times:** only one pre-dose sample per profile, just

before the dose, is accepted. **0 for a missing sample:** leave the cell empty, or write NS.

Duplicate time points: usually several analytes, matrices or periods in one column; filter the file, or map the Treatment and Period columns. **LLOQ left at 0:** processing is blocked when BLQ text is present. **Dates or clock times as time:** the Time column must be elapsed time since the dose. **Mixed or wrong units:** select the units that match the file; they drive the clearance and volume conversion. **Dose in mg/kg:** map the body-weight column (Units and Doses). **Subject IDs that restart in each sequence:** give every subject its own ID.

A cutoff with no sample near it: the protocol asks for AUC 0–30 min, but the samples are at 15 and 60 min. The app interpolates the concentration at 30 min, so the value depends on the trapezoidal method.

Excluding a subject by deleting rows: the Analysis Record then holds only the edited file. Leave data out in the app instead, with a reason (Leaving Out Data, and Exclusions, Outliers and Deviations in the chapter Working in a Regulated Environment).

Using the App

This chapter describes each path, with the labels the app shows. The Tutorials chapter takes you through them step by step; Reading Your Results explains the tables.

Path 1: Plan a Study

Calculates the sample size, or the power for a given number of subjects, with the PowerTOST package. The page has four cards. **Step 1: What do you want to calculate?** chooses between **How many subjects do I need? (find sample size)** and **What is the chance of success for my planned study? (calculate power)**. **Step 2: What kind of study?** chooses the method: standard average bioequivalence, highly variable drugs (EMA ABEL or FDA RSABE), or narrow therapeutic index drugs (FDA NTID). **Step 3: Study design** offers the designs of Bioequivalence except the paired comparison: 2×2, 2×2×3, 2×3×3, 2×2×4 and parallel groups for standard bioequivalence; the replicate designs for ABEL and RSABE; only 2×2×4 for NTID. **Parameters** holds the numbers. Click **Calculate**; a result is cleared when you change an input, so click Calculate again after every change.

The CV the planner asks for depends on the design and the method:

Design and method	CV asked for	Field label
Crossover or replicate, standard bioequivalence	Within-subject CV	Within-subject variability (CV %)
Parallel groups	Total CV (between plus within subjects)	Total CV, between + within subjects (CV %)
ABEL, RSABE, NTID	Within-subject CV of Test and of Reference	Within-subject CV of the Test product (CV %) and of the Reference product (%)

Enter the expected Test/Reference ratio (the default follows the method, as in PowerTOST: 95% for standard bioequivalence, 90% for highly variable drugs, 97.5% for narrow therapeutic index drugs) and, for sample size, the target power. The significance level (0.05) and, for standard bioequivalence, the acceptance limits can be changed under Advanced: acceptance limits & significance level. After a log-transformed bioequivalence analysis of your own data, a button such as Use the Peak Concentration (C_{max}) within-subject CV from my BE analysis (1.8%) offers its CV; with several metrics, Metric from the BE analysis chooses which. It offers the within-subject CV of a crossover for a crossover design, the total CV of a parallel analysis for a parallel design, and CV_{wT} and CV_{wR} of a replicate analysis for the scaled methods. The spread of C_{max} between subjects in batch results is not a within-subject CV, so the app does not offer it. The results are in three tabs: Power vs. Number of Subjects, Sample Size vs. Variability, and Calculation Summary.

For a pivotal partial AUC, plan with the within-subject CV of that partial AUC. An early window is often more variable than total AUC, so a CV taken from AUC gives too few subjects.

The planner makes no record. For a protocol, write down the method, design, CVs, expected ratio, power, significance level, the resulting N, and the versions of the app and PowerTOST (About page), and keep a screenshot of the Calculation Summary.

Path 2: Upload & Check Data

Choose **What kind of file?** (**Simple table (one row per sample)** or **CDISC ADNCA dataset**), upload it, and set the Delimiter and Decimal point for a text file. Check the column mapping (Subject ID, Time, Concentration, and where present Treatment / Formulation, Period, Sequence and Dose, and **Dose per kg: weight column** for a dose per kg), set the **LLOQ value** and the **BLQ handling rule**, and click **Process Data**. The app runs more than 30 quality checks: errors block processing, warnings do not. Some checks stop data the app cannot analyse safely: a CDISC-style file uploaded as a simple table, mixed units, dates or clock times in the Time column, profiles that do not start near time zero (time since the first dose), several profiles stacked in one column (duplicate times), subject IDs that restart in each sequence, a decimal point in a decimal-comma file, and a negative time other than one pre-dose sample just before the dose. A second dose inside one profile, or a multi-period file without a mapped Period, gives a warning. Processing a new file clears all previous results, and your exclusions after asking first. **Or use an example** loads one of the six example files with **Load example**; **Download** saves a copy so you can see how it is laid out. Afterwards a panel links to the next steps: Visualize Data, Single Subject, Batch Analysis and Bioequivalence. Reading Your Results explains the quality report.

The screenshot shows the 'Upload & Check Your Data' page in the NCA Assistant application. At the top, there's a navigation bar with 'NCA Assistant' and links to 'Statistical Methods', 'Data Guide', 'About', 'User Manual', and 'v1.8.0'. Below the navigation bar, the page title is 'Upload & Check Your Data'. A sub-header reads: 'Upload your concentration-time data, map the columns, and run a quality check. Once processed, your data is available for all analysis paths.'

The main content area is titled 'Upload PK Data'. It includes a warning: 'Not sure if your file is ready? Check the [Data Preparation Guide](#) for format instructions and example datasets.'

Under 'What kind of file?', there are two radio buttons: 'Simple table (one row per sample)' (selected) and 'CDISC ADNCA dataset'.

Below this, there's a 'Browse...' button and a text input field 'Choose CSV or Excel file'. To the right, there's a green checkmark icon, the filename 'example_biq.csv', and an 'Example dataset' badge. Below this, it says '96 rows, 3 columns. Columns: Subject, Time, Concentration'.

Under 'Or use an example:', there's a dropdown menu showing 'BLQ results <= 0.5" (Tutorial 2)'. To the right of the dropdown are two buttons: 'Load example' and 'Download'.

At the bottom, there's a section titled 'Map Your Columns'.

Figure 2. Upload & Check Data after Load example: the file, its Example dataset badge and the Or use an example controls.

Leaving Out Data

Some data should not enter the analysis: a sample the laboratory reported as invalid, or a period in which the subject vomited soon after dosing. After **Process Data**, click **Add an exclusion** in the **Exclusions** card, choose the profile, and choose whether to leave out samples or the whole profile. Pick a reason from the list. Other, protocol, dosing and adverse-event reasons also need a detail, such as the deviation number. Say whether the exclusion was pre-specified in the protocol, and give the section if it was.

An excluded sample is treated as if it was never collected. C_{max}, T_{max}, AUC and the half-life of that profile are calculated without it, and the BLQ rule is applied afterwards. The dialog warns when you pick the C_{max} sample, the last measurable sample or a pre-dose sample. An excluded profile keeps its NCA result, with a 1 in the **Excluded by Analyst** column, but is left out of summary statistics, mean curves and the bioequivalence comparison.

Nothing is deleted. The register lists each exclusion with its reason, who made it and when, and whether bioequivalence results had been seen. **Restore** brings the data back; on a controlled installation it asks for a reason, and both steps go into the audit trail. Every path shows a line **Left out by you** with a link to the register, and changing the exclusions clears the results. You can also start an exclusion from Half-Life Review, with **Exclude samples or this profile....** The exclusions go into the downloads and the Analysis Record, whose reproduction script applies them again. **Download exclusions** saves the register as a file, and **Load exclusions from a file** reads it back (or the analysis_settings.json of a record) for the same data, each exclusion keeping who made it and when. The register lives only in the browser session, so the browser asks before you close or reload the page while it holds changes you have not downloaded.

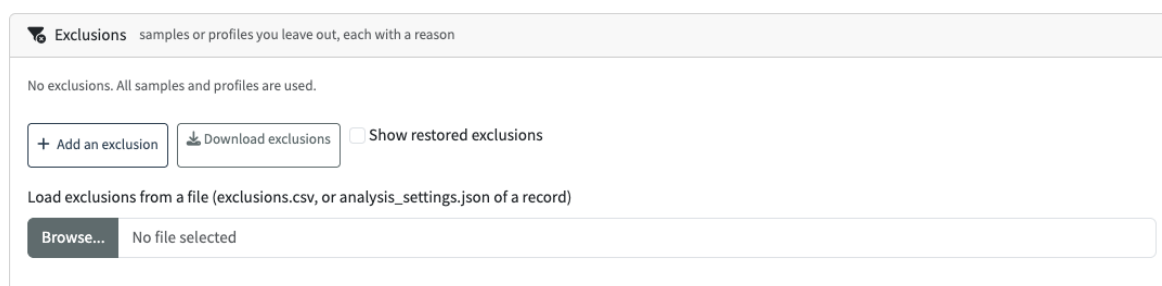


Figure 3. The Exclusions card on the Upload page, before any exclusion has been made.

Path 3: Visualize Data

Creates concentration–time figures from the processed data, from the hub or with the Visualize Data button on the Upload page. The settings are in the sidebar on the left; the plots are stacked on the right, and the **Plot Type** card decides which plots appear.

Plot Types

Plot Type has three options. Individual Profiles (spaghetti plot) draws one concentration–time line per profile (subject, treatment and period), coloured by a grouping of your choice; use it to look at individual variability, outliers and carry-over. **Summary Plot** (geometric mean $\times/\div$ geometric SD) shows the geometric mean at each time point, with error bars from the geometric mean divided and multiplied by the geometric SD. Treatment groups are overlaid automatically. Both plots shows the spaghetti plot above the summary plot.

Concentrations at a time point are usually skewed to the right, so the geometric mean describes the typical profile better than the arithmetic mean. The **Summary statistic** dropdown also offers the arithmetic mean $\pm$ SD, labelled as arithmetic and with a warning. Values of zero, for example BLQ results set to 0, cannot enter a geometric mean and are left out at that time point. Where more than half of the samples at a time point are zero or BLQ, no geometric mean is drawn there, because the remaining values would overstate the typical concentration. The figure legend counts the values left out. The mean profile groups

samples by exact time. With actual sampling times that differ between subjects it then shows single values, and the app says so; use nominal times for this figure.

Sidebar Controls

Y-axis scale (Linear or Log₁₀ (semi-log)): the semi-log view makes the terminal elimination phase linear, which shows biphasic profiles and half-life problems. Zero concentrations cannot be shown on a log scale; they are omitted with a note.

Colour lines by (spaghetti plot): Subject ID (always available), Treatment, Period and Sequence (when that column was mapped), and None. In a crossover, colour by Period to look for period effects, or by Treatment to compare formulations. **Show data points on lines** adds the samples to the spaghetti plot; **Overlay individual observations** adds them to the summary plot. **X-axis tick spacing** sets the distance between time labels (Auto by default).

Normalise to dose (show C/Dose), a checkbox that appears only when a Dose column was mapped: concentrations are divided by the dose, so the Y axis shows C/Dose (for example ng/mL per mg). A missing or zero dose gives a missing value. Change the Y-axis label to match, for example “C/Dose (ng·mL⁻¹·mg⁻¹)”.

Shade the partial AUC intervals of the last analysis shades on the summary plot the intervals of the last batch or bioequivalence run; the Figure Record reproduces the shading.

Colour palette: Default (app palette), Colorblind-safe (Okabe-Ito), or Grayscale. For print, choose Grayscale and export at 600 DPI. **Theme**: Clean (white background), recommended for reports, Minimal, or Classic. The X-axis label, Y-axis label, Plot title and Base font size (pt) (8–18) can be edited. The default Y-axis label is “Concentration”; add the drug and unit, for example “Cyclosporin (ng/mL)”. The default X-axis label is “Time (h)”.

Suggested Figure Legend

Above the plots, a card titled Suggested figure legend holds a ready-to-copy description of the current figure: number of subjects and observations, plot type, summary statistic, error bars, treatment names, the number of concentrations ≤ 0 left out of the geometric mean, **Y-axis scale** and colour grouping. It updates when you change settings. Paste it into your report and add the drug name and units.

Exporting a Figure

Export controls are at the bottom of the sidebar under **Export Figure**, always visible. A small “Will export” label above the button shows which plot will be downloaded; with Both plots, Which plot to export? chooses one. Format: PNG (Word documents, presentations, email), PDF (vector; recommended for journals that accept PDF figures), or SVG (to edit in Inkscape or Adobe Illustrator). Width (in) and Height (in) take 2–20 inches; journals usually ask for 3.0–3.5 in for a single column and 6.5–7.0 in for a double column. Resolution (DPI): 150 (screen and drafts), 300 (print, default) or 600 (high-resolution submission). **Download Figure** is grey until the first plot has appeared. The file name includes the plot type (spaghetti_plot.png or summary_plot.png), so both can be downloaded without overwriting each other.

The plots on screen are interactive (plotly): hover to see subject, time and concentration; drag to zoom; reset with the home icon. The exported file is a separate static figure (ggplot2) built from your sidebar settings.

Path 4: One Subject at a Time

Under **Where is your data?**, choose **Use my uploaded dataset (from step 2)** to step through the profiles with the arrow buttons, or **Enter data manually (type values directly)** to type times and concentrations, one per line, and click **Use This Data** (at least 3 points; Tutorial 1). Under **Analysis Settings**, set the Dose (filled in from the Dose column when there is one), the units, the **Route of administration** and the minimum R^2 , and click **Run PK Analysis**. The **PK Parameters** card shows the results; the **Half-Life Calculation** card shows the terminal-phase points, which you can change under **Points for half-life:** and **Recalculate**. The result belongs to the profile and settings it was computed with: selecting another profile, or changing a setting, clears it. Downloads: **Download Results (CSV)** and a Complete Analysis Record.

On a controlled installation, data you type in are marked in the audit trail and on the Records page, and the reviewer checks every value against the source document.

Partial AUCs (optional): set **Number of intervals** and enter the start and end of each, in the time unit of your data. Tick **End at the last measurable concentration (t)** for an interval that runs to the end of the profile, and **Also report Cmax and Tmax in this interval** when the protocol asks for the peak inside a window. An interval that reaches past the last measurable concentration stays empty, with a note saying so: partial AUCs are never extrapolated.

Path 5: All Subjects (Batch)

The sidebar has three cards. **Administration: Route of administration**, **Dose information (Same dose for all subjects, or Doses differ by subject or period (from Dose column in data))**, the dose, the infusion duration for an infusion, and **Steady-state (drug given repeatedly)** with the Dosing interval τ . **Units:** Dose unit, Time unit, Conc unit, and the molecular weight for molar units. **Calculation Method: Trapezoidal method (Linear-up / Log-down, the default, or Linear-up / Linear-down)**, **Minimum R^2 for half-life estimation** (default 0.70), the half-life rules below it (**Edit half-life rules...**), and **Calculate dose-normalised parameters** (at steady state also AUC_{τ} , C_{avg} , C_{min} and C_{τ}). Then **Partial AUCs (optional)** and **Run Analysis on All Subjects**.

The results are in five tabs. Concentration-Time Profiles shows all profiles and a Mean $\pm$ SD plot; that mean is arithmetic, so use Visualize Data for geometric summaries. **All Parameters (per profile)** is the table, with a short default view and **Show all parameters** for every column; AUC % Extrapolated above 20% is flagged in amber, and **Half-Life Flags** names any half-life rule a fit breaks. **Summary Statistics** gives the statistics per treatment. Individual Profiles (grid view) shows one small plot per profile, and **Half-Life Review** lets you choose the terminal-phase points of any profile: select it under Profile:, tick the points under **Points for half-life:** and click **Recalculate**. A selection stays with its profile; **Remove override (automatic fit)** returns it to the automatic fit. A CDISC parameter codes panel lists the official code of each parameter.

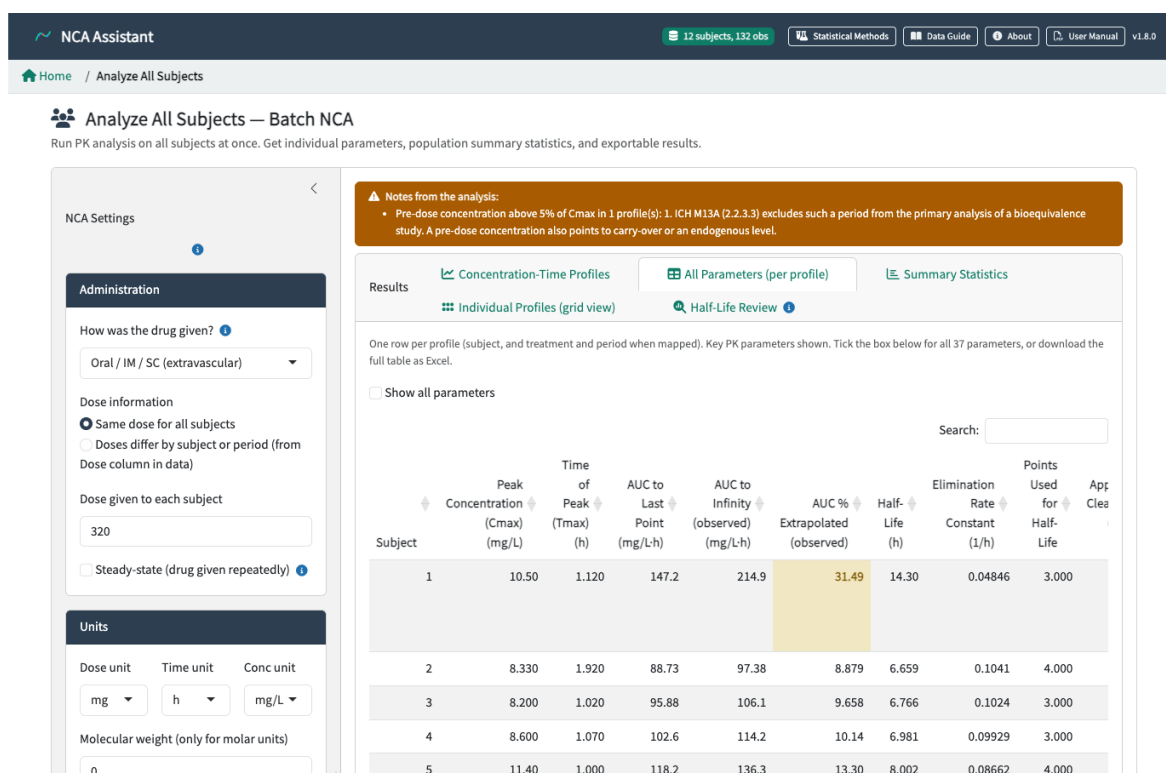


Figure 4. All Subjects (Batch) with the theophylline example of Tutorial 3: the settings, a note from the analysis and the parameter table, where the amber cell marks an AUC more than 20% extrapolated.

When the automatic terminal fit of a profile has an adjusted R^2 below the minimum, its half-life, AUC^∞ , CL/F , V_z/F and MRT are left empty and a note names the profile; select the points in **Half-Life Review** if you want a value. Profiles with fewer than 2 positive concentrations are left out, with an amber note listing them. A single-dose run also notes profiles whose pre-dose concentration is above 5% of C_{max} . At steady state, parameters extrapolated to infinity and the single-dose volume and MRT are left empty. Changing a setting clears the results, so run the analysis again.

Half-life rules. The line under the minimum R^2 shows the rules in force; **Edit half-life rules...** changes them or switches them off. By default a fit is flagged when its points span less than two half-lives, when AUC^∞ is more than 20% extrapolated (at steady state: more than 20% of AUC_T extrapolated past the last sample) or, after an IV bolus, more than 20% back-extrapolated. A flag changes no value and leaves nothing out. It marks the fits worth a second look: Half-Life Review adds (flagged) to those profiles in its list and shows each rule as pass or fail. All paths use the same rules, and the Analysis Record stores them. A profile without flags shows none, no half-life or rules off, so the column is never empty.

Partial AUCs join the table, the summary statistics and the downloads. An interval that reaches past the last measurable concentration of a profile stays empty for that profile, and a note names the profile and its last measurable time. With dose normalisation switched on, the partial AUCs and the C_{max} within an interval are divided by the dose as well; the time of that maximum is not.

Downloads: **Download Results (CSV)**, **Download Results (Excel)**, with summary statistics per treatment and a sheet of CDISC codes, and a Complete Analysis Record. For a study report, use the record: the plain downloads carry no app version, data hash or settings.

Path 6: Bioequivalence

Runs the NCA, then the ANOVA, then the confidence interval and a forest plot, and gives a verdict where one is justified.

Step 1: PK Analysis Settings

Dose information, dose, units, **Route of administration** (oral, IM or SC, or IV bolus; an IV infusion is not supported here), trapezoidal method and minimum R^2 , as in Path 5, and **Steady-state (drug given repeatedly)** with τ for steady-state data. At steady state AUC Within Dosing Interval replaces AUC to last point and AUC to infinity; after a first run, Minimum Concentration (Cmin) and Concentration at Tau (Ctau) can be added to the parameters to compare.

Step 2: Bioequivalence Settings

Study design: 2×2 standard crossover, 2×2×3 or 2×2×4 full replicate, 2×3×3 partial replicate, parallel groups, or paired comparison. The design found at upload is pre-selected, and a crossover chosen for one-period data is flagged. If all subjects received the treatments in the same order, the data are analysed as a paired comparison.

Reference treatment: the ratio is Test ÷ Reference. The run is blocked until one is chosen, and the results table shows the comparison, for example “New / Old”.

Statistical model: leave it at **Fixed effects (all terms fixed)** unless the protocol specifies the mixed model. The model has Sequence, Subject, Period and Treatment as factors; with fixed effects this is EMA Method A, and the mixed model (random subject intercept) is EMA Method B, not the FDA model for replicate designs. Fixed-Effects vs. Mixed-Effects Models, in Statistical Methods, explains the difference.

Acceptance approach (replicate designs only): **Standard (average bioequivalence)**, **EMA ABEL (Cmax, widened limits)** or **FDA RSABE (reference-scaled)**. The choice is your declaration before the run and belongs in the protocol; the app never switches by itself. ABEL is offered for 2×2×3, 2×3×3 and 2×2×4, RSABE for 2×3×3 and 2×2×4 only, because Appendix G of the FDA guidance has no steps for 2×2×3. With ABEL or RSABE the acceptance limits stay at 80 and 125: the method sets the limits, and they still apply to the metrics it does not scale. Reference-Scaled Approaches, in Statistical Methods, gives the rules.

Covariates (optional) (parallel groups only): choose baseline characteristics from the columns of your file that hold no other role. The line under the box says how each was read, numeric or categorical with its number of groups. **Advanced options** let you treat numeric codes such as 0 and 1 as categories and use the natural log of a covariate; they clear when you untick the box. Choose the covariates before you look at the results, and use values known before the first dose. The adjusted interval is the primary result, and the unadjusted interval is shown beside it. The run stops, with a message that names the problem, for a missing value, text such as “72 kg”, two different values for one subject, a constant or duplicated covariate, more than five covariates, fewer than 12 subjects per group or fewer than 10 residual degrees of freedom.

Log-transform (recommended), **Parameters to compare** (Cmax and AUC to last point by default, the primary parameters of ICH M13A; AUC to infinity, Tmax and half-life can be added), **Confidence interval** (80 to 99%; a verdict is given only at 90%) and the acceptance

limits **Lower (%)** and **Upper (%)**. Click **Run Complete BE Analysis**. Changing a setting clears the results; the parameters you chose are kept.

Widened and Narrowed Limits

With limits wider than 80–125%, two more settings appear. **Widened limits apply to** chooses Cmax only, Cmax and partial AUCs (not AUC to last point or to infinity), or All compared metrics (e.g. no-effect boundaries); the default is Cmax only. Also require the point estimate within 80.00–125.00% is ticked; untick it only for comparisons without such a rule, such as drug-interaction no-effect boundaries.

Widened limits are not ABEL or RSABE. Entering limits wider than 80.00–125.00% does not perform a reference-scaled analysis: the app does not check the conditions for widening (a pre-specified, clinically justified widening, CVwR of this study above 30%, the cap at 50%, a CVwR not driven by outliers; EMA 2010, section 4.1.10), and it does not compute the FDA scaled criterion. Use widened limits only when the protocol fixes them in advance, and document the justification outside the app. For a replicate design, choose EMA ABEL or FDA RSABE under Acceptance approach instead: the app then derives the limits, or applies the scaled criterion, from your data. A table also shows the within-subject variability of the Reference (and of the Test when replicated) and the EMA limits it would imply, for information.

For a narrow therapeutic index drug under the EMA approach, enter 90.00 and 111.11. Narrowed limits apply to every compared metric; the FDA NTID scaled test and its variability comparison are not available.

Running and Reading the Results

The tabs are **Test vs. Reference Comparison** (the table and the forest plot), PK Parameter Table, Statistical Details (ANOVA), Concentration-Time Profiles, Individual Profiles and **Half-Life Review**. Downloads: **Download Complete BE Report (Excel)**, **Download CI Table (CSV)** and a Complete Analysis Record. The verdict uses confidence limits rounded to two decimals. A design summary and an amber warning for subjects missing a treatment appear after the NCA. Reading Your Results explains the columns. If the selected design does not fit the data, the run stops with a message and shows no results.

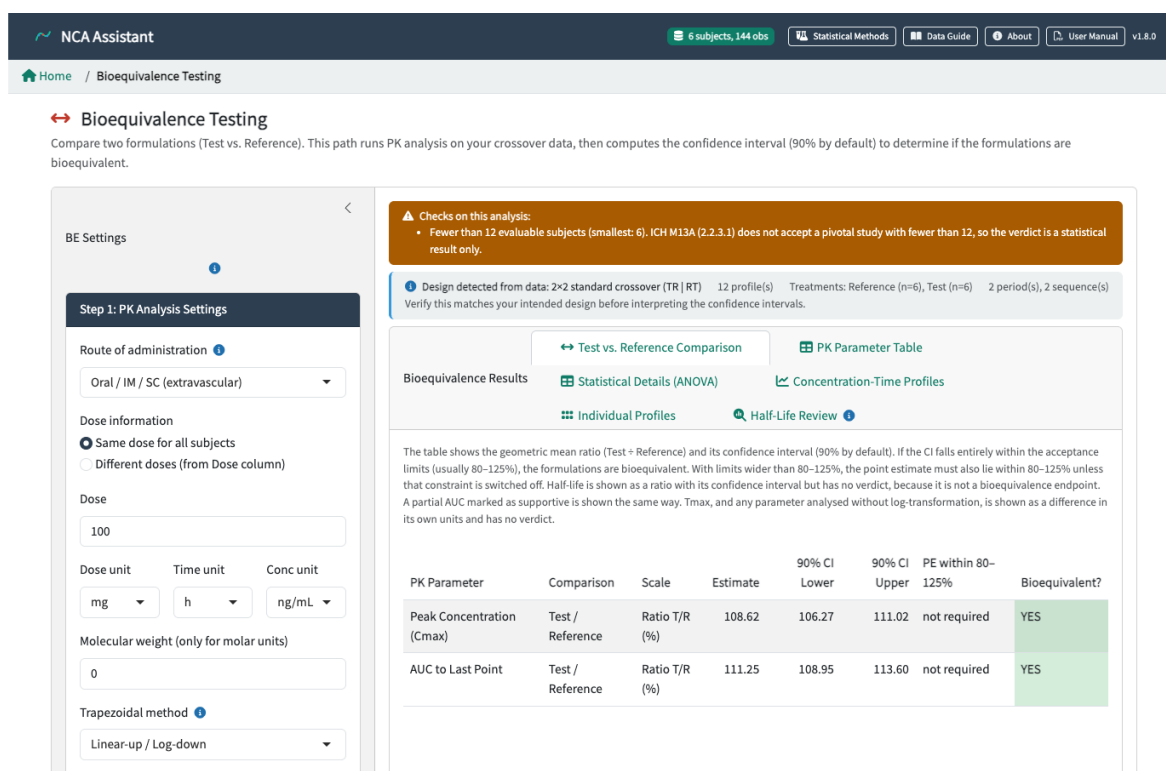


Figure 5. Bioequivalence with the crossover example of Tutorial 4: the checks, the design found in the data and the comparison table.

With the results the app lists points that ICH M13A asks you to check. A profile without any measurable concentration has no NCA result and is counted under Profiles missing; its subject leaves the comparison, which M13A (2.2.1.1) accepts only as a planned exception, in general for no more than one subject. A profile with a single measurable concentration keeps its Cmax, Tmax and AUC to last point and stays in the comparison. A period whose AUC to last point is below 5% of the geometric mean of that treatment is flagged as very low exposure, on the same terms. A single-dose study also lists a pre-dose concentration above 5% of Cmax, and AUC to last point covering less than 80% of AUC ∞ in more than 20% of the profiles; neither check runs at steady state. Fewer than 12 evaluable subjects is checked on the smaller of the Test and Reference counts.

A verdict is computed on all the data you uploaded. When the checks name a period with a pre-dose concentration above 5% of Cmax, ICH M13A (2.2.3.3) excludes it from the primary analysis, so the verdict shown is not the M13A primary analysis: exclude the period in the Exclusions card on the Upload page, with the protocol's reason (see Leaving Out Data), and run again. With fewer than 12 evaluable subjects, M13A (2.2.3.1) does not accept the study as pivotal, and the verdict is a statistical result only.

For parallel groups the results table also shows the Welch interval, which does not assume equal variances, next to the pooled interval. The pooled interval stays the primary result and decides the verdict; if your statistical analysis plan or a guidance asks for an unequal-variance analysis, the app's verdict is not your primary result.

With covariates, the line **Adjusted for** names them, the two **Unadjusted CI** columns give the interval without adjustment, and the card **Group balance on the covariates** shows mean (SD) or n (%) per group with the standardized difference. The card describes the

groups; it is not a test, and covariates are not to be chosen after looking at it. The Welch interval and the covariate coefficients are in the downloads (sheet BE_Covariates).

With ABEL or RSABE the table adds, after the verdict, the columns **Route** (Scaled or Standard), **s_WR**, **Limits** and, for RSABE, **Criterion bound**. Lines under the table say in words how each metric was judged, and the forest plot draws one pair of limits per metric. Below the switch (s_WR under 0.294) the route is Standard, which is not an error: nothing is scaled. The app warns below 24 evaluable subjects for RSABE, and says how many subjects were left out of the contrasts because they lack a period. The sensitivity analysis without exclusions uses the same approach. The downloads add a sheet BE_Scaled with the steps of the assessment.

With exclusions. The line **Left out by you** says what you left out, and the table counts excluded profiles per treatment. Under the table, **Sensitivity analysis: without your exclusions** repeats the comparison on all data, as ICH M13A expects. The M13A checks always use the data before exclusions, so leaving out a pre-dose sample does not hide its finding; a note says when the exclusions would change their outcome. The Excel report and the Analysis Record hold both results.

When There Is No Verdict

The Bioequivalent? column says why.

Situation	What the app shows
A crossover without a mapped Period column	No estimate and no verdict; the reason says to map Period
A confidence level other than 90%	The interval, without a verdict
The mixed model was chosen and cannot be fitted	The fixed-effects result for information, no verdict: changing the model after seeing the data is a post-hoc change (ICH M13A)
A paired comparison (a single treatment order)	Ratio and interval, no verdict: period and treatment cannot be separated
Half-life, Tmax, or a parameter without log transformation	A ratio (half-life) or a difference, no verdict: not bioequivalence endpoints
A supportive partial AUC	Ratio and interval, no verdict
A metric with a value of zero in any profile	No estimate and no verdict, with the zero counts per treatment

Tmax is compared as an untransformed difference with a t-based interval. Tmax is discrete and is usually reported as median and range, with a nonparametric test when the plan asks for one; do that outside the app, and ignore geometric statistics for Tmax.

Partial AUCs in Bioequivalence

Partial AUC intervals here also carry a role, set under **Role in bioequivalence**. A **Pivotal (verdict)** interval receives a bioequivalence verdict, a **Supportive (no verdict)** interval only its ratio and confidence interval. The metrics appear under **Parameters to compare**, ticked, as soon as the intervals are valid, and Cmax within an interval can be compared as well. A metric whose value is zero in any profile receives no estimate and no verdict, because zero cannot be log-transformed and the profiles that would drop out are the low-exposure ones.

That metric keeps its row in the table and in the forest plot, with the reason, the zero counts per treatment and the number of profiles that have no value.

Analysing Drug–Drug Interaction Studies

Bioequivalence (Path 6) can also analyse drug–drug interaction (DDI) studies: the design compares two conditions within subjects, with an ANOVA model and a confidence interval for the geometric mean ratio.

How to Set Up a DDI Analysis

Prepare the data as a crossover dataset (Scenario 4 or 5). Choose the victim drug given alone as the **Reference treatment**, so the ratio is (with perpetrator) ÷ (alone): the fold-change in exposure. For a single-dose victim, tick AUC to infinity under **Parameters to compare** as well; at steady state compare AUC Within Dosing Interval, Cmax and Cmin.

Run the analysis. A ratio of 200% means the perpetrator doubled the victim's exposure; 50% means it halved it. Check the Comparison column to confirm which treatment is in the denominator.

Interpreting DDI Results

A DDI study quantifies how large the interaction is; it is not a test of equivalence. ICH M12 (section 5.2) asks for the geometric mean ratio with and without the precipitant and its 90% confidence interval for AUC and Cmax (and Cmin when relevant), and for the range of individual ratios. The app does not give individual ratios: compute them from the exported PK Parameter Table. Interpret the result against no-effect boundaries for the victim drug, preferably based on its exposure-response relationship (section 5.3.1). The default no-effect boundaries are a 90% CI within 80–125%, which M12 calls generally acceptable but overly conservative for most drugs.

If the study used a sensitive index substrate, the point estimate also classifies the precipitant (M12 section 5.3.2): a strong, moderate or weak inhibitor raises the substrate's AUC at least 5-fold, 2- to 5-fold, or 1.25- to 2-fold; a strong, moderate or weak inducer lowers it by at least 80%, by 50 to 80%, or by 20 to 50%. These classes describe the precipitant, not the size of the interaction for other victim drugs.

Practical Considerations

To compare with no-effect boundaries, enter them as acceptance limits, choose All compared metrics under **Widened limits apply to**, and untick the point-estimate requirement if it does not apply. A YES in the Bioequivalent? column then means that the 90% CI lies within the boundaries you entered; it is not a bioequivalence conclusion.

For fixed-order DDI studies (all subjects receive the victim alone first, then with the perpetrator), select Paired comparison. The app reports the ratio and its confidence interval without a verdict, because period and interaction effects cannot be separated: read the interval against the boundaries yourself.

The forest plot and confidence interval table can go into the DDI section of a clinical study report when they were produced on a qualified installation and reviewed according to your procedures.

Complete Analysis Record

Available from One Subject at a Time, All Subjects and Bioequivalence, with **Generate Analysis Record**; Visualize Data produces a similar Figure Record. Fill in **Analyst name (optional)** and **Study name (optional)** first: they appear in the record, but they are typed, not signed in, so they do not establish who did the analysis. On a controlled installation, the analyst is your name from the sign-in, and the record is also stored on the server for review (Working in a Regulated Environment).

The zip contains results.xlsx, app_results_reference.csv, analysis_settings.json, nca_pipeline.R (the app's own data-processing code), reproduce_analysis.R, reproduction_check.txt, data_integrity.txt, analysis_summary.html and a copy of the data file. From the batch and bioequivalence paths, results.xlsx has a Checks sheet, and the summary a Checks and Notes section, with the data quality findings and the notes shown with the results, such as the ICH M13A checks. A bioequivalence record also holds app_be_reference.csv (the app's table with its unrounded estimates, and app_be_sensitivity_reference.csv when exclusions were made) and the app's bioequivalence code: be_analysis.R, be_scaled.R, designs.R and utils.R. data_integrity.txt holds SHA-256 hashes of the data, the settings, the results, the reference results, the script and the pipeline code, and for a bioequivalence record these files too. A SHA-256 hash is a 64-character fingerprint that changes if a single character in the file changes.

When it creates the record, the app runs the reproduction script in a separate R process and writes the outcome to reproduction_check.txt:

Outcome	Meaning
MATCH	The shipped pipeline, run again on the enclosed data and settings, gives the same NCA parameters (largest relative difference below 10^{-6}); for a bioequivalence record, also the same point estimates, confidence limits and verdict
CLOSE	The same within a relative difference of 10^{-3} , for example after a package update that changes rounding
DIFFERENT	Different results, or a data file, pipeline code, bioequivalence code or reference results that were changed, a changed setting or Reference treatment, or a parameter present on one side only
NOT COMPARED	The reference results (app_results_reference.csv) are missing

What the record shows and what it does not. MATCH shows that the analysis can be reproduced with the same code and packages. It is not independent verification. For a bioequivalence record the script recomputes the NCA parameters and then the point estimates, confidence limits and verdict, with the same code the app used. It does not compare the ANOVA tables or results.xlsx. The manifest is not signed: whoever changes a file can also recompute its hash. To make the record tamper-evident, store the zip, or its SHA-256, in a controlled system. The record is not an audit trail in the sense of 21 CFR Part 11 or EU GMP Annex 11; it documents one analysis. Timestamps come from the clock of

the computer that ran the app (UTC on shinyapps.io), and the summary names the time zone. On a controlled installation, the stored copy on the server is covered by the audit trail, and the reviewer's signature is bound to the record's SHA-256 (The Signature Sheet, in the chapter Working in a Regulated Environment).

If the uploaded file was no longer available when the record was made, the record holds the table as the app read it, written as CSV, with a different name and hash; the settings and the Checks say so. Keep the original source file with the record.

Partial AUC intervals, their roles and the BLQ share used for the notes are recorded, and the reproduction script recalculates the partial AUCs. Change an interval in the settings file afterwards and the check reports DIFFERENT.

Re-running a record. Unzip it into an empty folder. In R, set the working directory to that folder (`setwd("path/to/folder")`) and run `source("reproduce_analysis.R")`. The script checks the data hash, recomputes the results (for a bioequivalence record, also the confidence limits and the verdict) and prints the outcome. It also says whether R and the packages are the versions the record was made with. Under other versions, a result other than MATCH can come from the versions alone.

Reading Your Results

The Data Quality Report

After Process Data, the Upload page lists what the checks found, one line per finding.

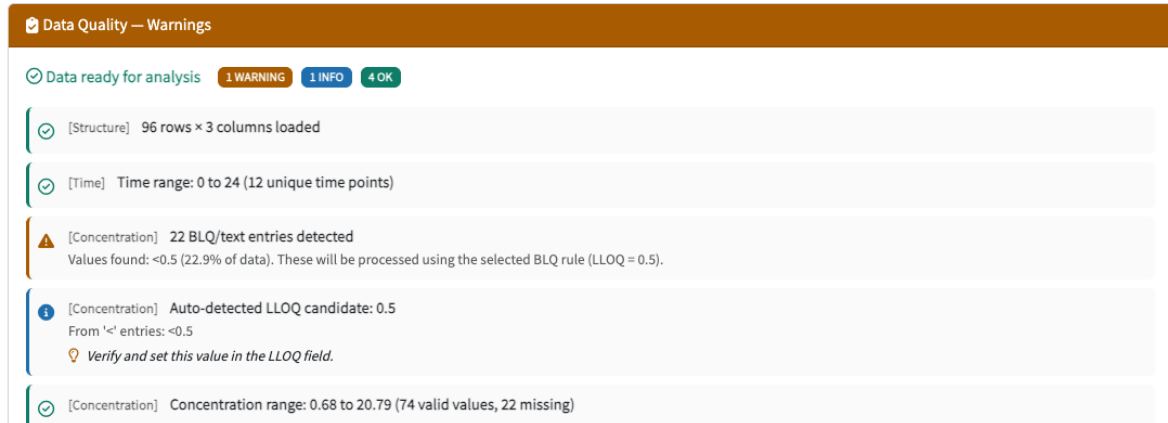


Figure 6. The data quality report for example_biq.csv with an LLOQ of 0.5 (Tutorial 2).

Column	Meaning
Severity	ERROR blocks processing; WARNING and INFO do not; OK means the check passed
Category	What was checked: Structure, Time, Concentration, Subjects, Design, Format and others
Message	What the check found, for example "22 BLQ/text entries detected"
Detail	The rows, subjects or values concerned
Action	What to do about it, where there is something to do

Fix every ERROR in the file or the mapping, and process the data again. Read the warnings too. They do not stop the analysis, but they can change what the results mean.

The Parameter Table

All Parameters (per profile) in Path 5, and PK Parameter Table in Path 6, have one row per profile: a subject under one treatment in one period. Subject, Treatment and Period identify the row; the other columns are parameters, labelled in words with their unit, for example "AUC to Last Point (ng/mL·h)". The default view shows the main parameters; Show all parameters shows every column. Appendix A lists each parameter with its definition, unit and when it is reported. An empty cell is not an error: it means the parameter could not be estimated for that profile, and a note above the table says why (for example a terminal fit below the minimum adjusted R^2). **Half-Life Flags** names any half-life rule the fit breaks (or says none, no half-life or rules off), and **Excluded by Analyst** is 1 for a profile you left out.

Summary Statistics

Summary Statistics gives, per treatment and parameter:

Column	Meaning
N	Number of profiles with a value
Geometric Mean	$\exp(\text{mean of } \ln \text{ values})$; the usual summary for AUC and Cmax
Geometric CV (%)	$100 \cdot \sqrt{\exp(\text{SD}^2 \text{ of } \ln \text{ values}) - 1}$
Median, Q1 (25%), Q3 (75%)	The middle value and the quartiles; the usual summary for Tmax
Arithmetic Mean, Std Dev, CV (%)	Mean, standard deviation, and SD as a percentage of the mean
Min, Max	The range

Geometric statistics need positive values; they are not meaningful for Tmax or for a parameter that can be zero.

The Bioequivalence Table and Forest Plot

The table on the Test vs. Reference Comparison tab has one row per compared metric.

Column	Meaning
PK Parameter, Comparison	The metric, and which treatment is divided by which ("Test / Reference")
Subjects (Test), Subjects (Reference)	Subjects that enter the comparison
Estimate	The geometric mean ratio Test ÷ Reference in percent (a difference for Tmax or untransformed parameters)
90% CI Lower, 90% CI Upper	The confidence interval, at the level you chose
Accept. Lower, Accept. Upper	The acceptance limits used for this metric
PE within 80–125%	Whether the point estimate lies within 80.00–125.00%, for widened limits
Bioequivalent?	YES or NO, or the reason why there is no verdict
Profiles missing, Zero values, Mostly BLQ (Test and Reference)	Profiles without a value, with a value of zero, or resting mainly on values set by the BLQ rule
Model	The model fitted, and whether a mixed model failed
Adjusted for, Unadjusted CI Lower and Upper	With covariates: the characteristics the estimate is adjusted for, and the interval without adjustment
Route, s_WR, Limits, Criterion bound	With ABEL or RSABE: Scaled or Standard; the within-subject SD of the Reference (FDA contrasts for RSABE, EMA model for ABEL); the limits in force; and, for RSABE, the criterion, met at 0 or below
Profiles excluded, Half-life flags (Test and Reference)	Profiles you left out, and profiles whose half-life fit breaks a rule; shown when there are any

Subjects without both treatments	In a 2×2 crossover or paired comparison: subjects left out of both models because they lack Test or Reference; shown when there are any
----------------------------------	---

The forest plot shows each estimate as a point with its confidence interval as a horizontal bar, and the acceptance limits as vertical lines: a bar entirely between the lines is a YES. The Residual Variance (MSE) and Degrees of Freedom are in the downloads (Download CI Table (CSV)); the within-subject CV of a crossover is $100 \cdot \sqrt{(\exp(\text{MSE}) - 1)}$.

Statistical Details (ANOVA) shows, per metric, the model terms (Sequence, Subject, Period, Treatment) with their degrees of freedom and F tests. A significant Sequence or Period term is a reason to look at the data, not a reason to change the analysis.

Judging the Terminal Phase

Half-life, AUC to infinity, CL/F, Vz/F and MRT all rest on the terminal slope λ_z . Before you report them, check each profile against these points, and follow your statistical analysis plan where it is stricter.

- **At least 3 points** in the terminal log-linear phase (ICH M13A, 2.1.8), after T_{max} and not counting C_{max} itself. The app accepts 2 points in a manual selection, but a 2-point fit has no measure of fit and should not be reported.
- **Span**: as a rule of thumb the points should cover at least two half-lives; ICH M13A (2.1.8) asks for sampling over at least three. The **Half-Life Flags** column marks a fit that spans less than two half-lives.
- **Adjusted R²**: the app's default minimum of 0.70 is a floor, not a recommendation; many analysis plans use 0.80 or 0.90.
- **AUC % Extrapolated** above 20% (amber in the table): the extrapolated part carries much of AUC to infinity. Report AUC to infinity, CL/F, Vz/F and half-life with a flag, or leave them out of the descriptive statistics, as the plan says. The **Half-Life Flags** column marks it too.

The app has no switch that leaves flagged values out of Summary Statistics: it uses them, and counts them in the column N Flagged (half-life rules). When your plan asks for statistics without them, download the results (Excel or CSV), remove the affected values where Flag: AUC Extrapolated Above Maximum or Flag: Span Below Minimum is 1, and compute the statistics from what is left. Say in the report which values were left out, and why.

Half-Life Review shows the fit the results use, on a log scale, with the chosen points marked.

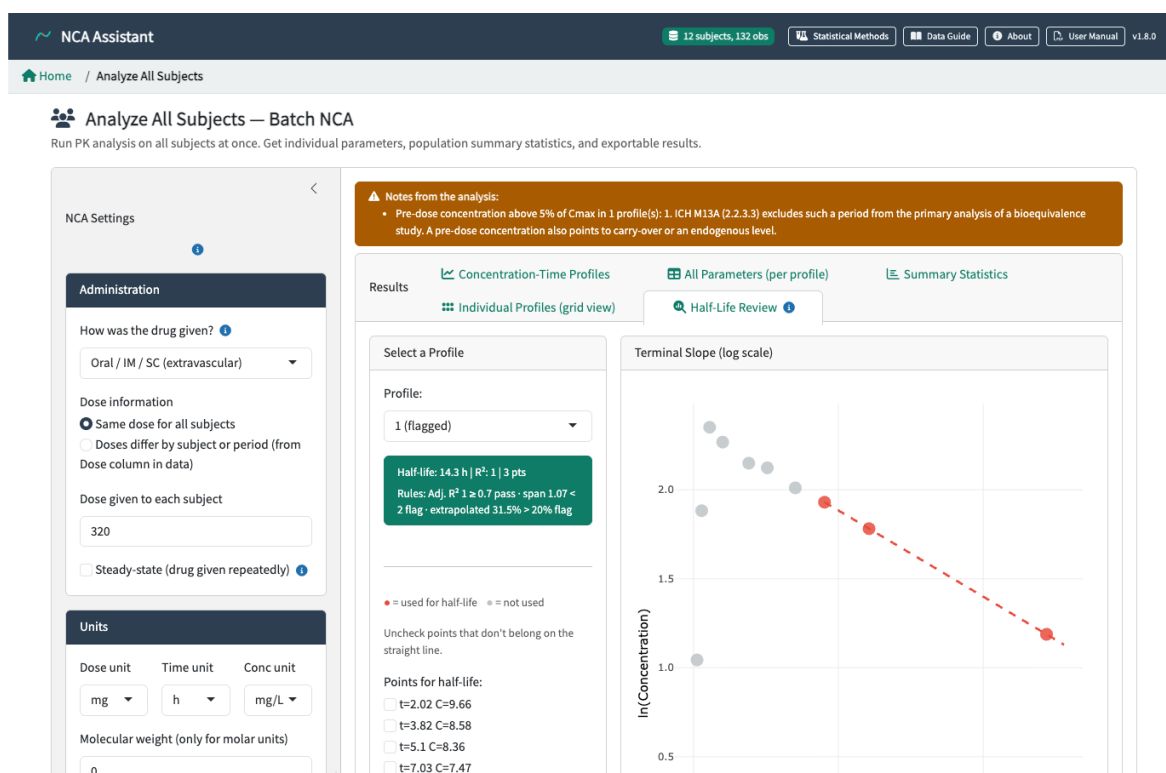


Figure 7. Half-Life Review for theophylline subject 1: the fitted points in red and the rules, with flags for a span of 1.07 half-lives and 31.5% extrapolated.

Checking Plausibility

- **Pre-dose concentration above zero:** carry-over, earlier dosing or an endogenous compound. In a bioequivalence study a value above 5% of C_{max} takes the period out of the primary analysis (the app says so).
- **C_{max} at the first sample after the dose:** absorption was faster than the sampling, and C_{max} may be underestimated. ICH M13A (2.1.8.1) asks for this to be discussed.
- **T_{max} at the last sample:** absorption was still going on, or sampling stopped too early; there is then no terminal phase and no λ_z.
- **A half-life much longer than expected:** check the number of points and the % extrapolated (Tutorial 3).
- **Flip-flop kinetics:** with slow absorption, such as a modified-release product, the terminal slope reflects absorption rather than elimination. Point selection cannot correct this: report the value as an apparent half-life. Only IV data can tell the two apart.
- **A profile without measurable concentrations (non-absorber):** it has no NCA result, and in bioequivalence it is counted as missing and named in the checks.
- **Very different results under two BLQ rules:** the result depends on the imputation; report the rule and consider a sensitivity analysis.

Working in a Regulated Environment

This chapter is for everyone who uses NCA Assistant for work that has to stand up to an audit or an inspection: analysts, reviewers, inspectors and the system owner. It explains what a controlled installation does, how its audit trail and its accounts work, and how a study is carried out on it, from loading the data to an approved and archived record. It also says what the app leaves to your organisation. Setting up the server is in the next chapter, Installing NCA Assistant for Regulated Use.

After this chapter you should be able to:

- say which work needs a controlled installation, and which rules apply;
- explain to an inspector what the audit trail records, and how a change to it would be detected;
- carry out a study analysis so that every step can be traced, from the source file to the signed record;
- review and sign a record, and review the audit trail;
- tell which controls the app provides and which your organisation must provide.

New to controlled mode? Read the sections in order, then do Tutorial 8 at the end, on a training installation.

Before You Rely on the App

Status and Scope

NCA Assistant is open-source software without warranty or supplier support, and it is not validated for your use. Its bioequivalence analysis covers average bioequivalence, EMA ABEL and FDA RSABE, not the FDA scaled test for narrow therapeutic index drugs. The public instance on shinyapps.io is outside the scope of any qualification: its package versions are set by the deployment, not by the validated release. For records under 21 CFR Part 11 or EU GMP Annex 11, use a controlled installation; a qualification of controlled mode holds only for the server on which it was carried out, not in general.

Which installation for which work. Electronic records and signatures that fall under 21 CFR Part 11 or EU GMP Annex 11 need a controlled installation. A qualified installation without controlled mode may be used for regulated work only where your procedures allow hybrid records: the results are filed and signed outside the app, with the procedural controls listed in What the App Does and What Your Organisation Does. The public instance on shinyapps.io is never used for regulated work.

Which Rules Apply

Bioequivalence and drug interaction studies are clinical trials, so good clinical practice applies: ICH E6(R3) and the EMA guideline on computerised systems and electronic data in clinical trials (EMA/INS/GCP/112288/2023). EU GMP Annex 11 and 21 CFR Part 11 apply where your organisation or the use falls under them. Your own procedure for validating computerised systems decides what is needed; the validation package gives the evidence to start from.

Qualifying an installation, from the approval of the protocol to the release by QA, is Step 8 of Installing NCA Assistant for Regulated Use.

How a Controlled Installation Works

What Is Different on a Controlled Installation

The analyses are the same. What changes is around them:

	Public app and standard installation	Controlled installation
Signing in	Not needed	Personal user ID and password; a new account starts with the starting password admin
Top bar	The usual buttons	Also a shield icon (hover over it to see the organisation and server), Records, Audit trail for reviewers and inspectors, and a menu with your name
Analyst name in a record	Typed by you, optional	Your name from the sign-in, not editable
After you create a record	You download it	You download it, and a copy is stored on the server for review
Data typed in by hand	Allowed	Allowed, and marked for the reviewer
Audit trail	None	Every sign-in, data load, analysis, download, record and signature, with time and user

What is recorded about you. The audit trail holds your user ID, name and role, and when you signed in, which data you loaded, which analyses you ran with which settings, what you downloaded and what you signed. Reviewers, inspectors and the system owner can read it. It is kept as long as the study records; your organisation's privacy notice says why.

Roles

Every user has one or more roles. The app checks the role on the server for every action, so a hidden button is never the only protection.

- **Analyst:** loads data, runs analyses and creates records. Sees only their own records on the Records page, and not the Audit trail page.
- **Reviewer:** sees all records and their history, approves or rejects the records of others, and reviews and signs the audit trail. A reviewer can never approve their own record, not even with the analyst role as well.
- **Inspector:** sees what a reviewer sees, and can export and verify, but cannot sign or change anything. Meant for an auditor or an inspector, for the length of an inspection.
- **System owner:** not a role in the app. The system owner manages the accounts on the server with `gxp/manage_users.R`, and each change is logged under their own name, with a reason.

Give each person the smallest role that fits. The app stops you from signing your own record, but not from signing a colleague's record for a study you worked on. That is for your procedure to prevent.

How the Audit Trail Works

The audit trail is a list of entries that only grows. Each entry records one event, with:

- a sequence number and the time in UTC, from the server's clock;
- the user ID, role and organisation;
- the event and what it concerned (a file, a record or an account), with the SHA-256 of that file or record where there is one;
- the details, such as the settings of an analysis run or the reason for a signature or an exclusion;
- the app version and the SHA-256 of the pipeline code in use.

What is recorded. Sign-ins, failed sign-ins and lockouts; password changes and failed attempts; every data file loaded; every analysis run, with its settings; every record created, downloaded, verified or signed; every exclusion added, restored or loaded from a file; downloads of results; reviews, exports and verifications of the trail itself; account changes by the system owner; security alerts; and every start and stop of the app. Moving between pages is not recorded.

How a change would show. Each entry also holds the fingerprint (the SHA-256 hash) of the entry before it, and its own fingerprint is computed over all of that. Think of numbered pages where every page repeats the seal of the page before. Change one word on page 40, and its seal no longer fits page 41, or any page after it. On top of that, the database refuses to change or delete an entry. Verify chain on the Audit trail page, or `manage_users.R` verify on the server, recomputes every fingerprint and reports the first entry that does not fit.

What it cannot see on its own. Someone with root access to the server could replace the whole trail with a new one that is consistent in itself. This is why the system owner files the fingerprint of the last entry (the head) outside the server at fixed times: a rebuilt trail would not contain the filed heads. The signature sheets in your study files and the backups would show it as well.

Fail-closed. If an entry cannot be written, for example because the disk is full, the action it belongs to does not happen, and the app says so. An analysis that is not in the trail was not done on this installation.

How Accounts Work

An account goes through the same stages from the first day to the last, and each stage is in the audit trail.

1. **Identity and training.** Before creating an account, the system owner checks who the person is and that their training is recorded.
2. **Creation.** The system owner runs `manage_users.R add` with the user ID, the printed name, the roles and a reason. The printed name is the one on records and signature sheets.

3. **First sign-in.** A new account's password is admin, and the app makes you choose your own straight away (Signing In for the First Time). The system owner creates the account only when you are ready to sign in, and checks afterwards that you did.
 4. **Passwords.** At least 12 characters, with a digit, a lower-case and an upper-case letter. A password expires after 90 days. Five wrong passwords lock the account, and only the system owner can open it again, with a reset.
 5. **Changes of role.** The system owner changes roles with manage_users.R role. A change applies at once: a reviewer whose role is removed can no longer sign, even in a session that is still open.
 6. **Deactivation.** When someone leaves, or no longer needs access, the system owner deactivates the account the same day. Accounts are never deleted and a user ID is never used again, so every entry in the trail keeps pointing to one person.
- The Users tab of the Audit trail page shows every account with its roles, status and history.

Signing In and Staying Signed In

Signing In for the First Time

1. The system owner gives you your user ID and tells you to sign in at once. The starting password is **admin**.
2. Open the address of your installation and sign in with your user ID and admin.
3. A window asks you to choose your own password: type admin as the current password, then the new one twice. Use at least 12 characters, with a digit, a lower-case and an upper-case letter; the app refuses admin and anything shorter.
4. Click Change password. The app opens, and the top bar shows your name.

If admin does not work for your new account, someone may have used it before you. Tell the system owner at once, so that the account is checked and reset.

Changing Your Password

Open the menu with your name in the top bar and choose **Change password**. Enter your current password, then the new one twice. You stay signed in. The password also expires after 90 days, and the app then asks for a new one when you sign in.

If you think someone else knows it, change it at once and tell the system owner. If you have forgotten it, ask the system owner for a reset: you then sign in with admin and choose a new password.

After five wrong passwords at sign-in, your account is locked. Only the system owner can unlock it.

Staying Signed In

After 15 minutes without activity you are signed out. Moving the mouse, typing and scrolling all count as activity. Two minutes before the end, a message counts down; click **Stay signed in**, or move the mouse. When you are signed out, the uploaded file, the settings and the results are gone, as when you close the browser: upload the file again. Records you created are safe, because they are stored on the server when you create them.

To sign out yourself, open the menu with your name and choose **Sign out**.

A Study, Step by Step

The steps below follow one study from start to finish, in the order the work is done. Where a step leaves a trace in the audit trail, the step says so. Your procedures may add steps; they should not leave any out.

Before the Study

- The protocol or the statistical analysis plan fixes every setting that changes results: design and Reference treatment, model, acceptance limits and what they apply to, BLQ rule and LLOQ, minimum R^2 and half-life rules, partial AUC intervals, trough window, and the criteria for leaving data out.
- Everyone who will analyse or review has an account with the right role, and has been trained on this chapter.
- The installation runs the qualified release. The About page shows the version.

Before Each Analysis

On a controlled installation, check that the top bar shows the shield icon and your name. Check on the About page that the version and the pipeline SHA-256 are those of your qualified release (the release manifest, validation/release_manifest.csv, lists the hash of every file). Record the SHA-256 of the source data file. Check the settings against the protocol and the statistical analysis plan: design, Reference treatment, model, acceptance limits and their scope, BLQ rule and LLOQ, partial AUC intervals and their roles.

Loading the Data

Upload the source file as it came from the laboratory or data management, unedited. The audit trail records its SHA-256, and so does the record, so anyone can later check that the file in the study file is the one that was analysed. If the data were typed from paper, or typed in by hand in One Subject at a Time, the record says so, and the reviewer checks every value against the source.

Exclusions, Outliers and Deviations

Leave data out in the app, with **Add an exclusion** on the Upload page (see Leaving Out Data), rather than by editing the file. ICH M13A (2.2.1) asks for the criteria to be defined in the protocol and for any exclusion to be documented before the bioanalysis. Extreme values should not be removed on statistical grounds alone; only contemporaneously documented protocol violations justify removal (M13A 2.2.1.1). A subject without measurable concentrations, or with very low ones (AUC below 5% of the geometric mean), may be excluded only as an exception, in general for no more than one subject. A period with a pre-dose concentration above 5% of C_{max} is excluded from the primary analysis (2.2.3.3).

A procedure that keeps the chain from source to result intact:

1. Upload the unedited source file; the Analysis Record holds its SHA-256.
2. Enter each exclusion with its reason and, where it applies, the protocol section. Enter pre-specified exclusions before you run the bioequivalence analysis: the register marks any

exclusion made after the results were seen. Download the register before you end the session; loaded back later, it keeps its timing.

3. Run the analysis and create the Analysis Record. It holds the register, the results with and without the exclusions, and a reproduction script that applies the same exclusions.

4. Report the sensitivity analysis where the plan asks for it.

5. If you still prepare a separate analysis file, file the SHA-256 of both files and the script that made it with the record.

Each exclusion, and each restore, is an entry in the audit trail with its reason.

Running the Analysis

Every run is recorded with its settings, including the runs you later discard. That is on purpose: the reviewer sees the history of the data and will ask why a setting changed. So change a setting only for a reason you can write down, write that reason in the study file, and do not try settings until a result looks right. A run on the wrong file or with a wrong setting is not a problem in itself. Leaving it unexplained is.

Creating a Record for Review

Create an Analysis Record as usual, with **Generate Analysis Record** (Complete Analysis Record, in the chapter on using the app). The record names you as the analyst, from your sign-in. After the download, a line under the button says **Stored for review**, with the first characters of the record's SHA-256, and **Open in Records**. A Figure Record is stored in the same way but is not reviewed: its line says **Stored**. The audit trail records the creation of the record, with its SHA-256.

On the **Records** page you see your own records and whether they are awaiting review, approved or rejected. A rejected record stays as it is. If the analysis must change, run it again and create a new record.

Data typed in by hand. In One Subject at a Time you can still type in times and concentrations. The record then says so, and the reviewer checks every value against the source document. The same applies to a file that someone typed from a paper source: tell the reviewer where the file came from.

Reviewing and Signing a Record

For reviewers. Open **Records**. The list starts at **Awaiting review, not mine**; choose **All records** or **My records** to see others. Your own records have no Approve or Reject buttons: someone else must review them.

Select a record. The details show the type, the study, who created it and when, the outcome of the reproduction check, and the SHA-256. Below them, **Audit history of the data** lists every run on the same data before the record was made, with the settings that changed between runs. A label says so when there was more than one run. **Open summary** shows the record's own summary; **Download** gives the zip.

The screenshot shows the 'Records' page of the NCA Assistant application. The top navigation bar includes links for Statistical Methods, Data Guide, About, User Manual, Records, Audit trail, and a user profile for Robin Reviewer. The main content area is titled 'Records' and includes a sub-header 'Every record stored on this installation. Select a record to see its details.' Below this is a table of records with columns for Created, Study, Type, Author, and Status. A single record is shown: '2026-09-28T20:10', 'Theophylline (training)', 'All subjects', 'Anna Analyst', and 'Approved'. To the right of the table is a 'Details' panel for the selected record, showing its title, status (Approved), type, study, creation time, reproduction check (MATCH), SHA-256 hash, and review status (Approved by Robin Reviewer). Below the details is an 'Audit history of the data' section showing a single entry for the analysis run. At the bottom of the page is a 'Verify a record file' section with a file upload button and a message indicating no file is selected.

Records
Every record stored on this installation. Select a record to see its details.

All records

Created	Study	Type	Author	Status
2026-09-28T20:10	Theophylline (training)	All subjects	Anna Analyst	Approved

Showing 1 to 1 of 1 entries

Previous 1 Next

Details

Analysis_Record_Theophylline__training__2026-09-28.zip
Approved

Type All subjects

Study Theophylline (training)

Created 2026-09-28T20:10:37.027Z (UTC), Anna Analyst (anna)

Reproduction MATCH

SHA-256 abe59e0e6355d0fb ...

Review Approved by Robin Reviewer (robin)

Signature valid: the stored record matches SHA-256 abe59e0e.

Audit history of the data

Time (UTC)	By	Analysis	Settings changed
2026-09-28T20:10:33	anna	Ran an analysis (all subjects)	

Open summary Download (with signature sheet) Signature sheet

Verify a record file

Upload any record zip to see whether this audit trail has it, and its signatures.

Browse... No file selected

Figure 8. The Records page of a reviewer: an approved record with its reproduction check, valid signature and the audit history of its data.

Before you approve, check:

- the settings match the protocol and the statistical analysis plan (design, Reference treatment, model, limits, BLQ rule and LLOQ, partial AUC intervals);
- the data SHA-256 matches the source file, and, where the data were typed in (the record is marked) or the file was typed from paper, every value matches the source document;
- earlier runs on the same data are explained, especially when a setting changed between runs;
- the reproduction check says MATCH;
- the checks and notes of the analysis have been dealt with.

Then click **Approve** or **Reject**. The dialog shows which record you sign and what your signature means. For a rejection, type the reason; **Sign** stays grey until you do. Type your user ID and password and click **Sign**, or press Enter. After a wrong password the message says how many attempts are left; the third failure signs you out. A record is reviewed once, and then stays approved or rejected.

Approve this record

Record	Analysis_Record_Theophylline__training__2026-09-28.zip
Study	Theophylline (training)
Type	All subjects
Created	2026-09-28T20:10:37.027Z (UTC) by Anna Analyst
SHA-256	abe59e0e

I have checked the settings of this analysis against the protocol, its results and its audit trail, and I approve this record.

User ID

robin

Password

Cancel Sign

Figure 9. The signing dialog: the record, what the signature means, and the user ID and password.

The Signature Sheet

Every reviewed record has a signature sheet, a web page you can print or file. It shows the record and its SHA-256, who created it and when, who reviewed it, when, with which meaning and reason, and the audit entry of the signature. All times are in UTC. **Download (with signature sheet)** gives a zip with the record and the sheet together; file both in the study file.

Before it shows a signature, the app checks that the stored record still matches its SHA-256. **Signature valid** means it does. **Signature INVALID**, in red, means the stored record has changed since it was signed: tell the system owner at once. To check a copy you hold, for example from the study file, upload it under **Verify a record file** on the Records page. The app lists the signatures of that exact file, or says that no record with its SHA-256 exists.

Reading Verdicts in a Regulated Context

- A verdict is computed on all the data you uploaded. When the checks name data that ICH M13A excludes, the verdict is not the M13A primary analysis (Path 6, Running and Reading the Results).
- Widened limits you type are not ABEL or RSABE (choose the approach for those), and All compared metrics also widens AUC to last point and AUC to infinity, which no guideline allows for bioequivalence. To widen a partial AUC, choose Cmax and partial AUCs.
- ICH M13A allows a fixed or a mixed model, but its primary analysis includes only subjects with data for both treatments (2.2.3.2). In a 2×2 crossover the app leaves subjects without both treatments out of both models; in replicate designs the mixed model still uses them, so use it there only when the plan specifies it.

- For parallel groups the pooled-variance interval is the app's primary result. ICH M13A, the EMA guideline and the FDA guidance of May 2026 do not say which interval to use; if the plan specifies an unequal-variance (Welch) analysis, that is your primary result.

The Analysis Record in a Regulated Context

The record documents one analysis and shows that it can be reproduced with the same code and packages (Complete Analysis Record, in the chapter on using the app). For a bioequivalence analysis it recomputes the confidence intervals and the verdict, but not the ANOVA tables. It is not signed, and whoever changes a file can recompute its hash. Store the zip, or its SHA-256, in a controlled system. The Analyst name and Study name fields are typed, not signed in: record who ran the analysis in your log. For a manual selection of half-life points, the record keeps the final selection but not the reason: write the reason in the study file. On a controlled installation the record is stored on the server, the analyst is the signed-in user, and the review signature is kept in the audit trail and shown on the signature sheet, bound to the record's SHA-256.

Reporting and Citing

Cite the article that describes NCA Assistant: ter Heine R. NCA Assistant: An Open-source R/Shiny Interface for Non-compartmental Pharmacokinetic Analysis, Bioequivalence Testing and Study Planning. The AAPS Journal. 2026;28(6):163. doi:10.1208/s12248-026-01316-w. To cite the exact version you ran, add: ter Heine R. NCA Assistant (v1.8.0). Radboud Applied Pharmacometrics, Radboudumc, Nijmegen, The Netherlands. <https://github.com/Robterheine/NCAassistant>. Add the pipeline SHA-256, the reference of your qualification, and the SHA-256 of the Analysis Record. The Example Methods Paragraphs on the Statistical Methods page describe the calculations; they say the analysis follows ICH M13A, EMA and FDA guidance, so check that your study design and analysis do before you keep that sentence. For a sample-size justification, record the planner's inputs and result (Path 1).

Archiving and Retention

File the source data file, the analysis file, the Analysis Record, the SHA-256 of the zip, the release tag, validation/renv.lock and the qualification reference together. Keep them as long as your procedures and the law require; in the EU the trial master file of a clinical trial is kept for at least 25 years (Regulation (EU) No 536/2014, Article 58). The app keeps nothing: download the record before you close the session. On a controlled installation, the system owner archives the audit trail and the records with `manage_users.R` archive (Appendix E), which also keeps the software needed to verify and restore them.

Reviewing the Audit Trail and Inspections

Reviewing the Audit Trail

For reviewers and inspectors. Open **Audit trail**. The line at the top says when the trail was last reviewed and how many entries have been added since.

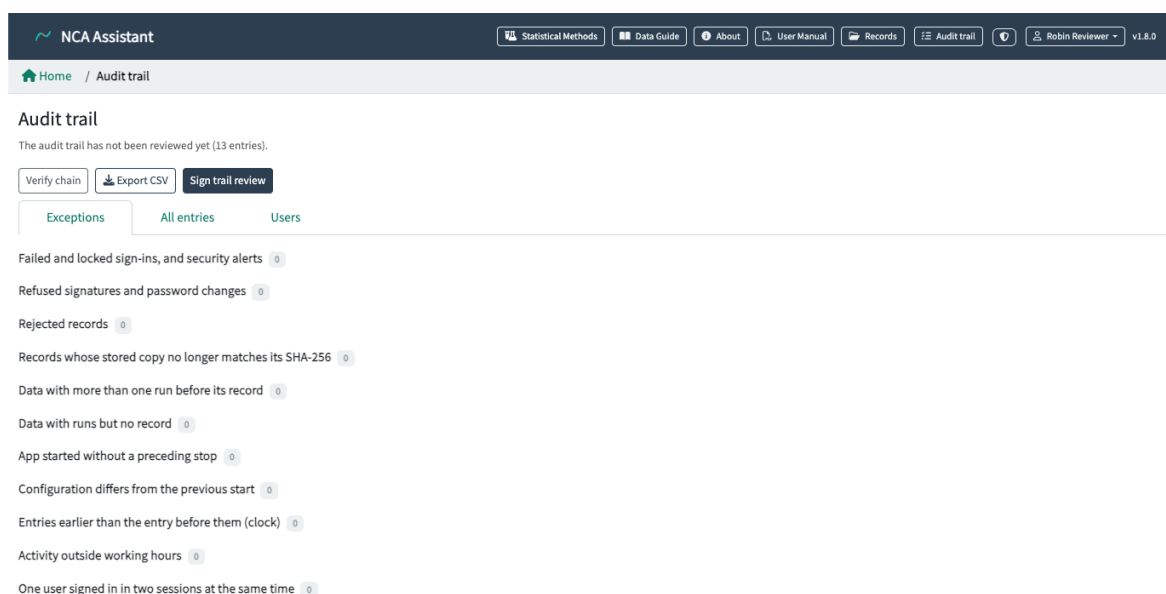


Figure 10. The *Exceptions* tab of the *Audit trail* page. Every count is zero here; a category with entries lists them below it.

- **Exceptions** lists what needs attention: failed and locked sign-ins and security alerts, refused signatures and password changes, rejected records, stored records that no longer match their SHA-256, data analysed more than once before a record, data analysed without a record, app starts without a preceding stop, changes in the configuration, entries out of time order, activity outside working hours, and one user signed in twice at the same time. Start here.

- **All entries** shows the whole trail in plain words, newest first. Filter by date, user or event, or search for a file name or SHA-256. Select an entry to see all of it, including the settings of a run.

- **Users** lists every account with its roles, status and role history.

Verify chain checks that no entry has been changed, inserted or removed. **Export CSV** saves the entries of the current filter, for example one study, with the result of the check at the top. After a review, a reviewer clicks **Sign trail review** and signs with user ID and password; this records the period reviewed. How often to review, and what else triggers a review, is set by your organisation's procedure.

For Inspectors

An inspector account can see every record, its history and signatures, the audit trail and the user overview, and can export and verify, but cannot sign or change anything. The system owner creates one per inspection and deactivates it afterwards.

Preparing for an Inspection

An inspector will usually want to follow one result back to its source. For each study, be ready to show:

1. the signed records with their signature sheets (Download, with signature sheet);
2. the CSV export of the audit trail for the study's data files, with the result of **Verify chain** at the top;

3. the source data files, whose SHA-256 matches the records;
4. the exclusion register with its reasons, and the notes in the study file that explain repeated runs;
5. the qualification summary report of the installation, and the release it covers;
6. the account list, and the most recent account review and trail review.

Create an inspector account for the inspection, and deactivate it afterwards. If the inspector asks how a change to the trail would be detected, show Verify chain and the heads filed in your quality system.

Good Practice: Data You Can Trust

Inspectors use the ALCOA+ principles to judge whether data can be trusted. The app supports each principle, but only when it is used as intended.

- **Attributable.** Every action carries a personal user ID. Never share an account or let someone else work in your session, and sign out when you leave your desk.
- **Legible and enduring.** Records are stored read-only with their SHA-256, and archived together with the software that opens them.
- **Contemporaneous.** Entries are written when the action happens, in UTC. Write down the reasons for your decisions at the time, in the study file, not afterwards.
- **Original.** The uploaded file is kept unchanged in the record. Analyse the original file, not a copy you edited.
- **Accurate.** The reproduction check and the reviewer's check support accuracy; values typed in by hand are checked against the source.
- **Complete and consistent.** Every run is in the trail, including those that were not used. Do not delete, rename or replace downloaded records to tidy up a study file.
- **Available.** An inspector account, the CSV export and the archive keep the data available for the whole retention period.

Things go wrong in real studies: a wrong file, a lost session, a record the reviewer rejects. Deal with them openly. Report anything that looks wrong with the installation to the system owner, and never work around the app, for example by analysing study data on the public instance.

What the App Does and What Your Organisation Does

A standard installation has no user accounts, no audit trail and no electronic signatures. Where your procedures allow a qualified standard installation for regulated work with hybrid records (see Status and Scope), put these controls in place around it:

- access to the qualified installation restricted to trained users;
- a log entry for every analysis: who ran it, when, on which data file (with its SHA-256), with which release, and for what purpose;
- a second person who checks the settings against the protocol and the results before they are used;

- archiving of the source files and the Analysis Records in a system with an audit trail.

A controlled installation provides personal accounts with roles, a password policy, an audit trail, stored records, review signatures, an audit trail page and archives (this chapter, Installing NCA Assistant for Regulated Use, and Appendix E). What stays with your organisation:

- the server as set up in Installing NCA Assistant for Regulated Use, reached only through two-factor sign-in from managed devices, with a monitored system log;
- a named system owner who is not an analyst or reviewer on the same studies, and root access only for IT;
- account management: an identity check before an account is given out, accounts created or reset only when the person signs in at once, reporting when admin does not work, a periodic review of accounts, and deactivation on the day someone leaves;
- training records, and a written policy that electronic signatures are binding (for submissions to the FDA, also the certification letter of 21 CFR 11.100(c));
- review of the audit trail at a fixed interval and before each study report, with extra reviews after failed sign-ins, a broken chain or repeated runs, and filing the trail head in your quality system;
- an investigation, with corrective action, of a broken chain or an unexplained gap between an app stop and start;
- producing the trail in an inspection: the CSV export of the study, the signature sheets, and on request the archive with its verification report, and an inspector account for the inspection;
- archives at the end of each study and when the server is retired, backups at a separate site, restore tests and a disaster recovery plan;
- a privacy notice for users, and patching, security testing and incident response for the server.

Limits you accept. Until its owner signs in, a new account can be taken over with the starting password admin by anyone who knows the user ID: create accounts only when the owner signs in at once. The audit trail shows any change, but someone with root access to the server could rebuild it; the heads you file, the signature sheets and the backups would reveal such a rebuild. Reviewers see every study on the server: separate sponsors need separate installations. A deactivated user's open session runs on until the 15-minute timeout, although signing is refused at once.

Upgrades, Data Protection and Problems

Upgrades and Change Control

Read the version history on the About page before you upgrade. Entries marked "Results can differ" and "Correctness fix" tell you which analyses a new version may change. A new release, a new R version or a new package version needs at least a targeted re-qualification (URS section 7); decide in advance whether completed analyses are repeated. On a controlled installation the audit trail continues across versions; each app start records the version and configuration in use.

Data Protection

Use a local installation for real study data. Pseudonymised data are still personal data under the GDPR (Article 4(5)), and pharmacokinetic data from a trial are health data (Article 9). Sending them to a hosted service needs a processor agreement and, for a host outside the EU, a transfer basis. Upload only synthetic, example or anonymised data to the public instance. The audit trail of a controlled installation holds personal data about its users: who did what, and when. Tell users in a privacy notice, and keep the trail as long as the records it belongs to.

Reporting Problems

Report problems at github.com/Robterheine/NCAassistant/issues, with a description of the data (never the data themselves, unless synthetic). Correctness fixes are listed in the version history under "Correctness fix", with the situations they affect.

Tutorial 8: From Analysis to Approved Record

Do this tutorial on a training installation, never on the production server: practice runs would end up in the real audit trail. A training installation in controlled mode runs on your own computer (Appendix E, A Training Installation). You need two accounts, an analyst and a reviewer; use two browsers, or a normal and a private window, to be signed in as both.

1. As the system owner, create an analyst and a reviewer account with `gxp/manage_users.R` add.
2. Sign in as the analyst and choose a new password. Open 4. One Subject at a Time, choose **Enter data manually (type values directly)**, click **Use This Data** and run the analysis with the settings of Tutorial 1.
3. Set **Min R² for half-life** to 0.9 and click **Run PK Analysis** again.
4. Type Tutorial 8 as Study name and click **Generate Analysis Record**. Read the line under the button.
5. Sign in as the reviewer in the other window. Open **Records**, select the record and read its details and the audit history of the data.
6. Click **Reject**, type the reason "Min R² changed from the protocol without a reason", and sign.
7. As the analyst, set **Min R² for half-life** back to 0.7, run the analysis again and create a new record.
8. As the reviewer, approve the new record, and download its signature sheet.

Questions: 1) What does the audit history of the first record show, and which setting changed? 2) Why is the data marked for the reviewer? 3) How many runs does the history of the second record list, and why? 4) What does the signature sheet say about who created the record and what the reviewer's signature means? 5) What happens when the reviewer, who also has the analyst role, wants to approve a record of their own?

Key Points

- Records under 21 CFR Part 11 or EU GMP Annex 11 need a controlled installation. A standard installation is used for regulated work only with hybrid records and the procedural controls of this chapter; the public app, never.
- Your user ID and password are your signature. Keep the password to yourself.
- Every run is recorded. Change a setting only for a reason, and write the reason down.
- Leave data out in the app, with a reason, never by editing the file.
- Someone else approves your record, after checking the settings against the protocol and the history of the data.
- If something looks wrong, stop and tell the system owner.

Installing NCA Assistant for Regulated Use

This chapter takes you from an empty server to an installation you can use for GCP and GxP work: controlled mode on, audit trail running, accounts issued, and the installation qualified and released by QA. It follows the order in which the work is done. Appendix E is the reference that goes with it, with the complete configuration files, the account commands and what to do when something goes wrong.

Three people's advice went into the steps. The regulatory reviewer's question is whether the evidence will hold up in an inspection or a submission; the system engineer's, whether the server, network and backups are safe. The R and Shiny programmer cares about one thing: that the app runs exactly the code and package versions that were validated. Their advice is in the steps where it applies.

The technical work takes days rather than weeks. The approvals and the performance qualification with your own data usually take longer.

What You Need Before You Start

People and their roles

Name these people before you begin, and write the names down in the installation record. One person can hold more than one role, with two exceptions: the system owner does not have root access to the server, and QA does not approve its own work.

- **System owner:** responsible for the installation and its use. Manages the accounts with `gxp/manage_users.R`, files the audit-trail heads and runs the periodic review.
- **IT administrator:** builds and patches the server, holds root access, runs the backups and the log monitoring. Does not use the app for analyses.
- **R administrator:** installs R and the validated package versions and checks that the app runs the released files unchanged. Often the same person as the IT administrator, sometimes a pharmacometrician with IT support.
- **Validation lead:** prepares and carries out the IQ/OQ/PQ protocol, and writes the summary report.
- **QA:** approves the plan, the protocol and the summary report, and releases the installation for use.
- **Users:** analysts, reviewers and inspectors, each with a personal account.

Documents and procedures

From the release you need the URS, the IQ/OQ/PQ protocol and this manual. From your own quality system you need, at minimum, procedures for:

- use of computerized systems, and training;
- account management (issuing, reviewing, deactivating);
- backup, restore and archiving;
- change control and periodic review;
- incidents and deviations;

- for studies that go to the FDA: the letter to the FDA certifying that your electronic signatures are the legally binding equivalent of handwritten ones (21 CFR 11.100(c)). It is sent once per organisation, before the signatures are used in a submission.

If a procedure does not exist yet, write it before go-live. An inspector will ask for it.

Two environments

Build two installations from this procedure: a test installation and a production installation. The test installation is where you run the manual tests, try upgrades and train users. Production holds only real work. Keep their controlled directories, keys and accounts separate. A test installation can be the training setup from Appendix E, but it is better on a server that matches production.

What controlled mode does not cover

Know the limits before you promise anything to a sponsor or a regulator:

- The app computes NCA parameters and average bioequivalence. It does not produce SDTM PP or ADaM ADPP datasets, XPT files or define.xml. For a submission, those are made by your own validated process from the app's results (app_results_reference.csv in the Analysis Record), with a reconciliation.
- The FDA scaled test for narrow therapeutic index drugs and the FDA model for replicate designs (separate Test and Reference variances and a subject-by-formulation interaction) are not available for analysis. EMA ABEL and FDA RSABE are.
- Open mode has none of the controls in this chapter. Use it for regulated work only with hybrid records and procedural controls, as Working in a Regulated Environment describes; never use the public app on shinyapps.io.

Step 1: Plan and Approve

1. Write down the intended use. Which studies, which parameters, which decisions rest on the results (for example "NCA of phase I studies for the clinical study report; average bioequivalence of 2×2 crossover studies"). Validation effort follows from it: the closer a result is to a regulatory decision, the more evidence you need.

2. Get one release, and only one. Use a tagged release from github.com/Robterheine/NCAassistant/releases, never a copy of the main branch. Write the version and the tag's commit in the installation record. A version number must stand for exactly one set of files: if you ever find two different sets of files with the same version, treat the later one as a new version under change control.

3. Read the URS and its risk assessment against your intended use. Add requirements of your own where your use needs them, and decide on the residual risks the product owner has accepted. One of them concerns you directly: every new account starts with the password admin, which the user must change at the first sign-in (FM-29 in the URS). Step 7 describes a procedure that closes this window in practice; record in the plan that you use it.

4. Plan the performance qualification. Choose data sets like your own studies (your routes of administration, steady state or single dose, BLQ patterns, designs). Choose an independent reference to compare with, for example results from another validated NCA program or a published data set, and fix the tolerances in advance (for example a relative

difference below 0.1% for AUC and Cmax). Your own comparison with an independent program is evidence the developer's tests cannot give you.

5. Fix the analysis defaults in your procedures or statistical analysis plans, because they change results: the minimum adjusted R^2 (the app's default of 0.70 is a floor; many plans use 0.80 to 0.90), the half-life rules, the BLQ rule, how flagged AUC to infinity values are reported, the trough window at steady state, and the criteria for excluding data.

6. Have QA approve the plan and the IQ/OQ/PQ protocol before anything is executed (protocol section 1.2). Tests carried out before approval do not count.

Step 2: Prepare the Server

These steps are for the IT administrator. The sizes are a starting point for a department with a handful of simultaneous users; check them against your own load.

1. Use a dedicated virtual machine with a supported Linux distribution that gets security updates (for example an Ubuntu LTS release or Red Hat Enterprise Linux). Two processor cores, 8 GB of memory and 50 GB of disk are enough to start. The controlled directory grows with every stored Analysis Record, so watch its disk space.

2. Harden it. Install only what the app needs. Allow incoming traffic on port 443 (and 80, which only redirects to 443), and nothing else from outside. Administrative access by SSH with keys only, from the IT network. Encrypt the disk. Switch on the operating system's audit logging (auditd or equivalent), so root actions are recorded outside the app. On Ubuntu, for example: `sudo ufw allow 443/tcp`, `sudo ufw allow 80/tcp`, `sudo ufw allow from <IT network> to any port 22`, then `sudo ufw enable`.

3. Put it behind your two-factor sign-in. Users reach the app only through your VPN or single-sign-on gateway with two factors, from managed devices. Set the browser policy on those devices so passwords are not saved: the app's password is part of the electronic signature.

4. Synchronise the clock with NTP. Every audit entry carries a UTC time from the server clock. Check with `timedatectl` (or the equivalent) that synchronisation is on (check IQ-GXP-09). Set the server's time zone to your own (for example `sudo timedatectl set-timezone Europe/Amsterdam`): the working hours in NCA_GXP_WORK_HOURS are read in that time zone. Audit entries stay in UTC.

5. Forward the system log to your monitoring, and name a person who receives the app's security alerts. The app writes them to the system log under the tag `nca-assistant` (lockouts, repeated failed signatures, attempts to go around the access rules).

6. Install the system programs: R in the version the release was validated with (the version in `validation/renv.lock`; 4.6.0 for this release), zip (Analysis Records are built with it), logger (from `util-linux` or `bsdutils`), Shiny Server (the open-source version) and nginx. On Debian and Ubuntu, R packages built from source also need `libssl-dev`, `cmake` and `libuv1-dev`. Install Shiny Server from the package for your distribution on Posit's download page. It starts the app only once R and the packages of Step 3 are in place.

7. Record everything in the installation record: host name, operating system and version, R version, Shiny Server and nginx versions, and who did what, on which date.

Step 3: Install the Release and the Validated Packages

These steps are for the R administrator. The server must run exactly the files and package versions that were tested, and nothing else.

1. **Download the release archive** of the tagged version and unpack it to `/srv/shiny-server/nca`. Make root the owner, and give the service account shiny read access only. Nobody edits files in this folder afterwards: installation check IQ-REL-01 compares every file with the release manifest, and a changed file fails it.

2. **Install the validated package versions into the site library**, as root, from the app folder:

```
cd /srv/shiny-server/nca
sudo Rscript -e 'install.packages("renv", repos = "https://cloud.r-project.org")'
sudo Rscript -e 'renv::restore(lockfile = "validation/renv.lock", library =
"/usr/local/lib/R/site-library", prompt = FALSE)'
```

`/usr/local/lib/R/site-library` is the site library on Debian and Ubuntu; on other systems, `Rscript -e '.Library.site'` prints it; use the first path it shows, or the one that `Rscript -e '.Library'` prints when it shows none.

Do not use `install_and_run.R` on the server: it also starts the app in a browser. Do not install packages in a personal library of the shiny account either, because R would then load those first. The lockfile includes shinymanager, DBI and RSQLite, which controlled mode needs (check IQ-GXP-01).

3. **Create the marker for controlled mode**, owned by root:

```
sudo touch /srv/shiny-server/nca/gxp/CONTROLLED
```

While this file exists, the app starts only in controlled mode. If the settings are missing, it refuses to start rather than fall back to open mode (check IQ-GXP-02).

Step 4: Configure Shiny Server and nginx

Appendix E, Example Configuration, has both files in full. The settings that matter:

1. **Shiny Server** (`/etc/shiny-server/shiny-server.conf`): `run_as shiny`; listen on `127.0.0.1` only, so that nobody can reach Shiny Server without going through nginx; and `app_idle_timeout 0`, so the app keeps running between users (otherwise every idle stop shows up as an exception in the audit trail). Check IQ-GXP-07.

2. **nginx** (`/etc/nginx/conf.d/nca.conf`): HTTPS only, with HTTP redirected; a certificate from your own certificate authority; TLS 1.2 or later. Three settings differ from nginx's defaults and are needed: WebSocket forwarding (the app talks to the browser over a WebSocket), `proxy_read_timeout 3600s` (so an open session is not cut off), and `client_max_body_size 60m` (the app accepts files up to 50 MB; nginx's default of 1 MB refuses most data files). Check IQ-GXP-06 and IQ-GXP-15.

On Red Hat Enterprise Linux with SELinux enforcing, nginx cannot connect to `127.0.0.1:3838` until you allow it with `sudo setsebool -P httpd_can_network_connect 1`. After the first start, check the SELinux log (`sudo ausearch -m avc`) for other refusals, for example on Shiny Server's log directory.

3. **Put the certificate's expiry date in your calendar** or monitoring. An expired certificate stops all work.

Step 5: Create the Controlled Directory and the Settings

The system owner and the IT administrator do this together.

1. **Create the controlled directory on the server's local disk**, never on a network share or a cloud-synchronised folder, because SQLite files can be damaged there:

```
sudo mkdir /srv/nca-gxp
sudo chown shiny:shiny /srv/nca-gxp
sudo chmod 0700 /srv/nca-gxp
```

It will hold the user store, the audit trail and the stored records. Only the service account may write to it (checks IQ-GXP-03 and IQ-GXP-04).

2. **Create a strong key for the user store**, for example with `openssl rand -base64 32`. Keep it in your password vault, with access for two people. Without it, an archive cannot be restored, even years from now.

3. **Create `/home/shiny/.Renviron`**, owner shiny, mode 0600, with `NCA_GXP_DIR`, `NCA_GXP_KEY`, `NCA_GXP_ORG` (your organisation's name, recorded in every audit entry) and `NCA_GXP_WORK_HOURS` (the working hours in the server's time zone; activity outside them is listed among the exceptions on the Audit trail page). Appendix E shows the file. The key is nowhere else on the server (check IQ-GXP-05).

4. **Add the sudoers rule** that lets the system owner run exactly `gxp/manage_users.R` as the service account, and nothing else (Appendix E; check IQ-GXP-10).

5. **Create the user store and the audit trail:**

```
sudo -H -u shiny Rscript /srv/shiny-server/nca/gxp/manage_users.R init
sudo -H -u shiny Rscript /srv/shiny-server/nca/gxp/manage_users.R verify
sudo -H -u shiny Rscript /srv/shiny-server/nca/gxp/manage_users.R head
```

The `-H` is needed: without it, R reads the system owner's own `.Renviron` and does not find the settings. `head` prints the anchor of the trail, a sequence number and a hash. File it in your quality system, outside the server (check IQ-GXP-11).

6. **Restart Shiny Server and reload nginx**. Open the app: it asks for a sign-in. The first `app_started` entry in the audit trail names the controlled directory and records the password policy in force (12 characters or more, a new password every 90 days, lockout after 5 failed sign-ins, sign-out after 15 minutes without activity).

Step 6: Protect the Audit Trail and the Records

The app makes the trail tamper-evident: every entry carries the hash of the one before it, and the database refuses changes and deletions. What it cannot do on its own is stop someone with root access from replacing the whole database. The following steps close that gap with your own controls.

1. **Back up the controlled directory every day** to a separate site and network, encrypted, and keep the backups for as long as your retention rules require. Restore one backup to another location and verify it (check IQ-GXP-13). Repeat the restore test at least once a year.

Make the copies consistent. An SQLite file copied while the app writes to it can be damaged in the copy. Back up a filesystem snapshot, or first copy the two databases with sqlite3 (package sqlite3) and back up those copies together with the records folder:

```
sudo mkdir -p /var/backups/nca
sudo sqlite3 /srv/nca-gxp/audit.sqlite ".backup '/var/backups/nca/audit.sqlite'"
sudo sqlite3 /srv/nca-gxp/users.sqlite ".backup '/var/backups/nca/users.sqlite'"
sudo rsync -a /srv/nca-gxp/records/ /var/backups/nca/records/
```

To check a backup, restore it into a folder the service account can read, and run verify on that folder:

```
sudo -H -u shiny Rscript /srv/shiny-server/nca/gxp/manage_users.R verify /srv/nca-restore-test
```

2. File the head of the trail regularly, outside the server: at least every week, at the end of every study, and before and after every change to the installation. Run `manage_users.R` head and keep the result in your quality system or on write-once storage. A later verify with the filed heads shows that no entry after them was removed or rewritten:

```
sudo -H -u shiny Rscript /srv/shiny-server/nca/gxp/manage_users.R verify 1520:ab12...
```

3. Test the security alert. Lock a test account with five wrong passwords and check that the alert reaches the named person through your log monitoring (check IQ-GXP-12).

4. Keep root access and app access apart. The IT administrator does not hold an app account with the reviewer role, and the system owner has no root access. Review the operating system's audit log for actions on `/srv/nca-gxp` and `/srv/shiny-server/nca` as part of the periodic review.

Step 7: Issue the Accounts

1. Check the identity of each person before you create an account, and check that their training is recorded. Give every person their own account; a user ID is never shared and never used again, even after deactivation.

2. Give the smallest role that fits: analyst, reviewer, both (analyst;reviewer), or inspector (read-only). A person who reviews and signs records should not also be the analyst of the same records; the app does not stop this, so your procedure does.

3. Close the starting-password window. A new or reset account has the password admin until its first sign-in. Create or reset an account only when the user is ready to sign in straight away, for example during a call or at the user's desk, and let them sign in and choose their own password at once. Then run `manage_users.R list` and check that the password has been changed. Do not create accounts in advance.

```
sudo -H -u shiny Rscript /srv/shiny-server/nca/gxp/manage_users.R add jsmith "Jane Smith" analyst "New analyst, training record TR-123"
sudo -H -u shiny Rscript /srv/shiny-server/nca/gxp/manage_users.R list
```

4. Deactivate an account on the day the person leaves or no longer needs access, and review all accounts at a fixed interval (for example every three months). Appendix E, Managing Accounts, lists all commands.

Step 8: Qualify the Installation

Follow the approved protocol. In outline:

1. Installation qualification. Run the automated suite on the server from the app folder. It writes its two result files into the validation folder, which only root may write to, so run it as root:

```
cd /srv/shiny-server/nca
sudo Rscript validation/validation.R
```

Move `validation_results.csv` and `validation_environment.txt` out of the app folder into the qualification file. They hold the results, the R and package versions and the hash of every tested file, and they are not part of the release manifest. IQ-REL-01 and IQ-REL-02 must pass: they show that the files and package versions are those of the release. Then carry out the server checks IQ-GXP-01 to IQ-GXP-15 in protocol section 2.1, and record the evidence (screenshots, command output) for each.

2. Operational qualification. The automated tests ran in the step above. Carry out the manual tests in protocol section 4 on the test installation, including the MAN-GXP tests for controlled mode, and record the outcome of each with name and date.

3. Performance qualification. Analyse the data sets chosen in Step 1 and compare the results with the independent reference, within the tolerances fixed in advance. Include every kind of analysis you will use: single dose and steady state, the BLQ rules you use, bioequivalence if you use it. Create an Analysis Record for each and check that `reproduction_check.txt` says MATCH.

4. Independent review. Have someone who did not carry out the tests review the evidence, including the PQ comparison. For work that goes to a regulator, an inspector will look for this review.

5. Deviations. Assess every failed test and every deviation from the protocol, including SUPPORTIVE failures, with its impact on the intended use (protocol section 7).

6. Summary report and release. The validation lead writes the summary report; QA approves it and releases the installation for its intended use, in writing (protocol section 8). The installation is not used for regulated work before this date.

Step 9: Go Live

1. Check that the procedures from "What You Need Before You Start" are in force and that every user has been trained on them and on this manual.

2. Issue the production accounts as in Step 7.

3. File a fresh head of the trail and record the go-live date and the released version in the installation record.

4. Tell users where to find help: this manual, the system owner, and the procedure for reporting problems.

Running the Installation

The system owner keeps a short calendar of recurring tasks. A suggested minimum:

- **Every working day:** the named person reads the security alerts from the log monitoring.
- **Every week:** run verify with the last filed heads and file a new head; check that the backups ran.

- **Every month:** a reviewer checks the Audit trail page for exceptions (failed sign-ins, activity outside working hours, restarts) and signs a trail review in the app.
- **Every three months:** review the accounts and deactivate those no longer needed.
- **Every year:** a periodic review of the system (open deviations, changes, incidents, account reviews, restore test, whether the intended use has changed), approved by QA; a restore test of a backup; the certificate's expiry date.
- **At the end of every study:** an archive of the study's records (see Retiring an Installation), and a filed head.

Changes and Upgrades

Every change goes through change control, with an impact assessment before it is made. What to repeat depends on what changes:

- **Operating system security updates** that do not touch R, the R packages, Shiny Server or nginx: install them in the normal patch cycle, then repeat the checks that could be affected (at least IQ-GXP-06, IQ-GXP-07, IQ-GXP-09 and IQ-GXP-15) and a short smoke test.
- **nginx, Shiny Server or certificate changes:** repeat the related IQ-GXP checks.
- **A new release of the app, or another R or package version:** this is a new version of the system. Install it first on the test installation, run the full automated suite, repeat the affected manual tests and a part of the PQ, and have QA release it. Archive the production installation before the upgrade and file a head. Records made with the old version stay reproducible, because each Analysis Record carries its own copy of the pipeline code and names its package versions.
- **Never edit a file in the app folder** on production, not even a comment. It fails IQ-REL-01 and breaks the link between the installation and its validation.

When Something Goes Wrong

Follow your incident procedure. Appendix E, When Something Is Wrong, describes what the app reports. In short:

- If verify, or Verify chain in the app, reports that the trail is not intact: stop regulated work on the installation, change nothing on the server, and investigate with the filed heads and the backups.
- If the app says an action could not be recorded in the audit trail, the action did not happen. Check the disk space and the permissions of the controlled directory before work goes on.
- A security alert about a locked account or repeated failed signatures: contact the user, and treat it as a possible attempt at misuse until you know otherwise.

Retiring an Installation

1. **Archive** with `manage_users.R` archive and a folder name. The archive holds the audit trail, the stored records, the app files and package lockfile that were running, the validation files, a verification report, restore steps (RESTORE.md) and a SHA-256 manifest of everything.

2. **Test the restore** on a clean computer, from a copy of the archive (check IQ-GXP-14; manual test MAN-GXP-01), and check that the trail verifies and the records open.
3. **Keep the archive, and the key in your vault**, for as long as your longest retention period for the records it holds.
4. **Decommission the server** under your IT procedure, and wipe its disks.

Checklist

Use this as a summary; the steps above have the details.

1. Roles named; procedures in force; test and production environments planned.
2. Intended use written; tagged release chosen; URS reviewed; residual risks accepted; PQ data sets, reference and tolerances chosen; analysis defaults fixed; plan and protocol approved.
3. Server built and hardened; two-factor access; NTP; log monitoring; system programs installed and recorded.
4. Release unpacked read-only; validated packages in the site library; gxp/CONTROLLED present.
5. Shiny Server on 127.0.0.1 with app_idle_timeout 0; nginx with HTTPS, WebSocket, read timeout and upload limit.
6. Controlled directory 0700 on local disk; key in the vault; .Renviron 0600; sudoers rule; trail created, verified, head filed.
7. Backups running and a restore tested; security alert tested; root and app roles separated.
8. Accounts issued with the starting-password procedure.
9. IQ, OQ and PQ carried out; evidence reviewed independently; deviations assessed; summary report approved; installation released by QA.
10. Go-live recorded; recurring tasks scheduled; change control in place.

Introduction to Non-Compartmental Analysis

What Is Non-Compartmental Analysis?

Non-compartmental analysis (NCA) is a method for describing the pharmacokinetics of a drug without assuming a specific structural model for how the drug distributes in the body. Instead of fitting a mathematical model, NCA calculates pharmacokinetic parameters directly from the observed concentration–time data using numerical integration and regression.

NCA is the standard approach for pharmacokinetic data in drug development. EMA, FDA and PMDA use it for bioequivalence studies, and it is routine in Phase I trials and dose-finding studies.

The Concentration–Time Curve

After a drug is administered, blood samples are collected at specified times and the drug concentration is measured. When plotted against time, the curve shows absorption (concentrations rising), distribution (transition), and elimination (concentrations declining). NCA extracts quantitative parameters from this curve without assuming its shape.

Key Pharmacokinetic Parameters

Parameter	Symbol	What it tells you
Peak concentration	C _{max}	The highest concentration observed. Relevant for efficacy and safety.
Time of peak	T _{max}	When the peak occurs. Reflects the rate of absorption.
AUC to last point	AUC _{0–t}	Exposure up to the last measurable concentration, from the observed data only. The primary AUC in bioequivalence.
AUC to infinity	AUC _{0–∞}	AUC _{0–t} plus the extrapolated tail, C _{last} /λ _z . Total exposure after a single dose.
AUC % extrapolated	%AUC _{extrap}	The share of AUC _{0–∞} that is extrapolated. Above 20% the value rests largely on the extrapolation.
Elimination rate constant	λ _z	The slope of the terminal phase on a log scale, from a regression on its last points.
Adjusted R ²	R ² _{adj}	How well the terminal regression fits; the app leaves λ _z -based parameters empty below a minimum.
Terminal half-life	t _{1/2}	ln 2/λ _z : the time for the concentration to halve during the terminal phase.

Apparent clearance	CL/F	Clearance divided by the bioavailability F: the volume of plasma cleared of drug per unit time.
Apparent volume of distribution	V _z /F	The terminal volume of distribution divided by F.
Mean residence time	MRT	The average time a drug molecule stays in the body; after an oral dose it includes the absorption time.
Partial AUC	AUC t ₁ –t ₂	Exposure over part of the profile, for example the first 30 minutes. It catches a difference in the shape of the curve that C _{max} and total AUC can miss.

Appendix A lists every parameter the app reports, with its unit.

Why NCA Is Used So Widely

NCA does not assume a compartmental structure, so the same calculations apply whether a drug follows one-compartment kinetics or multi-phasic elimination. Results depend on the observed data and the numerical method, not on model selection, which makes them reproducible across software and analysts.

Regulatory agencies accept NCA because the methods are well-defined and transparent. NCA is the standard approach for bioequivalence with pharmacokinetic endpoints, where C_{max} (rate) and AUC (extent) are the primary endpoints.

When to Use NCA vs. Compartmental Modelling

Use NCA when you need standard PK parameters from observed data: first-in-human studies, bioequivalence, drug interaction studies, dose proportionality, food effect studies. Use compartmental modelling when you need to simulate unobserved scenarios, when you have sparse sampling or irregular dosing, or when you need to understand mechanisms. The two approaches are complementary.

Theoretical Background

Trapezoidal Methods for AUC Calculation

The Linear Trapezoidal Rule

Between consecutive observations at t_i and t_{i+1} with concentrations C_i and C_{i+1} , the area is approximated as a trapezoid:

$$\text{AUC}(i) = (t_{i+1} - t_i) \times (C_i + C_{i+1}) / 2$$

This assumes linear concentration change. During elimination, drug concentrations typically decline exponentially, and the linear assumption leads to systematic overestimation: a straight line between two points on an exponential decay lies above the true curve.

The Logarithmic Trapezoidal Rule

If the concentration declines mono-exponentially between two points, the exact integral yields:

$$\text{AUC}(i) = (t_{i+1} - t_i) \times (C_i - C_{i+1}) / (\ln C_i - \ln C_{i+1})$$

This gives a more accurate estimate during the terminal phase. It requires both concentrations to be positive.

The Linear-Up / Log-Down Hybrid

Proposed by Yeh and Kwan (1978): the linear formula while concentrations rise and the logarithmic formula while they fall. It is the default in NCA Assistant. Choose the method in the protocol: ICH M13A (2.2.2.2) asks for it to be reported, and the two methods give slightly different results (in Tutorial 4, an AUC to last point ratio of 111.25% with linear-up/log-down and 111.16% with the linear method). A falling segment that ends at zero (a value set by a BLQ rule) has no logarithm, so the app integrates it with the linear formula.

Two Special Cases

A fall to zero. The logarithmic formula needs two positive concentrations, so it cannot be used when a BLQ sample set to 0 follows a measurable one. NonCompart then adds no area for that segment. The app uses the linear formula there instead, so the fall to zero is counted in AUC to last point, AUC_∞ and everything derived from them.

An IV bolus with a sample at time 0. After an IV bolus the concentration at time 0 is C_0 , the value just after the injection, not a pre-dose value. The app sets aside samples at or before time 0 and back-extrapolates C_0 log-linearly from the first two concentrations when they decline; otherwise C_0 is the first measured concentration.

Terminal Phase Characterisation and Half-Life Estimation

What Is the Terminal Phase?

After distribution is complete, the concentration declines as $C(t) = A \cdot e^{-\lambda_z t}$, where A is the intercept of the terminal phase. On a log-linear plot, this appears as a straight line with slope $-\lambda_z$. The half-life is $t_{1/2} = \ln(2) / \lambda_z$.

Automated Point Selection

NCA Assistant uses NonCompart's slope selection, the same for the results table and the Half-Life Review. It fits log-linear regressions to the last 3, 4, 5, ... points (after extravascular dosing only points after C_{max} are used, and never values set by a BLQ rule) and takes the fit with the highest adjusted R²; fits within 0.0001 of the best prefer more points. If that fit has an adjusted R² below the minimum (slider, default 0.70), the profile gets no half-life, and AUC_∞, CL/F, Vz/F and MRT are left empty too; a note names the profile. This happens when the terminal phase is poorly sampled, when late samples are near the LLOQ, or when elimination is multi-phasic. A lower minimum accepts noisier fits. Judging the Terminal Phase, in Reading Your Results, lists what to check before you report a half-life.

Manual Point Selection

Adjust the points yourself for multi-phasic elimination, enterohepatic recirculation or late samples near the LLOQ. In Half-Life Review, select a profile, tick the terminal-phase points and click Recalculate. Values set by the BLQ rule are not offered (the automatic fit does not use them either), and the label above the points says how many were left out. The slope must be negative, and although 2 points are accepted, a fit to be reported needs at least 3. NonCompart then recomputes half-life, AUC_∞, CL/F, Vz/F and MRT from those points, with its unit conversions. A hand-picked selection is used whatever its adjusted R², so judge the fit on the plot. Selections are kept when you run the analysis again and are listed in the Analysis Record. The half-life rules flag a hand-picked fit as well.

With flip-flop kinetics (absorption slower than elimination, as with many modified-release products) the terminal slope reflects absorption. No choice of points corrects that: report the value as an apparent half-life, and do not derive elimination from it. Only IV data can tell the two apart.

Derived Parameters

With λ_z and the observed last concentration C_{last}, NCA Assistant (through NonCompart) calculates:

$AUC_{0-\infty} = AUC_{0-t} + C_{last} / \lambda_z$ (observed); with the predicted last concentration from the regression, the "predicted" variants

$AUC \% \text{ extrapolated} = (AUC_{0-\infty} - AUC_{0-t}) / AUC_{0-\infty} \times 100$

$CL/F = \text{Dose} / AUC_{0-\infty}$ and $V_z/F = \text{Dose} / (\lambda_z \times AUC_{0-\infty})$, without /F after an IV dose

$MRT = AUMC_{0-\infty} / AUC_{0-\infty}$; after an infusion, minus half the infusion duration

The lag time T_{lag} is the last sampling time before the first measured concentration. After an IV bolus, AUC % back-extrapolated is the share of AUC between time 0 and the first sample, from the back-extrapolated C₀. The Statistical Methods page in the app gives the same formulas.

BLQ Handling: Pharmacological Rationale

Every assay has a lower limit of quantification (LLOQ). How BLQ samples are handled changes AUC. Setting early BLQ samples to 0 is usual before absorption; treating late BLQ samples as missing avoids adding area that was never measured. LLOQ/2 is a compromise that adds a little area wherever it is used. Rule 1, the default, follows the usual convention and ICH M13A, which treats BLQ values as zero. Rules 4 and 6 put LLOQ/2 into the absorption phase: that adds area during a lag, though the lag time itself is not changed,

because a value set by a rule is not a measured concentration. They belong in sensitivity analyses, or in a protocol that asks for them. Running two rules and comparing the results shows how much the choice matters. Beal (2001) compares methods for BLQ data in model-based (population) analysis, not in NCA.

Steady-State Considerations

At steady state AUC_{τ} replaces AUC_{∞} . You enter the dosing interval τ ; AUC_{τ} is the AUC from 0 to τ , interpolated between samples and extrapolated with λ_z when the last sample is missing or earlier than τ (not reported if that fit was rejected); the extrapolated share is the column AUC Within Dosing Interval % Extrapolated, with a note and, above the limit of the half-life rules, a flag. $C_{avg} = AUC_{\tau}/\tau$, fluctuation = $(C_{max} - C_{min})/C_{avg} \times 100$ and swing = $(C_{max} - C_{min})/C_{min}$. In fluctuation and swing, C_{max} is the highest observed concentration from 0 to τ . The C_{max} column is NonCompart's and covers the whole profile. When sampling goes on after τ and the highest value lies there, the two differ and a note names the profile; check the dosing times, because a sample after τ may follow the next dose. C_{min} is the lowest observed concentration from 0 to τ , the pre-dose sample included, and C_{τ} is the observed sample nearest to τ within the trough window; the two differ when the minimum falls after the dose. A BLQ result within the interval enters C_{min} at the value the BLQ rule gave it (0 or LLOQ/2). CL/F and V_z/F are calculated from AUC_{τ} . AUC_{∞} , its % extrapolated, $AUMC_{\infty}$, and the single-dose V_{ss} and MRT are not meaningful during repeated dosing and are left empty. A half-life from one interval is unreliable when the interval does not cover the terminal phase.

Partial AUCs: Interpolation and the End of the Interval

A partial AUC is the area over part of the profile. With both cutoffs on sampling times it is simply the sum of the trapezoids between them.

A cutoff between two samples has to be interpolated: linearly, or log-linearly on a falling segment when you chose linear-up/log-down. The two methods differ most where the curve bends, which is why a protocol that asks for a partial AUC should schedule a sample at the cutoff.

An interval can end at t , the time of the last measurable concentration. That time belongs to the profile, not to the sampling schedule: a later sample that came back below the limit of quantification does not extend it. The app calculates such an interval as AUC_{0-t} minus the AUC from 0 to the start, so an interval ending exactly at the last measurable time gives the same number.

Past that point nothing is reported. Extrapolating with λ_z would turn an observed exposure metric into a model-based one, and it would do so for the profiles that carry the least information.

Statistical Methods

The Two One-Sided Tests (TOST) Procedure

Conventional hypothesis testing ($H_0: \mu_T = \mu_R$) can only reject equality, not prove similarity. Schuirmann (1987) reformulated the problem as two one-sided tests:

$H_{01}: \mu(T)/\mu(R) \leq 0.80$ (ratio too low)

$H_{02}: \mu(T)/\mu(R) \geq 1.25$ (ratio too high)

Both rejected at $\alpha = 0.05 \rightarrow$ ratio lies within $[0.80, 1.25] \rightarrow$ bioequivalent. This is equivalent to: the 90% CI for the ratio falls entirely within $[80\%, 125\%]$. The limits are symmetric on the log scale: $\ln(0.80) = -0.2231$, $\ln(1.25) = +0.2231$.

The ANOVA Model for Crossover Designs

The standard model for a $2 \times 2 \times 2$ crossover:

$$\ln(Y) = \mu + \text{SEQ} + \text{SUBJ}(\text{SEQ}) + \text{PER} + \text{TRT} + \varepsilon$$

SEQ accounts for differences between sequence groups, SUBJ removes between-subject variability from the residual, PER accounts for changes between periods, and TRT is the effect of interest: its coefficient estimates the log-ratio of geometric means. NCA Assistant fits this model for all crossover designs, with Sequence, Subject, Period and Treatment as factors; with fixed effects this is EMA Method A.

Fixed-Effects vs. Mixed-Effects Models

Fixed effects (lm): all terms, including Subject, are fixed (EMA Method A). Subjects without both treatments do not contribute to the comparison. Residual df = $N - 2$ for a balanced $2 \times 2 \times 2$.

Mixed effects (nlme::lme): Subject is a random intercept, and denominator df use the containment method. In a 2×2 crossover and the paired comparison, subjects without both treatments are left out before either model is fitted, as the EMA guideline asks, and are counted in the column Subjects without both treatments. ICH M13A (2.2.3.2) allows a general linear or a mixed model; its primary analysis includes all data of the subjects with evaluable data for both products. In replicate designs the mixed model still uses subjects who received one treatment only, because their repeated administrations inform the within-subject variance. Use the model your protocol or statistical analysis plan specifies, for example the mixed model as a sensitivity analysis. If the mixed model was chosen and cannot be fitted, the app shows the fixed-effects result for information, says so in the Model column, and gives no verdict: the model is not changed after the data are seen. This is EMA Method B (one within-subject variance; results agree with replicateBE). It is not the FDA model for replicate designs, which has separate Test and Reference variances, a subject-by-formulation interaction and Satterthwaite degrees of freedom.

Confidence Interval Construction

$$90\% \text{ CI} = \exp(\bar{\delta} \pm t(0.95, \text{df}) \times \text{SE}(\bar{\delta})) \times 100\%$$

where δ is the treatment coefficient on the log scale, SE its standard error from the model, and df the denominator degrees of freedom. The limits are rounded to two decimals, with halves rounded up as in SAS, before comparison with the acceptance limits, as the EMA guideline (2010, section 4.1.8) and the FDA (2026) ask. Bioequivalence is concluded when the rounded interval lies within the limits (80.00–125.00% by default) and, for wider limits, the point estimate lies within 80.00–125.00%. A verdict is given only for the 90% interval, which corresponds to the two one-sided tests at $\alpha = 0.05$; at another level the app shows the interval without a verdict. Widened limits apply to Cmax only unless you choose otherwise (Path 6, Widened and Narrowed Limits).

Design-Specific Models

Paired comparison: $\ln(Y) = \mu + \text{SUBJ} + \text{TRT} + \varepsilon$, equivalent to a paired t-test; $df = N - 1$. Ratio and interval only, no verdict.

Parallel design: $\ln(Y) = \mu + \text{TRT} + \varepsilon$. $df = N_T + N_R - 2$. Differences between subjects enter the comparison, so more subjects are needed than in a crossover. The interval assumes equal variances in both groups (pooled variance). The results table also shows the Welch interval, which allows unequal variances, as supplementary information. When the analysis plan specifies an unequal-variance analysis, that is the primary result, not the app's verdict.

Covariates (parallel groups). When baseline covariates are chosen the model becomes $\ln(Y) = \mu + \text{TRT} + \text{covariates} + \varepsilon$: ordinary least squares, main effects only, pooled variance, and $N - 2 - p$ residual degrees of freedom, where p counts the covariate terms (a categorical covariate with k categories has $k - 1$). The treatment coefficient stays the log ratio of geometric means. The adjusted interval is primary, and the unadjusted pooled interval is shown beside it. ICH M13A (2.2.3.4) asks that stratification factors be accounted for in a parallel study, and the FDA (2026) says an applicant may consider prespecifying baseline prognostic covariates (see also Tsiatis et al., 2008). The covariates must be baseline characteristics named before the analysis. The app cannot check when a value was measured, and the Analysis Record lists the covariates. A missing value stops the analysis, so nobody is dropped or imputed.

Replicate designs: the same model, with Period at 3 or 4 levels; each administration enters once. Under the standard approach the confidence interval uses average bioequivalence limits; under EMA ABEL or FDA RSABE see Reference-Scaled Approaches. The within-subject variability table uses a separate model per treatment (Sequence, Subject and Period, fitted to subjects with that treatment at least twice), as in the replicateBE package; the app's results agree with replicateBE on its 30 reference data sets.

Tmax is compared as an untransformed difference with a t-based interval, without a verdict. Because Tmax takes only the sampling times, it is usually summarised as median and range and compared nonparametrically, outside the app.

Reference-Scaled Approaches (EMA ABEL and FDA RSABE)

For replicate designs the analyst chooses, before the run, between the standard approach, EMA ABEL and FDA RSABE. The $2 \times 3 \times 3$ and $2 \times 2 \times 4$ designs support both; the $2 \times 2 \times 3$ design supports ABEL only.

EMA ABEL (EMA 2010, section 4.1.10). C_{max} only; AUC is judged against 80.00–125.00%. CV_{wR} comes from the EMA model described under Replicate designs. Up to CV_{wR} 30% the limits stay at 80.00–125.00%. Above it they are $\exp(\pm 0.760 \cdot sWR)$, with $sWR = \sqrt{\ln(CV_{wR}^2 + 1)}$ and CV_{wR} (as a fraction) capped at 0.50, which gives 69.84–143.19% at the cap. Bioequivalence is concluded when the 90% confidence interval lies within the limits and the point estimate within 80.00–125.00%. The guideline asks the applicant to justify that the variability is reliable and not the result of outliers.

FDA RSABE (FDA 2026, Appendix G). Each log-transformed metric is assessed on the subjects with a value in every period. For each subject the administrations are put in period order and two contrasts formed: I, the mean of the Test values minus the mean of the Reference values, and D, the first minus the second Reference value. $I \sim \text{Sequence}$, by least squares, gives the estimate (the mean of the sequence means, each sequence weighted equally), its standard error and 90% limits from the t distribution with $n - m$ degrees of freedom (n subjects, m sequences). Let $x = \text{estimate}^2 - SE^2$ and bound_x the larger squared limit. $D \sim \text{Sequence}$ gives $s^2WR = MSE/2$ with $df(d)$ residual degrees of freedom. With $\theta = (\ln 1.25 / 0.25)^2 = 0.7967$, $y = -\theta \cdot s^2WR$ and bound_y = $y \cdot df(d) / \chi^2(0.95, df(d))$. The criterion bound is $(x + y) + \sqrt{((\text{bound}_x - x)^2 + (\text{bound}_y - y)^2)}$, the 95% upper confidence bound of the scaled criterion (Howe, 1974).

If sWR is 0.294 or higher the route is Scaled: RSABE is met when the criterion bound is at or below 0 and the point estimate, $100 \cdot \exp(\text{estimate})$, lies within 80.00–125.00%, rounded to two decimals at both limits. Below 0.294 nothing is scaled, and the ordinary test against 80.00–125.00% is used with the model you chose. The results then say that this model is not the FDA model for replicate designs (Appendix C), which has separate Test and Reference variances and Satterthwaite degrees of freedom. The limits implied by sWR are $\exp(\pm \ln 1.25 \cdot sWR / 0.25)$.

Two within-subject SDs are in play. sWR (EMA model) is the residual SD of the period-adjusted model fitted to the Reference data. sWR (FDA contrasts) comes from the difference between a subject's two Reference administrations. They are estimated differently and can differ, and the table names the one it shows.

The FDA recommends at least 24 evaluable subjects for RSABE; the app warns below 24 and still computes. The PowerTOST documentation reports an inflated type I error for the FDA method when the true CV_{wR} is at or below about 30%, around and below the switch. The partial replicate cannot estimate the Test variance, and with unequal variances its consumer risk can exceed 5% (Tothfalusi and Endrenyi, 2016). A subject who lacks a period leaves the RSABE contrasts. The app does not implement the exact method of Tothfalusi and Endrenyi or other corrections that neither agency has adopted.

Drug–Drug Interaction Analysis

DDI studies use the same statistics as BE: the ANOVA model, the geometric mean ratio and its 90% confidence interval. The difference is the interpretation: the ratio is the fold-change in exposure caused by the interacting drug, judged against no-effect boundaries for the victim drug (see Analysing Drug–Drug Interaction Studies) rather than BE acceptance limits.

ICH M12 on drug interaction studies (2024), published as guidance by EMA and FDA, asks for the geometric mean ratio of AUC and C_{max} (and C_{min} when relevant) with and without

the precipitant, with its 90% confidence interval (section 5.2). Crossover studies, including one-sequence studies, are preferred over parallel designs (section 3.2.1).

Power and Sample Size for Bioequivalence

Exact power via Owen's Q function for standard ABE. For ABEL, RSABE and NTID, power is simulated in PowerTOST: 100,000 simulated studies for the sample size, 10,000 for power and the power curve. The Bioequivalence path applies the same ABEL and RSABE criteria to a replicate study, so the planned and the analysed criterion agree; the NTID test is for planning only. Inputs: within-subject CV (for scaled methods, of Test and Reference), expected ratio, target power and design. A higher CV or a ratio further from 100% needs more subjects; a crossover needs fewer subjects than a parallel design.

Practical guidance: without data, a CV of 30% is a cautious assumption. The expected ratio defaults to 95% for standard bioequivalence, 90% for highly variable drugs and 97.5% for narrow therapeutic index drugs; target power 80%, or 90% for studies where failure is costly. Check the Sample Size vs. Variability tab to see how N changes with the CV.

Partial AUCs in Bioequivalence

Regulators ask for a partial AUC when two products can share C_{max} and total AUC while the shape of the curve differs. FDA product-specific guidances name the metric and the interval per product; Hopefl et al. (2025) list the guidances published between 2020 and 2024.

The statistics do not change. A partial AUC is log-transformed and analysed with the same model and confidence interval as C_{max} and AUC. The interval comes from the protocol; choosing it after seeing the data inflates the chance of a false conclusion of equivalence.

A pivotal metric has to meet the acceptance limits, a supportive one is reported without a verdict. With several pivotal metrics each has to pass on its own. That is an intersection-union test, so no correction for multiple testing applies.

Zeros need care. An early interval with only BLQ samples can be exactly 0, which has no logarithm. Leaving those profiles out would remove the lowest exposures and push the ratio up, so the app reports no estimate and no verdict for that metric and shows how many zero values each treatment had.

For every metric the table also counts, per treatment, the profiles that had no value, the profiles with a value of zero, and the profiles resting mainly on concentrations set by the BLQ rule. The last count matters because such a profile can pull a ratio on the log scale without ever being zero: under Rule 1 an early window may be exactly zero and get no verdict, while under Rule 6 the same sample becomes LLOQ/2 and the metric keeps a verdict that one imputed value helped to decide.

The EMA allows the acceptance criteria for a partial AUC to be widened in the same way as for C_{max} of a highly variable product (EMA/CHMP/EWP/280/96 Rev1, section 6.8.2.2). Under EMA ABEL the app widens C_{max} only. For a partial AUC it uses the limits you enter, which can be widened by hand when the protocol pre-specifies it. To widen C_{max} and the partial AUCs but not AUC to last point or to infinity, choose C_{max} and partial AUCs under Widened limits apply to. The same guideline recommends an early and a terminal partial

AUC, split at a pre-specified cut-off such as half the dosing interval, when a single-dose study replaces a multiple-dose study.

Plan such a study with the CV of the partial AUC itself. For loxapine inhalation powder the within-subject CV was above 30% for AUC 0–10 min but below 30% for AUC 0–30 min, which is why the 30-minute window became the pivotal metric (Hopefl et al., 2025).

Regulatory Guidelines

Which guideline applies. For immediate-release solid oral dosage forms, ICH M13A (in effect in the EU since January 2025) is the bioequivalence guideline. The EMA guideline of 2010 still applies to other dosage forms and to the widening of limits for highly variable drugs, which M13A does not cover; later parts of the ICH M13 series address such topics. In the USA, FDA guidances and product-specific guidances apply alongside M13A. Check the current status of each document before you cite it.

ICH M13A Bioequivalence for Immediate-Release Solid Oral Dosage Forms (Step 4, 23 July 2024; in effect in the EU from 25 January 2025; published by ICH, EMA and FDA)

EMA Guideline on the Investigation of Bioequivalence (CPMP/EWP/QWP/1401/98 Rev. 1, 2010): <https://www.ema.europa.eu/en/investigation-bioequivalence-scientific-guideline>

EMA product-specific bioequivalence guidance, and the EMA questions and answers on bioequivalence, which describe Methods A and B: on the EMA website

FDA Statistical Approaches to Establishing Bioequivalence (May 2026; replaces the 2001 guidance): <https://www.fda.gov/regulatory-information/search-fda-guidance-documents/statistical-approaches-establishing-bioequivalence>

FDA Bioequivalence Studies With Pharmacokinetic Endpoints for Drugs Submitted Under an ANDA, and **FDA product-specific guidances** for generic drugs, on the FDA website; **21 CFR Part 320** (bioavailability and bioequivalence requirements)

FDA Bioavailability Studies Submitted in NDAs or INDs: General Considerations (draft, 2022): <https://www.fda.gov/regulatory-information/search-fda-guidance-documents/bioavailability-studies-submitted-ndas-or-ind-s-general-considerations>

ICH M12 Drug Interaction Studies (2024; published by EMA and FDA): <https://www.ema.europa.eu/en/ich-m12-drug-interaction-studies-scientific-guideline>

EMA Guideline on the pharmacokinetic and clinical evaluation of modified release dosage forms (EMA/CHMP/EWP/280/96 Rev1, 2014)

ICH E6(R3) Good Clinical Practice (2025), and the **EMA Guideline on computerised systems and electronic data in clinical trials** (EMA/INS/GCP/112288/2023)

CDISC SDTM Controlled Terminology (release 2026-03-27, source of the parameter codes): <https://evs.nci.nih.gov/ftp1/CDISC/SDTM/>

Appendices

Appendix A: Output Parameters

Every column the NCA can produce, with its label in the app, its definition and its unit for a dose in mg, time in h and concentrations in ng/mL. The unit follows the units you select. "With λ_z " means the parameter needs an accepted terminal fit; "single dose" means it is left empty at steady state. Partial AUC columns (Partial AUC t1–t2, Cmax t1–t2, Tmax t1–t2) have the unit of AUC, concentration and time. CL/F, Vz/F and the bioequivalence comparison of AUC to infinity use the observed AUC0– ∞ .

Code	Label in the app	Definition	Unit (example)	Reported for
b0	Y-Intercept (ln scale)	Intercept of the terminal regression on the ln scale	(none)	All, with λ_z
C _{MAX}	Peak Concentration (C _{max})	Highest observed concentration	ng/mL	All
T _{MAX}	Time of Peak (T _{max})	Time of C _{max}	h	All
T _{LAG}	Lag Time	Last sampling time before the first measured concentration	h	Extravascular
CL _{ST}	Last Measurable Concentration	Last measurable concentration (C _{last})	ng/mL	All
CL _{STP}	Predicted Last Concentration	C _{last} predicted by the terminal regression	ng/mL	All, with λ_z
T _{LST}	Time of Last Measurable Conc	Time of C _{last} (t)	h	All
LAM _{ZHL}	Half-Life	Terminal half-life, $\ln 2/\lambda_z$	h	All, with λ_z
LAM _Z	Elimination Rate Constant	Terminal elimination rate constant λ_z	1/h	All, with λ_z
LAM _{ZLL}	Lambda_z Lower Time	First time point of the terminal regression	h	All, with λ_z
LAM _{ZUL}	Lambda_z Upper Time	Last time point of the terminal regression	h	All, with λ_z
LAM _{ZNPT}	Points Used for Half-Life	Number of points in the terminal regression	(none)	All, with λ_z
CORRXY	Correlation (r)	Correlation coefficient of the terminal regression	(none)	All, with λ_z

R2	R-squared	R ² of the terminal regression	(none)	All, with λz
R2ADJ	Adjusted R-squared	Adjusted R ² of the terminal regression; the minimum applies to it	(none)	All, with λz
LZSPAN	Half-Life Span Ratio	Time from the first to the last point of the terminal regression, divided by the half-life	(none)	All, with λz
FLAGS	Half-Life Flags	The half-life rules the fit breaks, in words; none, no half-life or rules off when it breaks none	(none)	All, with λz
FLAG_SPAN, FLAG_AUCPE, FLAG_AUCPBE	Flag: ...	1 when the fit breaks the span, extrapolation or back-extrapolation rule, 0 when it does not; empty when the rule is off or does not apply	(none)	All, with λz (FLAG_AUCPE: AUC $_{\infty}$ single dose, AUC $_{\tau}$ at steady state; FLAG_AUCPBE IV bolus)
AUCLST	AUC to Last Point	AUC from 0 to Tlast (AUC0–t)	ng/mL·h	All
AUCALL	AUC All (incl. trailing zero)	AUC from 0 to the last sample, including a trailing zero	ng/mL·h	All
AUCIFO	AUC to Infinity (observed)	AUC0–t + Clast/ λz	ng/mL·h	Single dose, with λz
AUCIFP	AUC to Infinity (predicted)	AUC0–t + Clast,pred/ λz	ng/mL·h	Single dose, with λz
AUCPEO	AUC % Extrapolated (observed)	Share of AUC0– ∞ (observed) that is extrapolated, %	(none)	Single dose, with λz
AUCPEP	AUC % Extrapolated (predicted)	Share of AUC0– ∞ (predicted) that is extrapolated, %	(none)	Single dose, with λz
AUMCLST	AUMC to Last Point	Area under the first-moment curve ($C \times t$) from 0 to Tlast	ng/mL·h ²	All

AUMCIFO	AUMC to Infinity (observed)	AUMC extrapolated to infinity with Clast	ng/mL·h ²	Single dose, with λz
AUMCIFP	AUMC to Infinity (predicted)	AUMC extrapolated to infinity with Clast,pred	ng/mL·h ²	Single dose, with λz
AUMCPEO	AUMC % Extrapolated (obs)	Share of AUMC0–∞ (observed) that is extrapolated, %	(none)	Single dose, with λz
AUMCPEP	AUMC % Extrapolated (pred)	Share of AUMC0–∞ (predicted) that is extrapolated, %	(none)	Single dose, with λz
VZFO	Apparent Volume (Vz/F)	Dose/(λz × AUC0–∞); at steady state from AUCτ	L	Extravascular, with λz
VZFP	Apparent Volume (pred)	As Vz/F, with the predicted AUC0–∞	L	Extravascular single dose
CLFO	Apparent Clearance (CL/F)	Dose/AUC0–∞; at steady state Dose/AUCτ	L/h	Extravascular
CLFP	Apparent Clearance (pred)	As CL/F, with the predicted AUC0–∞	L/h	Extravascular single dose
MRTEVLST	Mean Residence Time (to last)	AUMC0–t/AUC0–t, including absorption time	h	Extravascular single dose
MRTEVIFO	Mean Residence Time (to inf, obs)	AUMC0–∞/AUC0–∞ (observed), including absorption time	h	Extravascular single dose
MRTEVIFP	Mean Residence Time (to inf, pred)	As above, with the predicted values	h	Extravascular single dose
CMAx_DN	Dose-Normalised Cmax	Cmax divided by the dose of that profile	ng/mL per mg	With dose normalisation
AUCLST_DN	Dose-Normalised AUC Last	AUC0–t divided by the dose	ng/mL·h per mg	With dose normalisation
AUCIFO_DN	Dose-Normalised AUC Inf	AUC0–∞ (observed) divided by the dose	ng/mL·h per mg	With dose normalisation
AUCIFP_DN	Dose-Normalised AUC Inf (pred)	AUC0–∞ (predicted) divided by the dose	ng/mL·h per mg	With dose normalisation

AUMCLST_DN	Dose-Normalised AUMC Last	AUMC _{0-t} divided by the dose	ng/mL·h ² per mg	With dose normalisation
AUMCIFO_DN	Dose-Normalised AUMC Inf	AUMC _{0-∞} (observed) divided by the dose	ng/mL·h ² per mg	With dose normalisation
AUMCIFP_DN	Dose-Normalised AUMC Inf (pred)	AUMC _{0-∞} (predicted) divided by the dose	ng/mL·h ² per mg	With dose normalisation
AUCTAU	AUC Within Dosing Interval	AUC from 0 to τ ; extrapolated with λ_z when the last sample is before τ	ng/mL·h	Steady state
AUCTAU_PCTEXT	AUC Within Dosing Interval % Extrapolated	Share of AUC τ extrapolated with λ_z past the last sample, %; 0 when sampling reaches τ	(none)	Steady state
TAU	Dosing Interval (tau)	The dosing interval you entered	h	Steady state
CAVG	Average Concentration (Cavg)	AUC τ / τ	ng/mL	Steady state
CMIN_SS	Minimum Concentration (Cmin)	Lowest observed concentration from 0 to τ , pre-dose included	ng/mL	Steady state
CTAU_SS	Concentration at Tau (Ctau)	Observed sample nearest to τ within the trough window (default $\pm 10\%$ of τ , at most 1 time unit); empty when there is none	ng/mL	Steady state
CTAU_TIME	Time of Ctau	Time of the sample used as C τ	h	Steady state
FLUCTP	Peak-Trough Fluctuation (%)	(C _{max} - C _{min})/C _{avg} × 100, C _{max} and C _{min} from 0 to τ	(none)	Steady state
SWING	Swing ((C _{max} -C _{min})/C _{min})	(C _{max} - C _{min})/C _{min} , C _{max} and C _{min} from 0 to τ	(none)	Steady state
C0	Initial Concentration (C0)	Concentration at time 0, back-extrapolated from	ng/mL	IV bolus

		the first two samples		
AUCPBEO	AUC % Back-Extrapolated (observed)	Share of AUC _{0-∞} (observed) between 0 and the first sample, %	(none)	IV bolus single dose
AUCPBEP	AUC % Back-Extrapolated (predicted)	As above, with the predicted AUC _{0-∞}	(none)	IV bolus single dose
VZO	Volume of Distribution (V _z)	Dose/(λ _z × AUC _{0-∞}); at steady state from AUC _τ	L	IV, with λ _z
VZP	Volume of Distribution (pred)	As V _z , with the predicted AUC _{0-∞}	L	IV single dose
CLO	Clearance (CL)	Dose/AUC _{0-∞} ; at steady state Dose/AUC _τ	L/h	IV
CLP	Clearance (pred)	As CL, with the predicted AUC _{0-∞}	L/h	IV single dose
MRTIVLST	Mean Residence Time IV (to last)	AUMC _{0-t} /AUC _{0-t} (minus half the infusion time)	h	IV single dose
MRTIVIFO	Mean Residence Time IV (to inf, obs)	AUMC _{0-∞} /AUC _{0-∞} (observed; minus half the infusion time)	h	IV single dose
MRTIVIFP	Mean Residence Time IV (to inf, pred)	As above, with the predicted values	h	IV single dose
VSSO	Volume at Steady State (V _{ss})	Volume at steady state, MRT × CL	L	IV single dose
VSSP	Volume at Steady State (pred)	As V _{ss} , with the predicted values	L	IV single dose
EXCL	Excluded by Analyst (1 = yes)	1 for a profile you left out; its result is kept but not used in summaries or bioequivalence	(none)	All

Appendix B: Glossary

Term	Definition
------	------------

ABE	Average bioequivalence (80–125% limits by default)
ABEL	Average bioequivalence with expanding limits (EMA, for highly variable drugs): for C _{max} , limits that widen with the within-subject variability of the Reference above CV _{wR} 30%, capped at 50%. Chosen under Acceptance approach in replicate designs.
Adjusted and unadjusted interval	Adjusted: the confidence interval from a model that contains covariates. Unadjusted: the interval from the model with Treatment only. The adjusted interval is the primary result, and the unadjusted one is shown beside it.
Adjusted R ²	Goodness of fit of the terminal regression, corrected for the number of points; the minimum in the app applies to it
ADNCA	CDISC ADaM dataset for non-compartmental analysis
ANOVA	Analysis of variance
AUC	Area under the concentration–time curve
AUC % extrapolated	The share of AUC to infinity that lies beyond the last measurable concentration
AUC to infinity (AUC _{0–∞})	AUC to last point plus C _{last} /λ _z
AUC to last point (AUC _{0–t})	AUC from 0 to the last measurable concentration
AUC _τ	AUC within the dosing interval (steady state)
Audit trail	On a controlled installation: the record of who did what, when and on which data, kept on the server; entries cannot be changed through the app, and any change made otherwise shows
Baseline characteristic	A characteristic of a subject that is known before the first dose, such as age, weight or sex. Only baseline characteristics are suitable as covariates.
BLQ	Below the limit of quantification
C _{avg}	Average concentration over the dosing interval at steady state: AUC _τ / τ
CL/F	Apparent clearance: clearance divided by the bioavailability F
C _{max}	Maximum observed concentration
C _{min} (steady state)	Lowest observed concentration from 0 to τ, the pre-dose sample included
Controlled installation	NCA Assistant on your organisation's server in controlled mode, with sign-in, audit trail, stored records and review signatures
Covariate	A subject characteristic that may influence a PK parameter and is included in the model. NCA Assistant offers baseline covariates for parallel-group studies.

Criterion bound	In RSABE, the 95% upper confidence bound of the scaled criterion. RSABE requires it to be at or below 0.
Crossover	A design in which each subject receives the treatments in turn, in periods separated by a washout
C _T	Concentration at the end of the dosing interval: the observed sample nearest to τ within the trough window
CV%	Coefficient of variation
CV _{wR} / CV _{wT}	Within-subject coefficient of variation of the Reference / Test
DDI	Drug–drug interaction
EMA Method A / B	The fixed-effects (A) and mixed-effects (B) ANOVA of the EMA questions and answers on bioequivalence
Exclusion register	The list of samples and profiles left out of an analysis, each with its reason, author and time; it can be restored, downloaded and loaded again
Extravascular	Given other than into a vein (oral, intramuscular, subcutaneous), so the drug is absorbed first
GCP / GxP	Good clinical practice; GxP is the collective name for the good-practice rules (clinical, laboratory, manufacturing) that apply to regulated work
Geometric mean	$\exp(\text{mean of the } \ln \text{ values})$; the usual summary for AUC and C _{max}
Geometric CV	$100 \cdot \sqrt{(\exp(\text{SD}^2 \text{ of the } \ln \text{ values}) - 1)}$
GMR	Geometric mean ratio (Test/Reference)
Half-life rules and flags	Checks on each terminal fit (span of the fitted points, % of AUC extrapolated or back-extrapolated) that flag a fit without changing any value
Head (of the audit trail)	The sequence number and fingerprint of the last entry, filed outside the server so that a later check can show that no entry was removed or rewritten
HVD	Highly variable drug (within-subject CV of the Reference above 30%)
Hybrid records	Records kept partly electronically and partly on paper, with the signatures on paper
ICH M13A	ICH guideline on bioequivalence for immediate-release solid oral dosage forms (2024)
Inspector	Account that can see records, the audit trail and the user overview, but cannot sign or change anything
Intersection-union test	A test that passes only if each of several tests passes; no multiplicity correction is needed

Intra-subject contrast	In the FDA RSABE method, a difference of log values within one subject: the Test mean minus the Reference mean (I), and the first minus the second Reference administration (D).
IQ / OQ / PQ	Installation, operational and performance qualification: the tests that show a system is installed correctly, works as specified, and does what your own use needs
Point estimate constraint	The requirement, for widened or scaled limits, that the point estimate also lies within 80.00–125.00%.
Route	In a scaled run, Scaled when the method scaled the limits for that metric, and Standard when the ordinary 80.00–125.00% test applied (for example sWR below 0.294, or AUC under ABEL).
Scaled limit	The acceptance limits implied by the variability of the Reference: $\exp(\pm 0.760 \cdot \text{sWR})$ for ABEL, $\exp(\pm \ln 1.25 \cdot \text{sWR} / 0.25)$ for RSABE.
Standardized difference	A measure of imbalance between groups on a baseline characteristic: the difference in means divided by the pooled SD (for a category, the difference in proportions divided by the pooled SD). It describes the groups and is not a test.
Switch	The value of sWR at which RSABE starts to scale, 0.294 (about CVwR 30%). Below it the ordinary test applies.
sWR	The within-subject standard deviation of the Reference on the log scale. The EMA model and the FDA contrasts estimate it differently, so the table says which one it shows.
λ_z	Terminal elimination rate constant
LLOQ	Lower limit of quantification
MRT	Mean residence time
MSE	Mean squared error: the residual variance of the ANOVA model
NCA	Non-compartmental analysis
NTID	Narrow therapeutic index drug
Open mode	The app without controlled mode: no sign-in, audit trail or electronic signatures
Owen's Q	The function PowerTOST uses for exact power of the two one-sided tests
Partial AUC	AUC over part of the profile, for example AUC 0–30 min; the interval comes from the protocol or a product-specific guidance
PE	Point estimate (geometric mean ratio)
Period	One administration in a crossover, with its own profile
Pivotal metric	A bioequivalence metric that has to meet the acceptance limits; it receives a verdict

Power	The probability that a study shows bioequivalence when the products are bioequivalent as assumed
Precipitant (perpetrator)	The drug that causes a drug interaction
Profile	The concentrations of one subject under one treatment in one period
Reference	The treatment in the denominator of the ratio, chosen by the user
Reviewer	Account that approves or rejects records made by others and signs reviews of the audit trail
RSABE	Reference-scaled average bioequivalence (FDA, for highly variable drugs): when the within-subject SD of the Reference (sWR) is 0.294 or higher, the criterion scales with sWR, and the point estimate must also lie within 80.00–125.00%. Chosen under Acceptance approach in 2×3×3 and 2×2×4 designs.
Sequence	The order in which a subject receives the treatments, for example TR or RT
SHA-256	A 64-character fingerprint of a file; it changes if a single character in the file changes
Signature sheet	Page showing a record's signatures: printed names, UTC times, meanings and the audit entries
Starting password	The password admin that a new or reset account has until its first sign-in, when the user must replace it
Supportive metric	A bioequivalence metric reported with its ratio and confidence interval, without a verdict
System owner	The person who manages accounts and archives on the server
t (last measurable)	The time of the last measurable concentration of a profile; an interval written as 4 h–t ends there
$t_{1/2}$	Terminal elimination half-life
Tmax	Time of the maximum observed concentration
TOST	Two one-sided tests
Trough window	The time around τ within which a sample counts as C_{τ} ; by default $\pm 10\%$ of τ , at most 1 time unit
URS	User Requirement Specification: the list of what the system must do, with its risk assessment
Victim	The drug whose pharmacokinetics a drug interaction changes
Vz/F	Apparent terminal volume of distribution: the volume divided by F
Washout	The time between periods of a crossover, long enough for the drug to disappear
Welch interval	A confidence interval for a parallel design that does not assume equal variances in the two

	groups; shown next to the pooled interval as supplementary information
--	--

In this manual, "missing" and "excluded" for a BLQ value mean the same: the value is left out of the calculation.

Appendix C: Recognised Column Names

Meaning	Recognised names (case-insensitive)
Subject	Subject, SubjectID, USUBJID, ID, Pat, PatID, Proband, Teilnehmer
Time	TAD (preferred over TIME), Time, Hours, Hour, Tpt, NTIM, APTS, Zeit, Tid
Concentration	Conc, Concentration, DV, Cp, Result, Konz, Plasma, or a name containing ng/mL or ug/L
Treatment	Treatment, Trt, Formulation, Drug, Arm, Behandlung
Period	Period, Per, Prd, Phase, APERIOD, Occasion, OCC (not Visit)
Sequence	Sequence, Seq, Grp
Dose	Dose, Amt, Amount, Dosis

When no name matches, the app falls back to a column by position, so always check the suggested mapping. When the file has another column that could be the time or the concentration, a note says which one was chosen and names the other.

Appendix D: Troubleshooting

Opening and Installing

“Disconnected from the server”: the public app was idle too long, the analysis exceeded memory or time limits, or the app met an unexpected error. Reload the page, upload the file again and click Process Data; results that were not downloaded are lost. If it happens again, run the app locally and check the console.

The installation stops with an error: see Installing the App on Your Own Computer.

Uploading and Processing

Comma decimals not recognised: set Decimal point to comma (and the Delimiter, often semicolon) in the upload settings.

A decimal point in a decimal-comma file: the file mixes decimal marks (a value such as 0.250 also counts). Correct the values, or set the decimal mark to period if the file uses points throughout.

“N BLQ entries detected but LLOQ is not set”: set the LLOQ and click Process Data again; for “<X” entries, the button Set LLOQ to X fills in the value (Tutorial 2). If the “<” entries give different limits, there is no button: take the LLOQ from the assay report.

“Profile(s) do not start near time zero”: time is probably measured from the first dose of the study. Use the time since the dose of each period (in CDISC terms ARRLT or NRRLT, not AFRLT).

“N negative time values detected”: only one pre-dose sample per profile, no earlier than 10% of the profile's duration before the dose, is analysed at time 0. Check whether time counts from the dose of each period.

“A weight column for doses per kg needs a Dose column”: map the Dose column as well, or set Dose per kg: weight column back to (no).

Subject IDs that restart in each sequence: give every subject an ID that is unique in the study, for example TR-01 and RT-01.

“The data give concentrations in ...”: the unit column of the file says one unit and the analysis settings another. Select the unit the file states, or correct the file.

Large sampling gap warning: for studies longer than 48 h the threshold is 48 h, so daily trough samples are not flagged; for shorter studies it is 24 h. A gap close to the threshold is usually normal sampling.

“Refused: ...” in the CDISC ADNCA upload: the message says why (for example derived records, several units, AFRLT as time, or AVAL missing without a BLQ result). Fix the dataset upstream or make the choice the message asks for.

“... exclusion(s) match no sample or profile of these data”: the register names profiles or times that are not in this file, for example after the columns were mapped differently. Restore those exclusions in the Exclusions card, or map the columns as before.

NCA

Unit error before the NCA runs: the combination of dose, time and concentration units is not supported, or molar units are combined with a mass dose without a molecular weight. Select supported units and enter the molecular weight.

“NCA failed”: the message gives the reason, usually no profile with at least 2 positive concentrations, or a mapping problem. Check the mapping and the quality report.

Half-life is empty: the automatic fit is below the minimum adjusted R^2 , or fewer than 3 usable points remain. Review the profile in Half-Life Review, or lower the minimum if your analysis plan allows it.

Half-life differs from other NCA software: check the minimum R^2 , the point selection in Half-Life Review, the BLQ rule and the trapezoidal method; programs differ in how they choose the terminal-phase points.

My partial AUC is empty: the interval reaches past the last measurable concentration of that profile, or starts after it. The note above the results names the profile and its last measurable time. Partial AUCs are not extrapolated. In bioequivalence such an interval gives no estimate and no verdict for that metric. The profiles that would drop out are the low-exposure ones, so the ratio of the others would be biased.

Bioequivalence

“Choose the Reference treatment”: the run is blocked until you choose the Reference treatment in the Bioequivalence settings. The app pre-selects it only for unambiguous names (R, Ref, Reference, Comparator, Innovator, Originator, RLD).

No verdict: see When There Is No Verdict, in Path 6. The Bioequivalent? column gives the reason.

No verdict for my partial AUC: the interval is marked supportive, or the metric has a value of zero in at least one profile. The table gives the reason and the number of zero values per treatment.

Route says Standard although I chose RSABE: for that metric sWR is below 0.294, so nothing is scaled and the ordinary test against 80.00–125.00% applies. This is not an error. Under ABEL the same happens for AUC, which the EMA never widens, and for Cmax when CVwR is 30% or lower.

Warning about fewer than 24 subjects: the FDA recommends at least 24 evaluable subjects for RSABE. The app computes the result anyway; discuss the sample size.

Subjects left out of the RSABE contrasts: a subject who lacks a value in one or more periods cannot form the two contrasts. The number of subjects used is in the downloads (the Subjects columns and the sheet BE_Scaled).

The run stopped with a covariate message: the message names the column and the subjects. The usual causes are a missing value, text such as “72 kg” or a decimal comma, two values for one subject and a constant column. Fix the file and upload it again, or remove the covariate.

My approach was cleared: changing any setting clears the results. The approach stays selected unless the design no longer offers it.

“Treatment column must have exactly 2 levels”: see Studies with Three or More Treatments.

“The analysis was not run”: the selected design does not fit the data. The message says what the data have, for example four sequences, and what the design expects. Choose the design that matches the study, or check the Period and Sequence columns on the Upload page. If the sequence labels carry a group, as in TR-G1 and TR-G2, the study was run in several groups, which the app does not model.

Figures

Download Figure is grey: it becomes active once the first plot has appeared; wait for it.

Time points are missing from my summary plot: the summary plot averages per time value. With actual sampling times (0.95 h and 1.05 h for a nominal 1 h), use a copy of the file with nominal times for the figure; keep the actual times for the NCA.

“N observations with concentration ≤ 0 excluded”: zeros, for example BLQ values set to 0 by Rules 1, 2 or 5, cannot be log-transformed, so they are left out of the geometric mean at that time point. They stay in your data.

Too many colours to distinguish: try the Colorblind-safe (Okabe-Ito) palette. With more than eight groups, switch to Grayscale and annotate the exported figure in a graphics program.

Controlled Installation

I cannot sign in: check your user ID; after five wrong passwords the account is locked, and only the system owner can unlock it (reset).

admin does not work for my new account: someone may have used it first. Tell the system owner at once.

I think someone knows my password: change it at once (Change password in the menu with your name) and tell the system owner.

I forgot my password: ask the system owner for a reset, then sign in with admin and choose a new password.

I was signed out while working: there was no activity for 15 minutes. Upload the file again; records you created are stored on the server.

"This action was not carried out because it could not be recorded in the audit trail": nothing was done and nothing was lost. Tell the system owner.

My record was rejected: read the reason on the Records page, change the analysis, and create a new record.

I signed the wrong record: a signature cannot be undone. Tell the reviewer or the system owner; the analyst creates a new record, which is reviewed again. The mistaken signature stays visible in the audit trail.

Signature INVALID: the stored copy has changed since it was signed. Tell the system owner at once.

Records and Exports

The reproduction check says DIFFERENT: a file in the record was changed after it was made, a parameter is present on one side only, or the results differ. `reproduction_check.txt` gives the reason.

The record holds a CSV instead of my Excel file: the uploaded file was no longer available; see Complete Analysis Record.

Appendix E: Setting Up and Administering a Controlled Installation

For IT and the system owner. This appendix goes with the chapter Installing NCA Assistant for Regulated Use, which has the steps in order. Here you find the configuration files in full, the account commands, archiving, what to do when something is wrong, and a training installation. The installation checks of the IQ/OQ/PQ protocol (section 2.1) test the server against that chapter and these files, so follow them as written. Controlled mode is qualified per installation: a changed step is a deviation.

Example Configuration

Adapt the names, paths and address to your server. Lines starting with `#` are comments.

`/etc/shiny-server/shiny-server.conf:`

```
run_as shiny;
server {
    listen 3838 127.0.0.1;          # only nginx on the same server can reach it
    location / {
```

```

    site_dir /srv/shiny-server;
    log_dir /var/log/shiny-server;
    app_idle_timeout 0;           # keep the app running between users
  }
}

```

/etc/nginx/conf.d/nca.conf (the app is then at <https://nca.example.org/nca/>):

```

map $http_upgrade $connection_upgrade { default upgrade; '' close; }

server {
    listen 80;
    server_name nca.example.org;
    return 301 https://$host$request_uri;
}

server {
    listen 443 ssl;
    server_name nca.example.org;
    ssl_certificate /etc/ssl/certs/nca.crt;
    ssl_certificate_key /etc/ssl/private/nca.key;
    ssl_protocols TLSv1.2 TLSv1.3;
    client_max_body_size 60m;      # the app accepts files up to 50 MB

    location / {
        proxy_pass http://127.0.0.1:3838/;
        proxy_http_version 1.1;
        proxy_set_header Upgrade $http_upgrade;
        proxy_set_header Connection $connection_upgrade;
        proxy_set_header Host $host;
        proxy_set_header X-Forwarded-Proto $scheme;
        proxy_read_timeout 3600s;  # do not cut off an open session
        proxy_buffering off;
    }
}

```

/home/shiny/.Renviron (owner shiny, mode 0600):

```

NCA_GXP_DIR=/srv/nca-gxp
NCA_GXP_KEY=<a long passphrase, kept only here and in your password vault>
NCA_GXP_ORG=Your Organisation
NCA_GXP_WORK_HOURS=07:00-19:00

```

/etc/sudoers.d/nca (edit with visudo -f /etc/sudoers.d/nca; replace jdoe by the system owner's account, and /usr/bin/Rscript by the path that which Rscript shows):

```

jdoe ALL=(shiny) /usr/bin/Rscript /srv/shiny-server/nca/gxp/manage_users.R *

```

After a change to these files, restart Shiny Server and reload nginx. Then check that the app starts in controlled mode: the audit trail's first app_started entry names the controlled directory, and a file of 5 MB or more uploads without an error.

Managing Accounts

The system owner runs every command as: `sudo -H -u shiny Rscript /srv/shiny-server/nca/gxp/manage_users.R` command. Every change is logged under the system owner's own name, with the reason given. There is no delete: a user ID is never used again.

Command	What it does
add id "Printed Name" roles "reason"	Creates an account with the starting password admin. Roles: analyst, reviewer,

	analyst;reviewer or inspector. Create it only when the person can sign in at once
role id roles "reason"	Changes the roles
reset id "reason"	Sets the password back to admin and unlocks the account; the user must sign in at once and choose a new password
deactivate id "reason"	Blocks sign-in and signing; the account and its history stay
list	Shows every account, its roles, status and history
verify [folder] [seq:hash ...]	Checks the audit trail; verify folder checks an archived copy, also on another computer. Add the heads you filed, for example verify 1234:<hash>, to check that those entries are still there: without them, entries removed from the end of the trail do not show.
head	Prints the last entry of the trail as seq:hash: file it in your quality system at a fixed interval, and give the filed heads to verify
archive folder	Writes a read-only archive (below)

Before you give out an account, check the person's identity. Deactivate accounts on the day someone leaves or no longer needs access, and review the account list at a fixed interval.

Archiving and Restoring

manage_users.R archive folder writes a dated, read-only folder with a consistent copy of the audit trail, the stored records, the app files that are running, validation/renv.lock, the release manifest, the validation files, a verification report, restore steps (RESTORE.md) and a SHA-256 manifest of everything. Archive at the end of a study and when the server is retired, and keep the archive as long as the records. The restore steps bring the archived release back on a clean computer, on a copy of the archive; test this once during installation qualification.

When Something Is Wrong

If verify, or Verify chain in the app, reports that the trail is not intact, stop using the installation for regulated work, keep the server as it is, and start an investigation under your procedure. Compare the trail with the heads you filed and with the backups to find out what changed. If the app says an action could not be recorded in the audit trail, check the disk space and the permissions of the controlled directory.

A Training Installation

For training and for Tutorial 8, run controlled mode on your own computer. Use a separate copy of the app folder for it: the settings go in a file in that folder, so the app runs in controlled mode only when you start it from there, and your normal installation stays as it is. A training installation is not a qualified setup, because you own the folder that holds its audit trail.

On macOS and Linux:

1. Install the app as described in Installing the App on Your Own Computer. The installation script also installs shinymanager, DBI and RSQLite.
2. Copy the app folder, for example to ~/nca-training-app, and create an empty folder for the controlled data, for example ~/nca-training-data. Keep both on your local disk, not in OneDrive, Dropbox or a network folder: the audit trail can be damaged there.
3. In ~/nca-training-app, create a text file named .Renviron with these three lines (use the full path of your data folder):

```
NCA_GXP_DIR=/Users/yourname/nca-training-data
NCA_GXP_KEY=training-only-passphrase
NCA_GXP_ORG=Training
```

4. Open a terminal, go to the app folder and create the audit trail and two accounts:

```
cd ~/nca-training-app
Rscript gxp/manage_users.R init
Rscript gxp/manage_users.R add trainee1 "Trainee One" analyst "training"
Rscript gxp/manage_users.R add trainer1 "Trainer One" reviewer "training"
```

5. Start the app from the same folder: `Rscript -e "shiny::runApp(launch.browser = TRUE)"`. (In RStudio, open the folder as a project with File > New Project > Existing Directory: R then starts in this folder and reads its .Renviron file.)
6. Sign in as trainee1 with admin and choose a new password. Open a second browser, or a private window, and sign in there as trainer1. You can now do Tutorial 8.

On Windows the steps are the same, with these differences. Use forward slashes in the path in .Renviron (NCA_GXP_DIR=C:/Users/yourname/nca-training-data). Create .Renviron in the app folder with Notepad: choose All files as the file type, so it is not saved as .Renviron.txt. Run the commands in a Command Prompt in the app folder; if Rscript is not found, use its full path, for example "C:\Program Files\R\R-4.6.0\bin\Rscript.exe". Security alerts are written to security_alerts.log in the data folder, because Windows has no logger.

After the training, delete both folders. They hold only training data; your normal installation is untouched.

Appendix F: References

Beal SL. Ways to fit a PK model with some data below the quantification limit. *J Pharmacokinet Pharmacodyn*. 2001;28(5):481–504.

CDISC. SDTM Controlled Terminology, release 2026-03-27. NCI Enterprise Vocabulary Services.

Davit BM, Chen ML, Conner DP, et al. Implementation of a reference-scaled average bioequivalence approach for highly variable generic drug products by the US Food and Drug Administration. *AAPS J*. 2012;14(4):915–924. doi:10.1208/s12248-012-9406-x.

European Medicines Agency. Guideline on the investigation of bioequivalence. CPMP/EWP/QWP/1401/98 Rev. 1, 2010.

European Medicines Agency. Guideline on the pharmacokinetic and clinical evaluation of modified release dosage forms. EMA/CHMP/EWP/280/96 Rev1. Adopted 20 November 2014, effective 1 June 2015.

European Medicines Agency. Guideline on computerised systems and electronic data in clinical trials. EMA/INS/GCP/112288/2023.

Fuglsang A, Schütz H, Labes D. Reference datasets for bioequivalence trials in a two-group parallel design. *AAPS J.* 2015;17(2):400–404. doi:10.1208/s12248-014-9704-6.

Food and Drug Administration. Statistical Approaches to Establishing Bioequivalence. Guidance for industry, May 2026 (Appendix G: reference-scaled average bioequivalence).

Hopefl R, et al. A 2024 update on US FDA implementation of partial area under the curve into bioavailability and bioequivalence assessments. *Clin Pharmacol Ther.* 2025;117:1185–1193.

Howe WG. Approximate confidence limits on the mean of $X+Y$ where X and Y are two tabled independent random variables. *J Am Stat Assoc.* 1974;69:789–794.

International Council for Harmonisation. ICH E6(R3): Guideline for Good Clinical Practice. Step 4, 2025.

International Council for Harmonisation. ICH M12: Drug Interaction Studies. Step 4 guideline, 21 May 2024.

International Council for Harmonisation. ICH M13A: Bioequivalence for Immediate-Release Solid Oral Dosage Forms. Step 4 guideline, 23 July 2024.

Kim H, Han S, Cho YS, Yoon SK, Bae KS. Development of R packages: ‘NonCompart’ and ‘ncar’ for noncompartmental analysis (NCA). *Transl Clin Pharmacol.* 2018;26(1):10–15.

Labes D, Schütz H. replicateBE: Average Bioequivalence with Expanding Limits (ABEL). R package version 1.1.3 (used to test the replicate-design results).

Labes D, Schütz H, Lang B. PowerTOST: Power and Sample Size for (Bio)Equivalence Studies. R package version 1.5.7.

Owen DB. A special case of a bivariate non-central t-distribution. *Biometrika.* 1965;52(3/4):437–446.

Pinheiro JC, Bates DM. *Mixed-Effects Models in S and S-PLUS*. Springer, 2000.

R Core Team. R: A Language and Environment for Statistical Computing. R Foundation for Statistical Computing, Vienna, Austria; 2026 (version 4.6.0).

Schuirmann DJ. A comparison of the two one-sided tests procedure and the power approach. *J Pharmacokinet Biopharm.* 1987;15(6):657–680.

Smith BP, Vandenhende FR, DeSante KA, et al. Confidence interval criteria for assessment of dose proportionality. *Pharm Res.* 2000;17(10):1278–1283.

Tothfalusi L, Endrenyi L. An exact procedure for the evaluation of reference-scaled average bioequivalence. *AAPS J.* 2016;18(2):476–489. doi:10.1208/s12248-016-9873-6.

Tsiatis AA, Davidian M, Zhang M, Lu X. Covariate adjustment for two-sample treatment comparisons in randomized clinical trials: a principled yet flexible approach. *Stat Med.* 2008;27(23):4658–4677. doi:10.1002/sim.3113.

Yeh KC, Kwan KC. A comparison of numerical integrating algorithms by trapezoidal, Lagrange, and spline approximation. *J Pharmacokinet Biopharm.* 1978;6(1):79–98.

Appendix G: Document Control

Document ID NCA-UM-001. The PDF of this manual is part of each release of the app; its SHA-256 is listed in validation/release_manifest.csv, so an installation can be checked to hold the manual of its release.

Version	Date	Author	Changes
1.0	March 2026	Rob ter Heine	Initial release
1.1	April 2026	Rob ter Heine	Added theoretical background (Part V), expanded statistical methods (Part VI), DDI study analysis, decimal separator guidance
1.2	April 2026	Rob ter Heine	Quick Start data preparation callout, grid plot fix documentation, single-subject Analysis Record, DDI interpretation thresholds, regulatory DDI guideline links, troubleshooting updates
1.3	April 2026	Rob ter Heine	New chapter 3 (Visualize Data); chapters 3–5 renumbered to 4–6; hub layout updated to 3+3 grid; attribution updated
1.4	April 2026	Rob ter Heine	Factual corrections (Rule 6 label, AUC extrapolation flag); new feature documentation (steady-state in BE, balance check, design summary, Apply LLOQ button, excluded profiles note); beginner clarity additions (LLOQ/BLQ definitions, extravascular explained, R ² failure consequences, SHA-256 explained, half-life review how-to, AUC % Extrapolated meaning)
1.5	September 2026	Rob ter Heine	Updated for NCA Assistant v1.4.0: CDISC ADNCA upload and parameter codes, safety checks on upload, decimal comma and semicolon files, doses per period, replicate designs and the within-subject variability table, choosing the Reference treatment, the minimum R ² rule, planning with Test and Reference CV, the reproduction check in the Analysis Record. Exercise 1 answers corrected; guideline references updated.
1.6	September 2026	Rob ter Heine	Updated for NCA Assistant v1.5.0: partial AUCs in Paths 4, 5 and 6, with intervals ending

			at a time or at the last measurable concentration, Cmax and Tmax within an interval, pivotal and supportive metrics in bioequivalence, and interval shading in Visualize Data. New Exercise 4, new sections in Parts V and VI, and new glossary, troubleshooting and reference entries.
1.7	September 2026	Rob ter Heine	Updated for NCA Assistant v1.6.0, after a review of the whole app. New: intended use and data notice; ICH M13A (Cmax and AUC to last point compared by default, BLQ values kept out of the half-life, checks listed with the results, no verdict where the guideline excludes one); widened limits for Cmax only; Rules 4 and 6 changed; IV bolus samples at time 0 set aside; thousands separators in decimal-comma files; unique subject IDs; units taken from the data; settings kept and results cleared; file integrity in the reproduction check; the Welch interval; excluding subjects. Exercise answers checked against v1.6.0.
1.8	25 September 2026	Rob ter Heine	For NCA Assistant v1.7.0, after a review by a tutor, a clinical pharmacologist and a regulatory compliance officer. New: Getting Started with local installation, seven tutorials with an answer key, study types (IV bolus and infusion, single and multiple ascending dose, food effect, metabolites, bioavailability, three or more treatments), Reading Your Results, a chapter on regulated use, derived-parameter formulas, an appendix of output parameters, and troubleshooting by stage. Corrected throughout for the app changes of v1.7.0.
1.9	7 October 2026	Rob ter Heine	For NCA Assistant v1.9.0. New: chapter Working on a Controlled Installation (sign-in, records for review, review signatures, signature sheets, the audit trail page), Tutorial 8, Appendix E on setting up and

			administering a controlled installation, troubleshooting for controlled installations. The chapter on regulated use describes the controls of a controlled installation and the organisation's responsibilities. Revised within v1.8.0: exclusions with a reason (Leaving Out Data), half-life rules, the sensitivity analysis without exclusions, and one-click example data in the tutorials; after an adversarial review of the app: doses per kg, the trough window for C_T, AUC_T extrapolation, complete subjects in Method B, pre-dose samples at a negative time, loading the exclusion register again; new chapter Installing NCA Assistant for Regulated Use, with Appendix E as its reference; the chapters Working on a Controlled Installation and Using NCA Assistant in a Regulated Environment merged into one chapter, Working in a Regulated Environment, with new sections on roles, the audit trail, accounts, a study step by step, inspections and good practice; after a review by a PhD student, a senior pharmacometrician, a system engineer and a regulatory compliance officer: How to Use This Manual, 14 glossary terms, which installation for which work, a worked steady-state example, reporting without flagged values, an optional Tutorial 3 question, and server notes on backups, TLS, SELinux, the firewall and the time zone; ten screenshots (Figures 1 to 10). Also added to v1.8.0: the Welch interval beside the pooled interval for parallel groups, checked against the eleven published datasets of Fuglsang et al. (2015); the citation of the NCA Assistant article; covariate adjustment for parallel groups; EMA ABEL and FDA RSABE for replicate designs (Acceptance approach); Tutorials 4b and 4c with two
--	--	--	--

			example files and four figures; and the scope statement, glossary, troubleshooting and references revised. Also: the Analysis Record of a bioequivalence run recomputes the confidence intervals and the verdict, and a design that does not fit the data is refused (Complete Analysis Record, Path 6, Appendix D).
--	--	--	--